

BFS - unweighted

Shortest Path

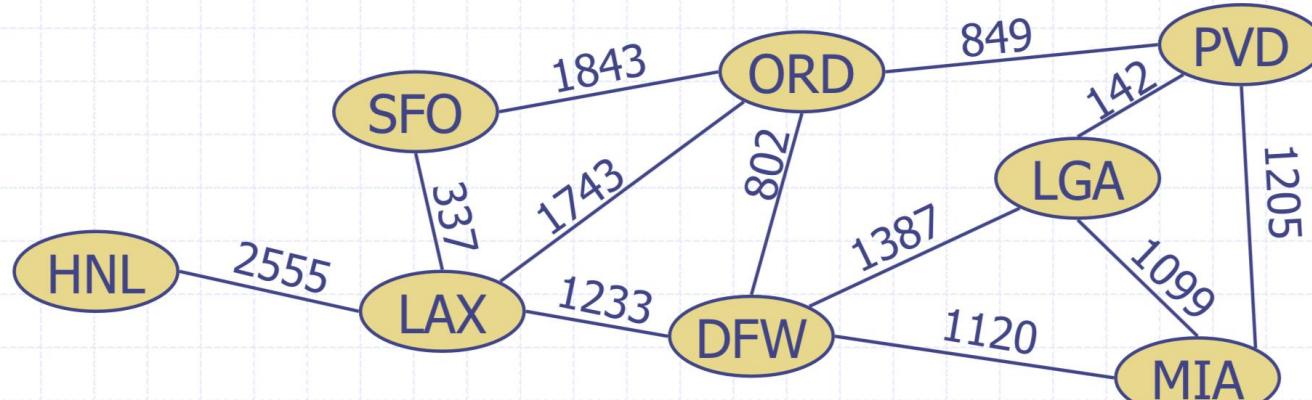
- Weighted Graphs
- Shortest Path Problem
- Dijkstra's Algorithm

→ DAG-based

Weighted Graphs



- ◆ In a weighted graph, each edge has an associated numerical value, called the weight of the edge
- ◆ Edge weights may represent, distances, costs, etc.
- ◆ Example:
 - In a flight route graph, the weight of an edge represents the distance in miles between the endpoint airports

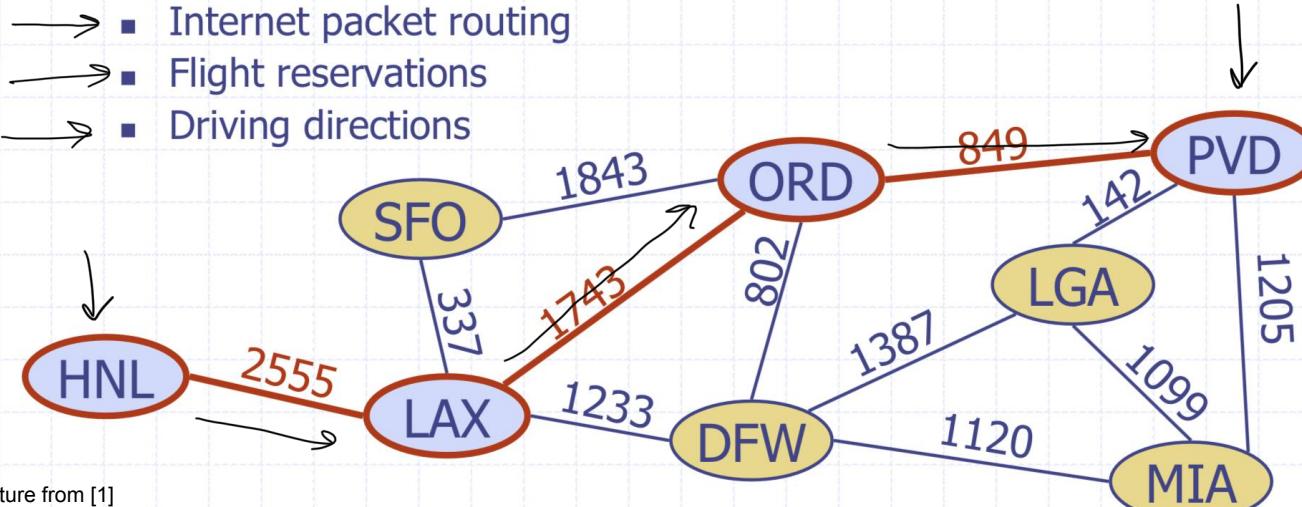


Picture from [1]

Shortest Path Problem



- ◆ Given a weighted graph and two vertices u and v , we want to find a path of minimum total weight between u and v .
 - Length of a path is the sum of the weights of its edges.
- ◆ Example:
 - Shortest path between Providence and Honolulu
- ◆ Applications
 - ▪ Internet packet routing
 - ▪ Flight reservations
 - ▪ Driving directions



Shortest Path Properties



Property 1:

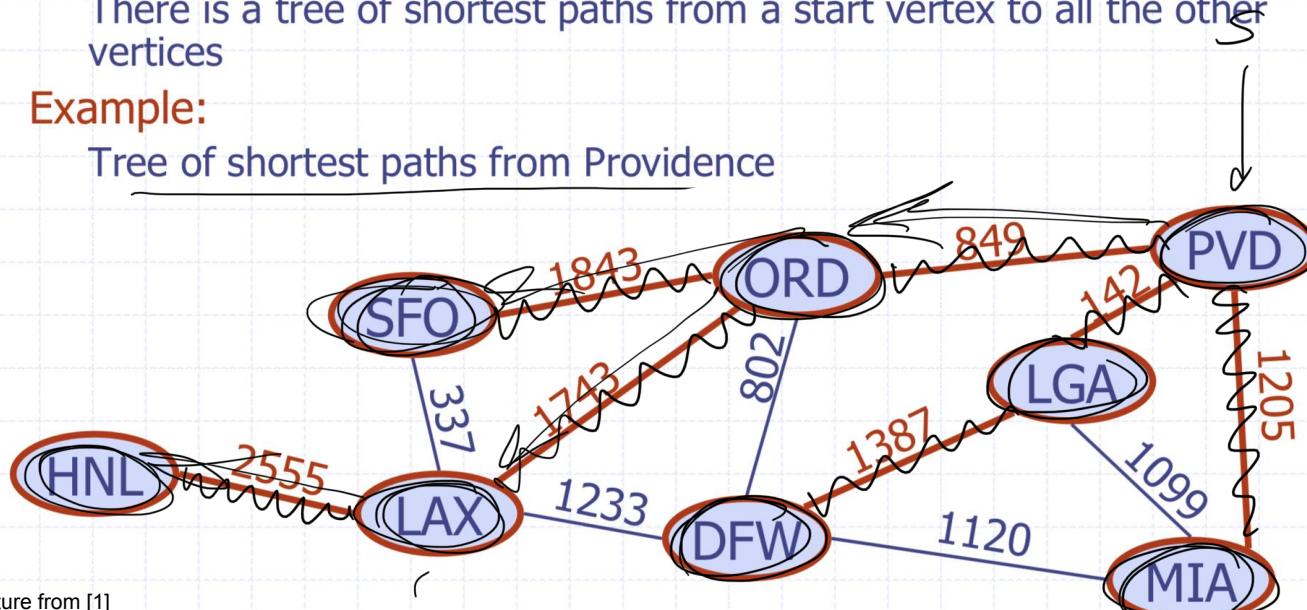
A subpath of a shortest path is itself a shortest path

Property 2:

There is a tree of shortest paths from a start vertex to all the other vertices

Example:

Tree of shortest paths from Providence



Dijkstra's Algorithm: A Greedy Approach

- The *distance* of a vertex v from a vertex s is the length of the shortest path from s to v
 - Dijkstra's Algorithm computes the distances of all the vertices from the start vertex s
 - Assumptions:
 - the graph is connected
 - the edge weights are nonnegative
- We grow a “cloud” of vertices beginning with s that will eventually cover all vertices.
 - We keep track of the current shortest distance from s to each vertex in the subgraph made up of the cloud and its adjacent vertices.
 - At each step, we consider a new vertex adjacent to the “cloud.”

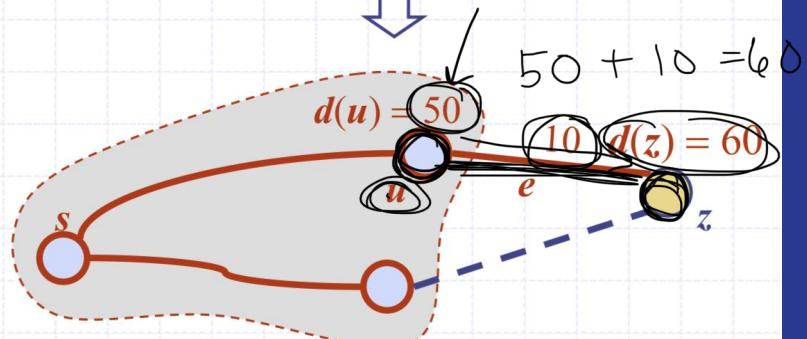
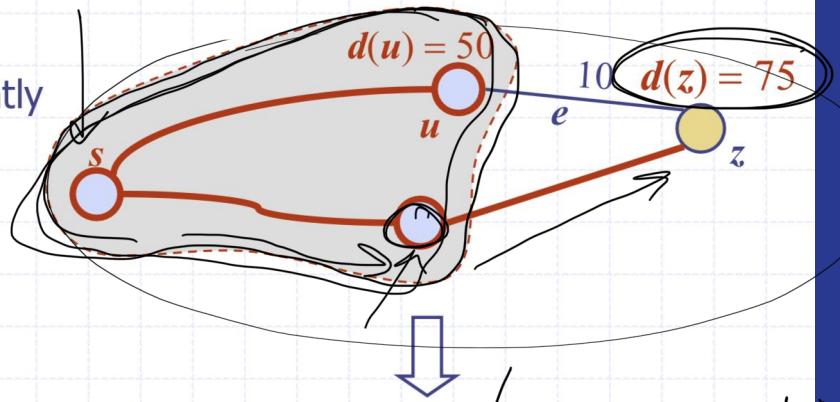
Edge Relaxation



- ◆ Consider an edge $e = (u, z)$ such that
 - u is the vertex most recently added to the cloud
 - z is not in the cloud

- ◆ The relaxation of edge e updates distance $d(z)$ as follows:

$$d(z) \leftarrow \min\{d(z), d(u) + \text{weight}(e)\}$$



Algorithm *dijkstrasShortestPath*(G, s) **Input:** a weighted graph $G = (V, E)$

with non-negative edge weights and a source vertex s

Output: the paths with minimum total weight from s to all other vertices in

G

$\boxed{dist} :=$ an array of size $|V|$

$\boxed{prev} :=$ an array of size $|V|$

$\boxed{Q} :=$ a min-heap-based priority queue

$dist[s] \leftarrow 0$

-for each vertex v in G :

 if $v \neq s$

$dist[v] \leftarrow \infty$

$prev[v] \leftarrow -1$

$Q.add(dist[v], v)$

while $!Q.isEmpty()$:

$u \leftarrow Q.extractMin()$

 for each vertex v in $G.adj(u)$ that is still in Q

$alt \leftarrow dist[u] + length(u, v)$

 if $alt < dist[v]$

$dist[v] \leftarrow alt$

$prev[v] \leftarrow u$

$Q.set(alt, v)$

return $dist, prev$

- **dist** array keeps track of the shortest distance from **s** to each vertex
- **prev** array keeps track of the previous vertex on the shortest path to each vertex
- **Q** provides an efficient way to extract the vertex with the current minimum path

Algorithm *dijkstrasShortestPath*(G, s) **Input:** a weighted graph $G = (V, E)$ with non-negative edge weights and a source vertex s
Output: the paths with minimum total weight from s to all other vertices in G

$dist :=$ an array of size $|V|$

$prev :=$ an array of size $|V|$

$Q :=$ a min-heap-based priority queue

$dist[s] \leftarrow 0$

for each vertex v in G :

if $v \neq s$

$dist[v] \leftarrow \infty$

$prev[v] \leftarrow -1$

$Q.add(dist[v], v)$

while $!Q.isEmpty()$:

$u \leftarrow Q.extractMin()$

for each vertex v in $G.adj(u)$ that is still in Q

$alt \leftarrow dist[u] + length(u, v)$

if $alt < dist[v]$

$dist[v] \leftarrow alt$

$prev[v] \leftarrow u$

$Q.set(alt, v)$

return $dist, prev$

- We set the source vertex distance to 0.
- All the other distances are set to INFINITY.
- We set the previous values to -1 (or something else to indicate they are UNDEFINED).
- We add all the vertices into the priority queue with the distances as the keys.

Algorithm *dijkstrasShortestPath*(G, s) **Input:** a weighted graph $G = (V, E)$ with non-negative edge weights and a source vertex s
Output: the paths with minimum total weight from s to all other vertices in G

$dist :=$ an array of size $|V|$

$prev :=$ an array of size $|V|$

$Q :=$ a min-heap-based priority queue

$dist[s] \leftarrow 0$

for each vertex v in G :

if $v \neq s$

$dist[v] \leftarrow \infty$

$prev[v] \leftarrow -1$

$Q.add(dist[v], v)$

while $!Q.isEmpty()$:

$u \leftarrow Q.extractMin()$

for each vertex v in $G.adj(u)$ that is still in Q

$alt \leftarrow dist[u] + length(u, v)$

if $alt < dist[v]$

$dist[v] \leftarrow alt$

$prev[v] \leftarrow u$

$Q.set(alt, v)$

 update

edge relaxation

return $dist, prev$

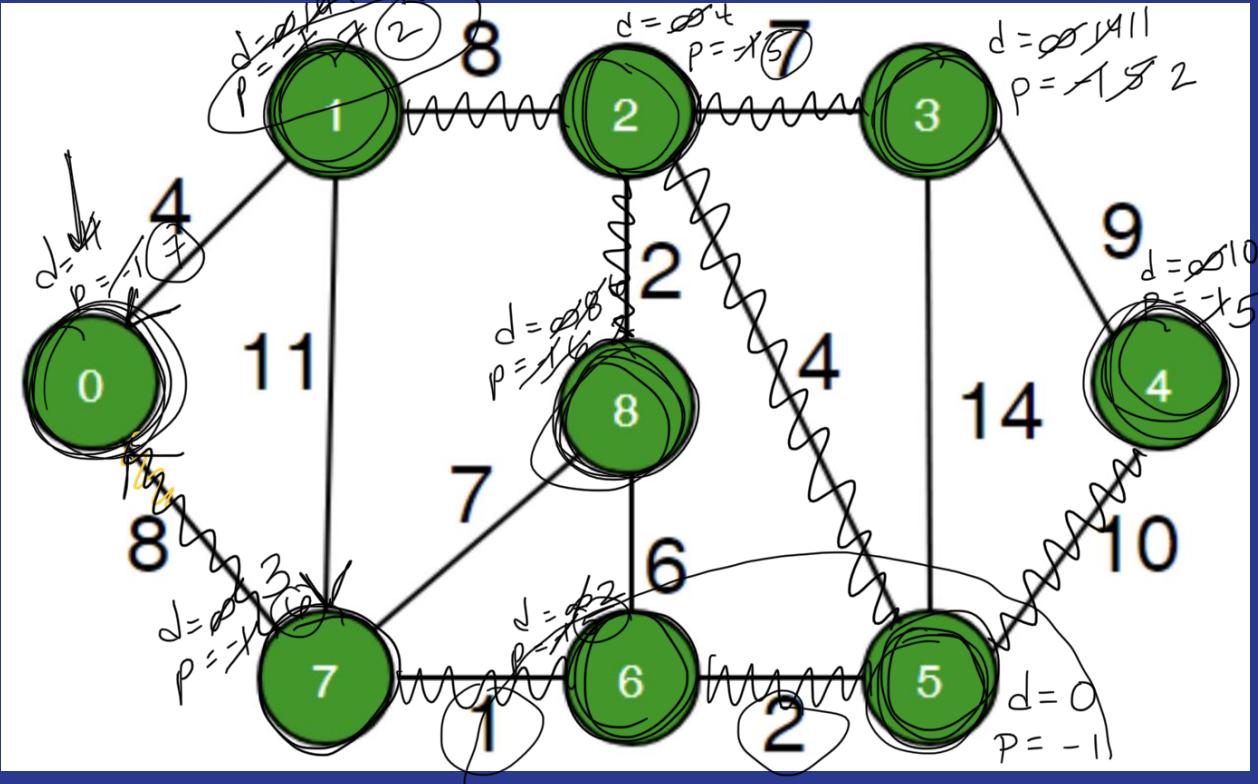
update

The Main Loop:

Until the priority queue is empty we...

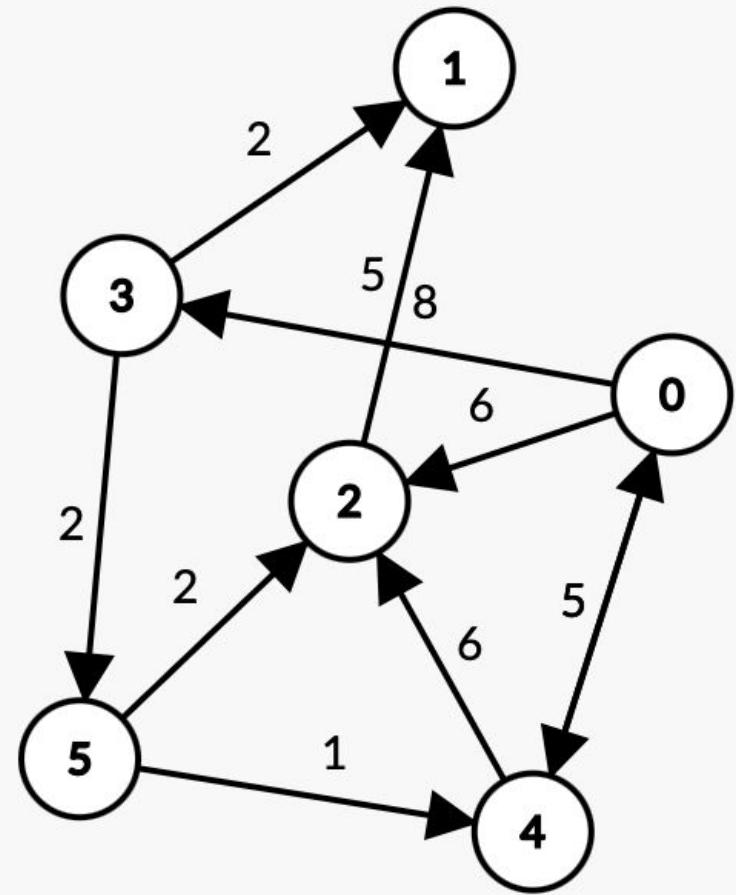
- get the minimum vertex
- check its adjacent vertices to find the “shortest” path to the new vertex
- update **dist**, **prev**, and the queue as necessary

Example. Show Dijkstra's Algorithm on the graph below.



P	Q
$(\infty, 5)$	$(11, 0)$
(∞, ∞)	$(11, 1)$
$(\infty, 1)$	$(14, 1)$
$(\infty, 2)$	$(4, 2)$
$(\infty, 3)$	$(14, 3)$
$(\infty, 4)$	$(10, 4)$
$(\infty, 6)$	$(2, 6)$
$(\infty, 7)$	$(3, 7)$
$(\infty, 8)$	$(8, 8)$
$(\infty, 9)$	$(6, 9)$

Example. Show Dijkstra's Algorithm on the graph below.



Analysis: What is the runtime?

Algorithm *dijkstrasShortestPath*(G, s) **Input:** a weighted graph $G = (V, E)$ with non-negative edge weights and a source vertex s
Output: the paths with minimum total weight from s to all other vertices in G

$dist :=$ an array of size $|V|$

$prev :=$ an array of size $|V|$

$Q :=$ a min-heap-based priority queue

$dist[s] \leftarrow 0$

for each vertex v in G :

if $v \neq s$

$dist[v] \leftarrow \infty$

$prev[v] \leftarrow -1$

$Q.add(dist[v], v)$

while $!Q.isEmpty()$:

$u \leftarrow Q.extractMin()$

for each vertex v in $G.adj(u)$ that is still in Q

$alt \leftarrow dist[u] + length(u, v)$

if $alt < dist[v]$

$dist[v] \leftarrow alt$

$prev[v] \leftarrow u$

$Q.set(alt, v)$

return $dist, prev$

- Runtime for setting all the values in the two arrays: $\underline{\underline{O(V)}}$
- Runtime for adding all the items to the priority queue: $O(\sqrt{V} \log V)$

Algorithm *dijkstrasShortestPath*(G, s) **Input:** a weighted graph $G = (V, E)$ with non-negative edge weights and a source vertex s
Output: the paths with minimum total weight from s to all other vertices in G

$dist :=$ an array of size $|V|$

$prev :=$ an array of size $|V|$

$Q :=$ a min-heap-based priority queue

$dist[s] \leftarrow 0$

for each vertex v in G :

if $v \neq s$

$dist[v] \leftarrow \infty$

$prev[v] \leftarrow -1$

$Q.add(dist[v], v)$

while $!Q.isEmpty()$:

$u \leftarrow Q.extractMin()$

for each vertex v in $G.adj(u)$ that is still in Q

$alt \leftarrow dist[u] + length(u, v)$

if $alt < dist[v]$

$dist[v] \leftarrow alt$

$prev[v] \leftarrow u$

$Q.set(alt, v)$

return $dist, prev$

- Runtime for setting all the values in the two arrays: **O(V)**
- Runtime for adding all the items to the priority queue: **O(VlogV)**

Algorithm *dijkstrasShortestPath*(G, s) **Input:** a weighted graph $G = (V, E)$ with non-negative edge weights and a source vertex s
Output: the paths with minimum total weight from s to all other vertices in G

$dist :=$ an array of size $|V|$

$prev :=$ an array of size $|V|$

$Q :=$ a min-heap-based priority queue

$dist[s] \leftarrow 0$

for each vertex v in G :

if $v \neq s$

$dist[v] \leftarrow \infty$

$prev[v] \leftarrow -1$

$Q.add(dist[v], v)$

while $\neg Q.isEmpty()$:

$u \leftarrow Q.extractMin()$ $O(\log V)$

for each vertex v in $G.adj(u)$ that is still in Q

$alt \leftarrow dist[u] + length(u, v)$

if $alt < dist[v]$

$dist[v] \leftarrow alt$

$prev[v] \leftarrow u$

$Q.set(alt, v)$

return $dist, prev$

- Runtime for setting all the values in the two arrays: $O(V)$
- Runtime for adding all the items to the priority queue: $O(V \log V)$

- Runtime for resetting values in the two arrays: $O(E)$
- Runtime for resetting values in the queue: $O(E \log V)$

Analysis: What is the runtime?

- In all, the runtime is $O(V) + O(V\log V) + O(E) + O(E\log V) =$
 $O((V+E)\log V)$
- Since the graph is connected, we can also say $O(E\log V)$ since every vertex is connected to at least one edge...

Further explanations:

- At most, the values $\text{dist}[v]$ and $\text{prev}[v]$ for any vertex v are updated $\deg(v)$ times, and the sum of all $\deg(v)$ (for all vertices) is $2E$, which explains the $O(E)$ runtime for the updates in the arrays.
- The same idea explains the $O(E\log V)$ runtime for the queue updates

Analysis: Note on Priority Queue Implementation

- The $O((V+E)\log V)$ runtime depends on the implementation of the algorithm regarding the vertices in the priority queue.
- Recall that a binary heap allows **insert** and **delMin** (for a min-heap) in $O(\log V)$ time, but **search** in a binary heap is $O(V)$ in the worst case.
- This analysis depends upon the method for updating the a key in the priority queue in $O(\log V)$ time.
- In order to do that, the implementation of the algorithm must maintain a link between vertices and their positions in the queue (e.g. an array of size V that gets updated whenever a vertex changes position in the queue).
- This would also allow $O(1)$ checking whether a vertex is still in the queue
- When a key gets updated, at most $\log V$ vertex positions will have to be updated as the heap is rearranged--so updating a key can be done in $O(\log V)$ time.

Proof of Correctness for Dijkstra's Algorithm

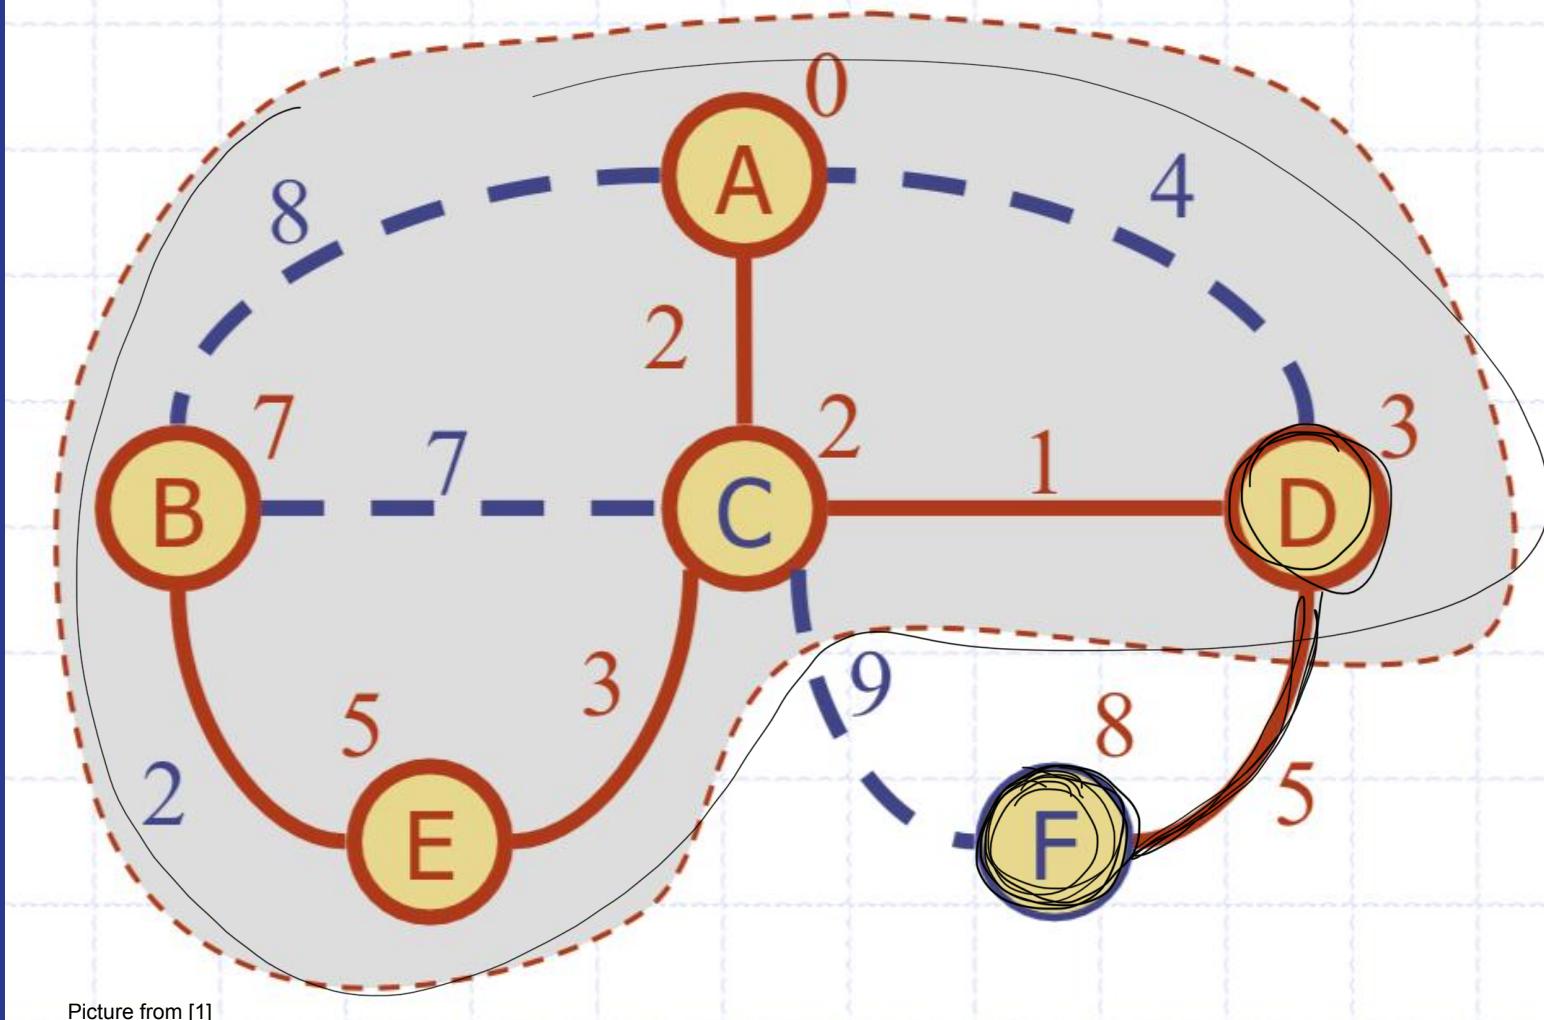
Conjecture: Dijkstra's Algorithm works on all valid inputs

Assume there is some input that results in incorrect results.

Why does Dijkstra's Algorithm work? Proof by Contradiction

Dijkstra's Algorithm is a **greedy** approach to the shortest path problem--it adds nodes by choosing the shortest distances first using the vertices that have been visited.

- Suppose it didn't find all the shortest distances. Suppose the distance it chose from **s** to **F** was NOT the shortest using the previously visited vertices, and **F** is the first vertex chosen incorrectly.
- Let **D** be the vertex previous to **F** on the chosen path. Since **F** was chosen incorrectly, that means that there must be a shorter path to **F** from **s** either using **D** or not using **D**.
 - If that shorter path uses **D**, then the path from **s** to **D** was not the shortest, which is a contradiction because we assumed **F** was the first to be chosen incorrectly.
 - If that shorter path does not use **D**, then the edge from **D** to **F** was not *relaxed*, which is a contradiction because Dijkstra's Algorithm only adds edges once they are relaxed.

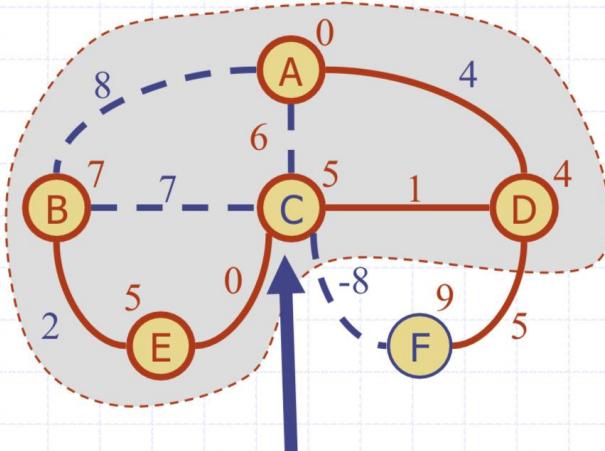


Why It Doesn't Work for Negative-Weight Edges



- Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.

- If a node with a negative incident edge were to be added late to the cloud, it could mess up distances for vertices already in the cloud.



C's true distance is 1, but it is already in the cloud with $d(C)=5$!

DAG-based Shortest Path

Algorithm $DAGDistances(G, s)$

Input: $G = (V, E)$, a weighted, directed, acyclic graph and s , the source vertex

Output: the paths with minimum total weight from s to all other vertices in G

$dist[] :=$ an array of size $|V|$

$prev[] :=$ an array of size $|V|$

for each vertex $v \in V$:

$dist[v] = \infty$

$prev[v] = -1$

$dist[s] = 0$

perform a topological sort of the vertices,
assigning values 1 to $n = |V|$

for $u \leftarrow 1$ to n **do:** //in topological order

for each edge $e = (u, v)$ with weight w going out
 of u **do:**

 //relax edge e

if $dist[u] + w < dist[v]$:

$dist[v] := dist[u] + w$

$prev[v] := u$

return $prev[], dist[]$

DAG-based Shortest Path

DAG-based Shortest Path

Algorithm $DAGDistances(G, s)$

Input: $G = (V, E)$, a weighted, directed, acyclic graph and s , the source vertex

Output: the paths with minimum total weight from s to all other vertices in G

$dist[] :=$ an array of size $|V|$

$prev[] :=$ an array of size $|V|$

for each vertex $v \in V$:

$dist[v] = \infty$

$prev[v] = -1$

$dist[s] = 0$

perform a topological sort of the vertices,
assigning values 1 to $n = |V|$

for $u \leftarrow 1$ to n **do:** //in topological order

for each edge $e = (u, v)$ with weight w going out
 of u **do:**

 //relax edge e

if $dist[u] + w < dist[v]$:

$dist[v] := dist[u] + w$

$prev[v] := u$

return $prev[], dist[]$

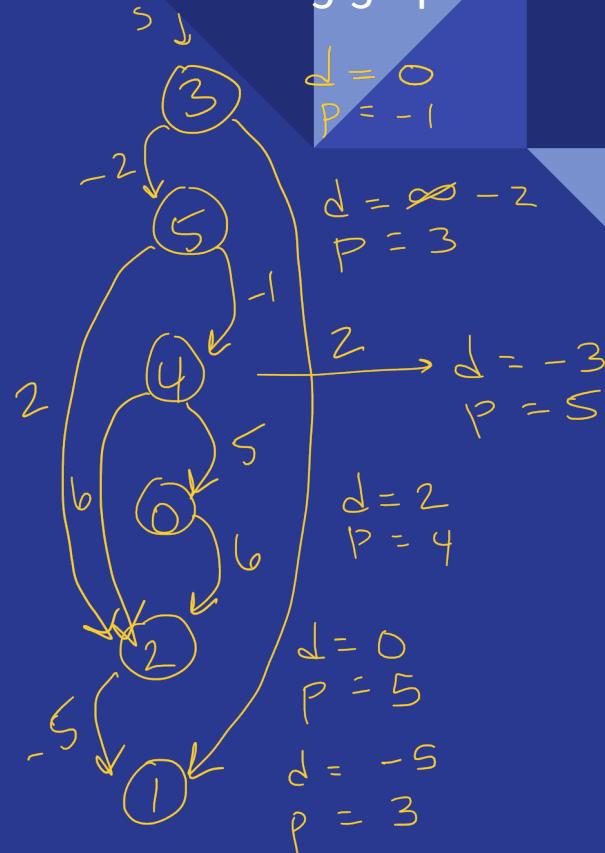
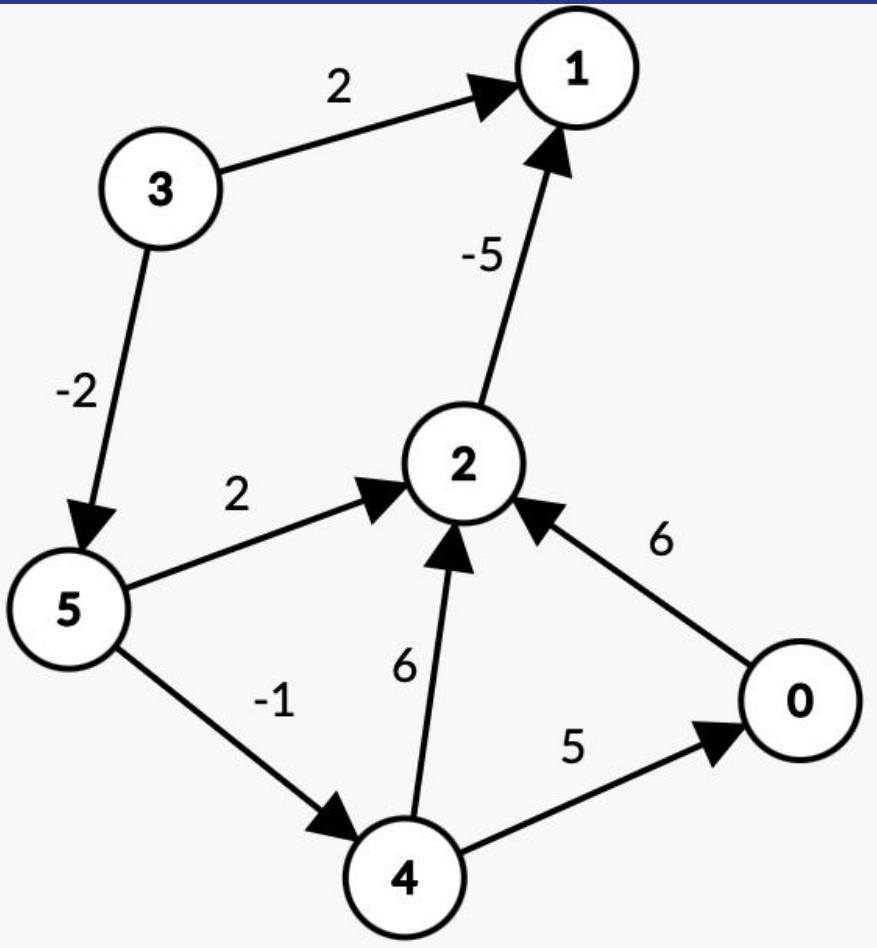
Setting the values: $O(|V|)$

Topological sort with DFS:
 $O(|V| + |E|)$

Relaxation of edges:
 $O(|V| + |E|)$

Total: $O(|V| + |E|)$

Example. Run the DAG-based Shortest Path Algorithm on the following graph.



References

1. Goodrich & Tamassia
2. Sedgewick & Wayne