

Minimum Spanning Tree

- Terminology and Properties
- Prim's Algorithm
- Kruskal's Algorithm


Imagine...

You are employed at a brand new airline, and your first task is to determine which flights the airline should provide according to the following guidelines:

- There needs to be a way to get between each pair of cities/airports in the following list: San Francisco, CA (SFO); Vancouver, BC (YVR); Albuquerque, NM (ABQ); Chicago, IL (ORD); Atlanta, GA (ATL); New York, NY (LGA)
- The route between two cities does not have to be a direct flight.
- The total number of flights should be minimized.
- The total distance covered by all flights should be minimized.
- Assume that if a flight exists, it goes both ways (e.g. SFO→YVR and YVR→SFO)

How would you go about solving this problem?

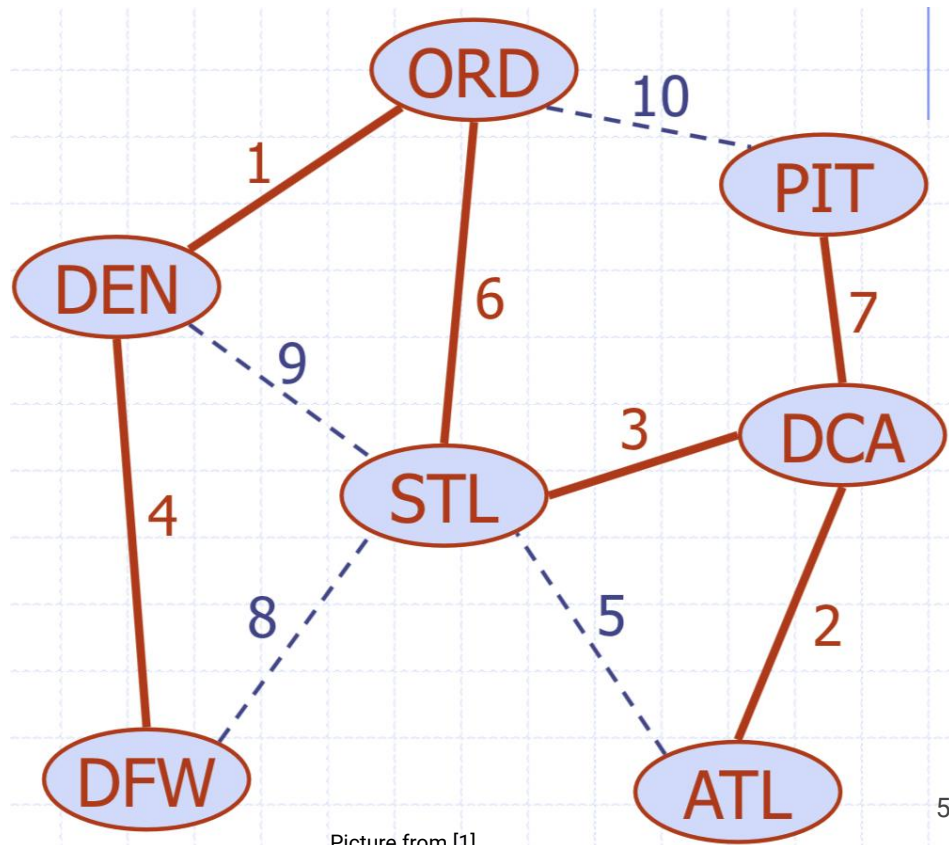
	SFO	YVR	ABQ	ORD	ATL	LGA
SFO	0	1280	443	2960	3440	4139
YVR	1280	0	2072	2820	3600	3920
ABQ	443	2072	0	1799	2048	2940
ORD	2960	2820	1799	0	951	1163
ATL	3440	3600	2048	951	0	1219
LGA	4139	3920	2940	1163	1219	0



What we really want is a minimum spanning tree (MST)...

A minimum spanning tree of a graph G...


- ...is a spanning subgraph (i.e. it contains all the vertices of G)
- ...is a tree (i.e. it has no cycles)
- ...is minimal weight-wise (i.e. the total weight of all the edges it uses is minimized so that no other spanning tree has a smaller total weight)
- ...does NOT guarantee the shortest path between two vertices!



Property 1: The Cycle Property

- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T
- Let C be the cycle formed by adding e to T
- Claim: For every edge f of C :
 $\text{weight}(f) \leq \text{weight}(e)$


Proof by Contradiction: ???



Property 1: The Cycle Property

- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T
- Let C be the cycle formed by adding e to T
- Claim: For every edge f of C :
 $\text{weight}(f) \leq \text{weight}(e)$


Proof by Contradiction: Given all the properties above, assume that for some edge f in C , $\text{weight}(f) > \text{weight}(e)$. Then replacing f with e will produce a spanning tree T' such that the total weight of T' is smaller than the total weight of T . But that contradicts the definition of T as an MST of G .



Property 2: The Partition Property

- Consider a partition of a graph \mathbf{G} into subsets \mathbf{U} and \mathbf{V}
- Let \mathbf{e} be an edge of minimum weight across the partition (i.e. \mathbf{e} has one endpoint in \mathbf{U} and one endpoint in \mathbf{V})
- Claim: There is a minimum spanning tree of \mathbf{G} that includes edge \mathbf{e}


Proof: ???



Property 2: The Partition Property

- Consider a partition of a graph G into subsets U and V
- Let e be an edge of minimum weight across the partition (i.e. e has one endpoint in U and one endpoint in V)
- Claim: There is a minimum spanning tree of G that includes edge e


Proof: Let T be an MST of G and let the definitions above be true. If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition. By the cycle property, $\text{weight}(f) \leq \text{weight}(e)$. Thus, $\text{weight}(f) = \text{weight}(e)$, and we obtain another MST by replacing f with e .



Questions to think about...

1. Given an undirected, weighted graph \mathbf{G} , let \mathbf{T} be an MST of \mathbf{G} . Is \mathbf{T} unique?
2. Does an MST also give you the shortest path between a pair of vertices?
3. Does every weighted undirected graph have an MST?
4. Can you find an MST of a weighted undirected graph if there are negative weight edges?
5. Can you find an MST in a directed graph?
6. How would you find an MST of a weighted undirected graph?

graph




MST



A 250-vertex Euclidean graph (with 1,273 edges) and its MST


Prim's Algorithm (aka Jarnik's Algorithm)

- Grow the tree one edge at a time
- Start with a single vertex (counts as a tree)
- Add $|V| - 1$ edges to it by adding the minimum-weight edge that connects a new vertex to the tree (crossing the partition that is defined by the current tree vertices--called a *crossing edge*)
- NOTE: When you add an edge to the tree, you are also adding a vertex to the tree.
- Like Dijkstra's Shortest Path, this is a greedy algorithm.



Prim's Algorithm (aka Jarnik's Algorithm)

Proof of Correctness: Follows directly from the Partition Property because we are choosing the minimum weight edge across the tree-defined partition.



Implementation of Prim's Algorithm

- **marked[]**: an array of booleans to keep track of vertices on the tree
- **edgeTo[]**: an array to keep track of the lightest edge connecting a new vertex to the tree
- **distTo[]**: an array to keep track of the weight of the lightest edge connecting a new vertex to the tree
- **pq**: a minimum priority queue to keep track of eligible crossing edges (key is the weight of the edges)

Algorithm *PrimsMST*(G)

Input: $G = (V, E)$, a weighted, undirected graph

Output: A minimum-weight spanning tree of G

$edgeTo[]$:= a $|V|$ -sized array to store the edge connecting a vertex to the tree

$distTo[]$:= a $|V|$ -sized array to keep track of the distance of the edge connecting a vertex to a tree

$marked[]$:= a $|V|$ -sized array to keep track of which vertices have been visited

pq := a heap-based min priority queue with weights as keys and vertices as values

//initialize structures

for all $v \in V$:

$distTo[v] = \infty$

$distTo[0] = 0$

$pq.insert(0, 0)$

//main loop

while ! $pq.isEmpty$:

$v = pq.delMin()$

$marked[v] = T$

for each edge $e = (v, u)$ adjacent to v :

if $marked[u]$

continue

if $e.weight() < distTo[u]$:

$edgeTo[u] = e$


$distTo[u] = e.weight()$

if $pq.contains(u)$:

$pq.changeKey(u, distTo[u])$

else:

$pq.insert(u, distTo[u])$



```
//main loop
```

```
while !pq.isEmpty:
```

```
    v = pq.delMin()
```

```
    marked[v] = T
```

```
    for each edge e = (v, u) adjacent to v:
```

```
        if marked[u]
```

```
            continue
```

```
        if e.weight() < distTo[u]:
```

```
            edgeTo[u] = e
```

```
            distTo[u] = e.weight()
```

```
        if pq.contains(u):
```

```
            pq.changeKey(u, distTo[u])
```

```
        else:
```

```
            pq.insert(u, distTo[u])
```


Algorithm *PrimsMST*(G)

Input: $G = (V, E)$, a weighted, undirected graph

Output: A minimum-weight spanning tree of G
 $edgeTo[]$:= a $|V|$ -sized array to store the edge connecting a vertex to the tree
 $distTo[]$:= a $|V|$ -sized array to keep track of the distance of the edge connecting a vertex to a tree
 $marked[]$:= a $|V|$ -sized array to keep track of which vertices have been visited
 pq := a heap-based min priority queue with weights as keys and vertices as values

//initialize structures

for all $v \in V$:

$distTo[v] = \infty$

$distTo[0] = 0$

$pq.insert(0, 0)$

$O(|V|)$

//main loop

while ! $pq.isEmpty$:

$v = pq.delMin()$

$marked[v] = T$

for each edge $e = (v, u)$ adjacent to v :

if $marked[u]$

continue

if $e.weight() < distTo[u]$:

$edgeTo[u] = e$

$distTo[u] = e.weight()$

if $pq.contains(u)$:

$pq.changeKey(u, distTo[u])$

else:

$pq.insert(u, distTo[u])$

```
//main loop
```

```
while !pq.isEmpty:
```

```
    v = pq.delMin()
```

```
    marked[v] = T
```

```
    for each edge e = (v, u) adjacent to v:
```

```
        if marked[u]
```

```
            continue
```

```
        if e.weight() < distTo[u]:
```

```
            edgeTo[u] = e
```

```
            distTo[u] = e.weight()
```

```
            if pq.contains(u):
```


```
                pq.changeKey(u, distTo[u])
```

```
            else:
```

```
                pq.insert(u, distTo[u])
```


Assuming adjacency list representation, total time is **$O((V + E)\log V) = O(E\log V)$** for a connected graph

- delMin in a heap-based PQ: $O(\log V)$
- checking if PQ contains a vertex and changing a key is $O(\log V)$ as long as there is an auxiliary data structure keeping track of positions in the queue
- key of any vertex v is updated at most $\deg(v)$ times and sum of all degrees is $O(E)$




Kruskal's Algorithm

- Process edges, adding edges to the tree smallest-weight first **as long as the new edge does not form a cycle.**
- The result is that the algorithm creates a forest that eventually merges into a single tree.
- What would be some of the challenges in implementing this?
- What data structures would be useful?
- How would you represent the graph?



Kruskal's Algorithm

- Can be implemented using disjoint sets for each separate component
- Start with $|V|$ disjoint sets, each containing a single vertex
- Combine sets (union) by adding edges, reducing the number of separate components until there is only one



Algorithm *KruskalsMST*(G)

Input: $G = (V, E)$, a connected, weighted, undirected graph

Output: A minimum-weight spanning tree of G

$pq :=$ a minimum priority queue of edges where keys are weights

for each edge $e = (u, v) \in E$:

$pq.insert(e, e.weight())$

$A = \emptyset$

for each $v \in V$:

$makeSet(v)$

while $|A| < |V| - 1$

$e = (u, v) = pq.delMin()$

if $findSet(u) \neq findSet(v)$:

$A = A \cup (u, v)$

$union(u, v)$

return A

Algorithm *KruskalsMST*(G)

Input: $G = (V, E)$, a connected, weighted, undirected graph

Output: A minimum-weight spanning tree of G

$pq :=$ a minimum priority queue of edges where keys are weights

for each edge $e = (u, v) \in E$:

$pq.insert(e, e.weight())$

$A = \emptyset$

for each $v \in V$:

$makeSet(v)$

while $|A| < |V| - 1$

$e = (u, v) = pq.delMin()$

if $findSet(u) \neq findSet(v)$:

$A = A \cup (u, v)$

$union(u, v)$

return A

- **$O(E \log E)$** for doing **E** inserts into the **E** -sized PQ
- Alternatively, we could just sort the edges in **$O(E \log E)$** time

Algorithm *KruskalsMST*(G)

Input: $G = (V, E)$, a connected, weighted, undirected graph

Output: A minimum-weight spanning tree of G

$pq :=$ a minimum priority queue of edges where keys are weights

for each edge $e = (u, v) \in E$:

$pq.insert(e, e.weight())$

$A = \emptyset$

for each $v \in V$:

$makeSet(v)$

while $|A| < |V| - 1$

$e = (u, v) = pq.delMin()$

if $findSet(u) \neq findSet(v)$:

$A = A \cup (u, v)$



$union(u, v)$


return A

- $makeSet$ is an $O(1)$ operation, so the total time is **$O(V)$**
- The total time required for **E union** and $findSet$ operations is **$O(E \log V)$**
- **Total time: $O(E \log E + E \log V) = O(E \log E)$**


Prim's vs. Kruskal's


- Space: $|V|$ vs. $|E|$
- Time: $|E|\log|V|$ vs. $|E|\log|E|$








More Examples











References

- [1] Goodrich and Tamassia
- [2] Sedgewick and Wayne
- [3] en.wikipedia.org