

Searching III

Binary Search Trees

Ordered Dictionary:

Keeps the keys in order

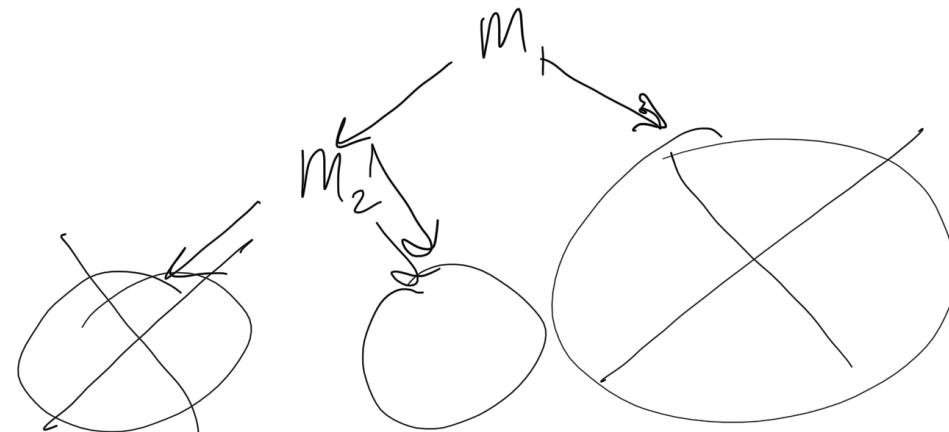
public class ST<Key extends Comparable<Key>, Value>	
ST()	<i>create an ordered symbol table</i>
void put(Key key, Value val)	<i>put key-value pair into the table (remove key from table if value is null)</i>
Value get(Key key)	<i>value paired with key (null if key is absent)</i>
void delete(Key key)	<i>remove key (and its value) from table</i>
boolean contains(Key key)	<i>is there a value paired with key?</i>
boolean isEmpty()	<i>is the table empty?</i>
int size()	<i>number of key-value pairs</i>
Key min()	<i>smallest key</i>
Key max()	<i>largest key</i>
Key floor(Key key)	<i>largest key less than or equal to key</i>
Key ceiling(Key key)	<i>smallest key greater than or equal to key</i>
int rank(Key key)	<i>number of keys less than key</i>
Key select(int k)	<i>key of rank k</i>
void deleteMin()	<i>delete smallest key</i>
void deleteMax()	<i>delete largest key</i>
int size(Key lo, Key hi)	<i>number of keys in [lo..hi]</i>
Iterable<Key> keys(Key lo, Key hi)	<i>keys in [lo..hi], in sorted order</i>
Iterable<Key> keys()	<i>all keys in the table, in sorted order</i>

Intuition for Binary Search Tree

Combine the useful elements of linked lists with the useful elements of ordered arrays.

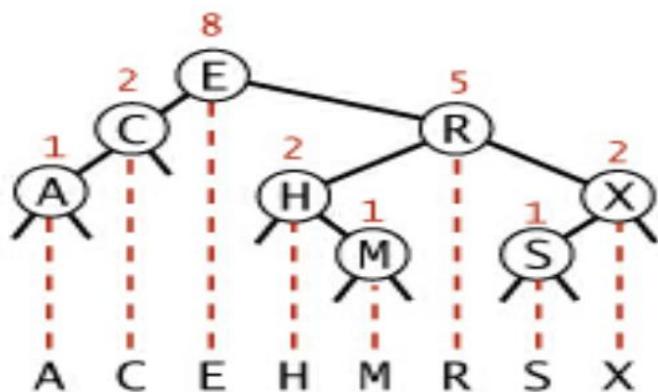
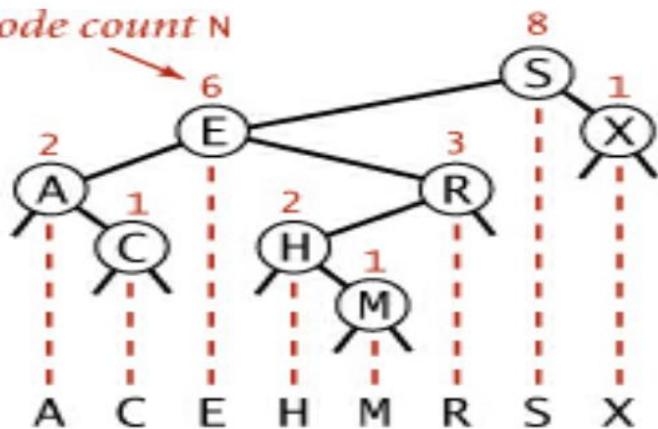
linked list : easy to grow (insert, delete)

ordered array : indexing - binary search



Binary Search Tree

node count N

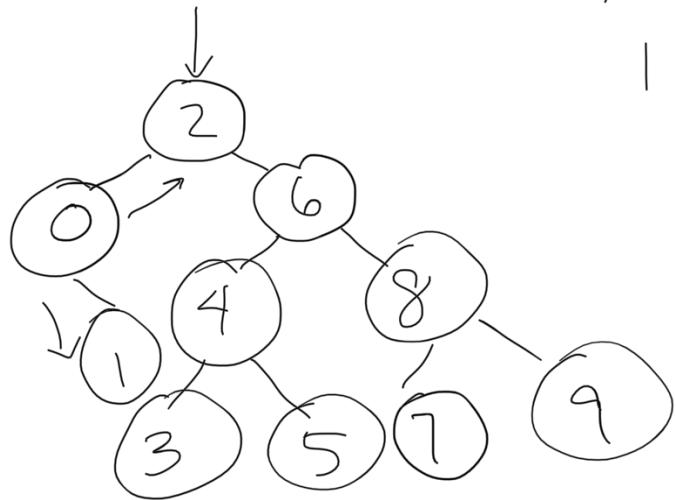


**Two BSTs that represent
the same set of keys**

Same keys, different trees...

HOW DOES THIS HAPPEN?

random order : 2, 6, 4, 0, 8, 5,
1, 7, 3, 1, 9



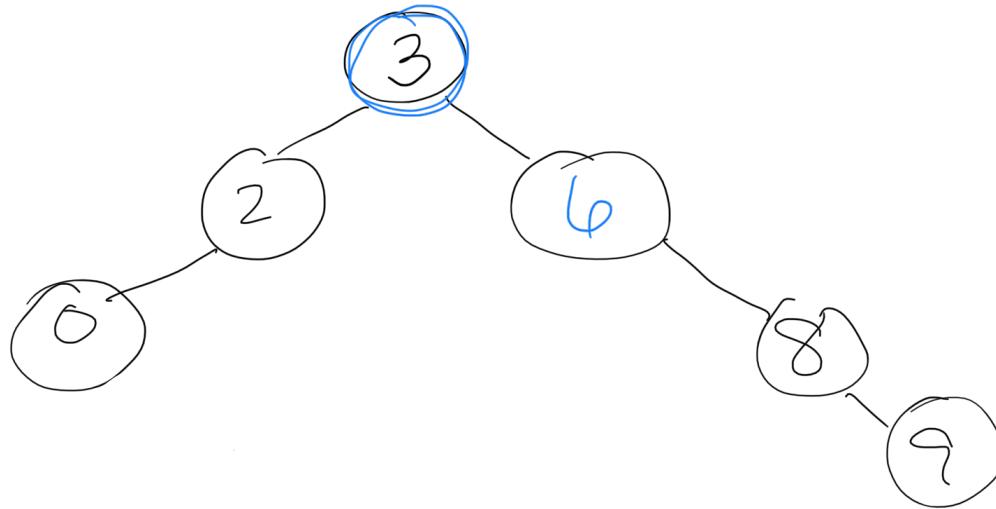
best case :

delete 1

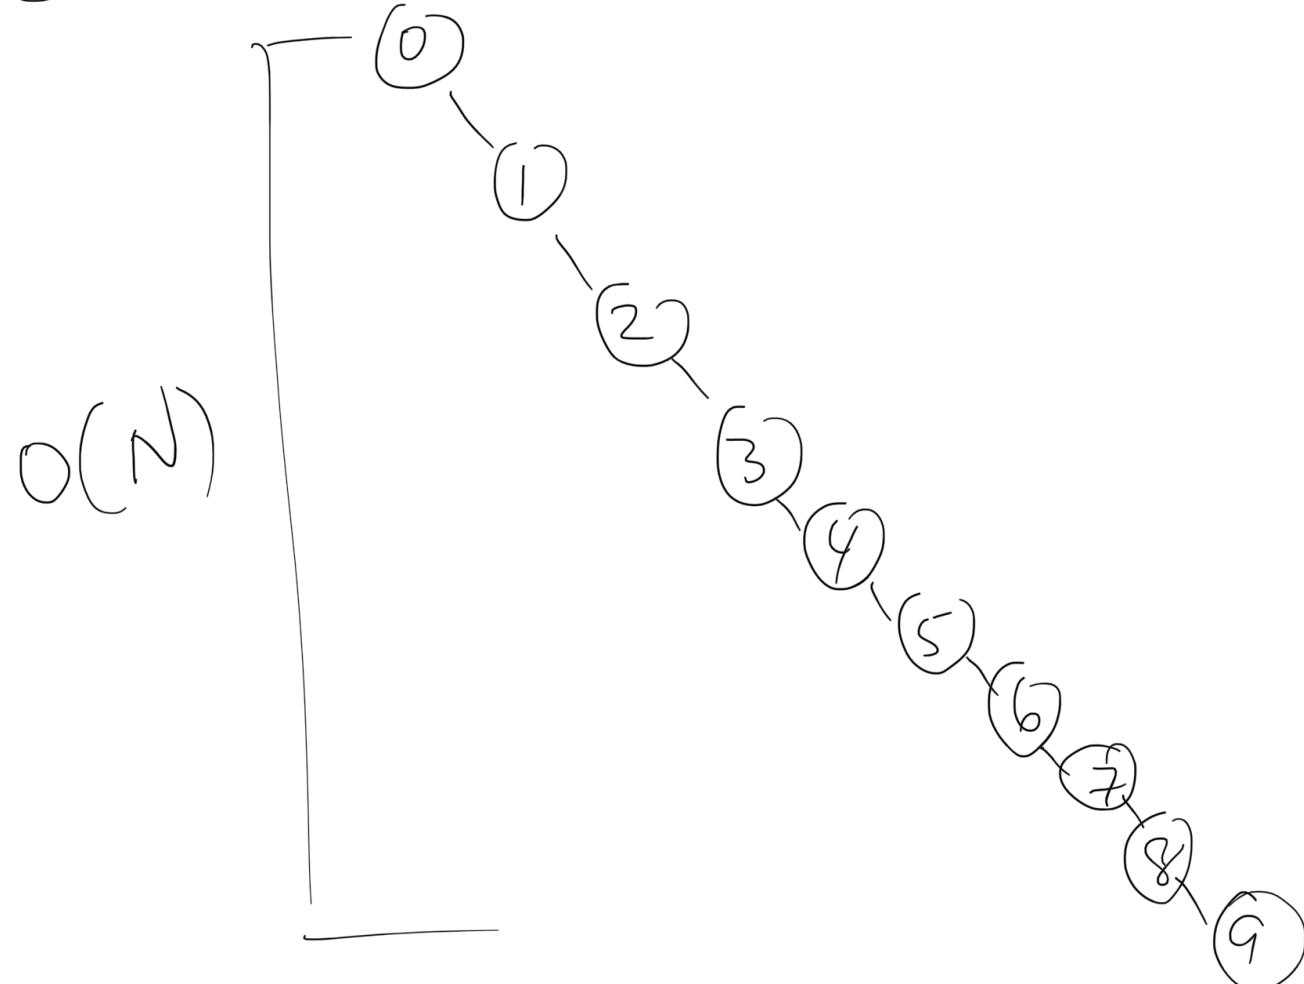
delete 4

delete 5

delete 7



worst case :



Main Operations

put -

get

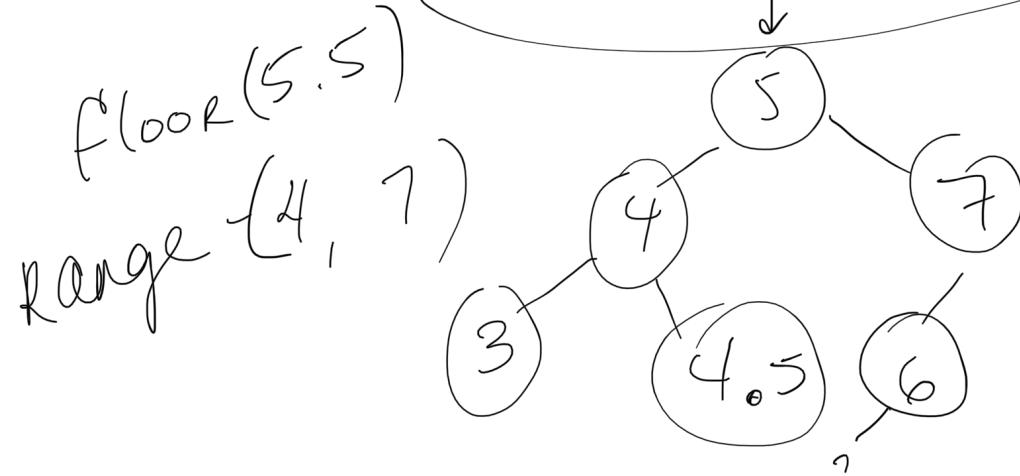
delete

Other Operations

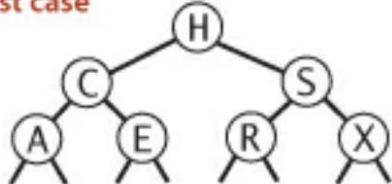
min

max

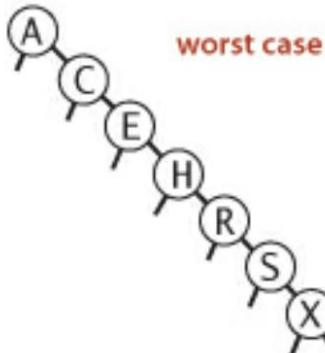
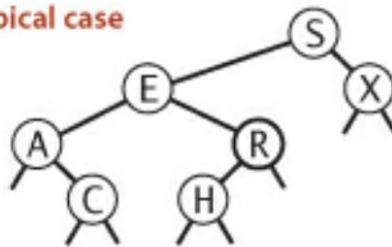
floor (key) - largest key \leq key



best case



typical case



BST possibilities

algorithm
(data structure)

worst-case cost
(after N inserts)

average-case cost
(after N random inserts)

efficiently
support ordered
operations?

*sequential search
(unordered linked list)*

search

insert

search hit

insert

no

*binary search
(ordered array)*

$\lg N$

N

$\lg N$

$N/2$

yes

*binary tree search
(BST)*

N

N

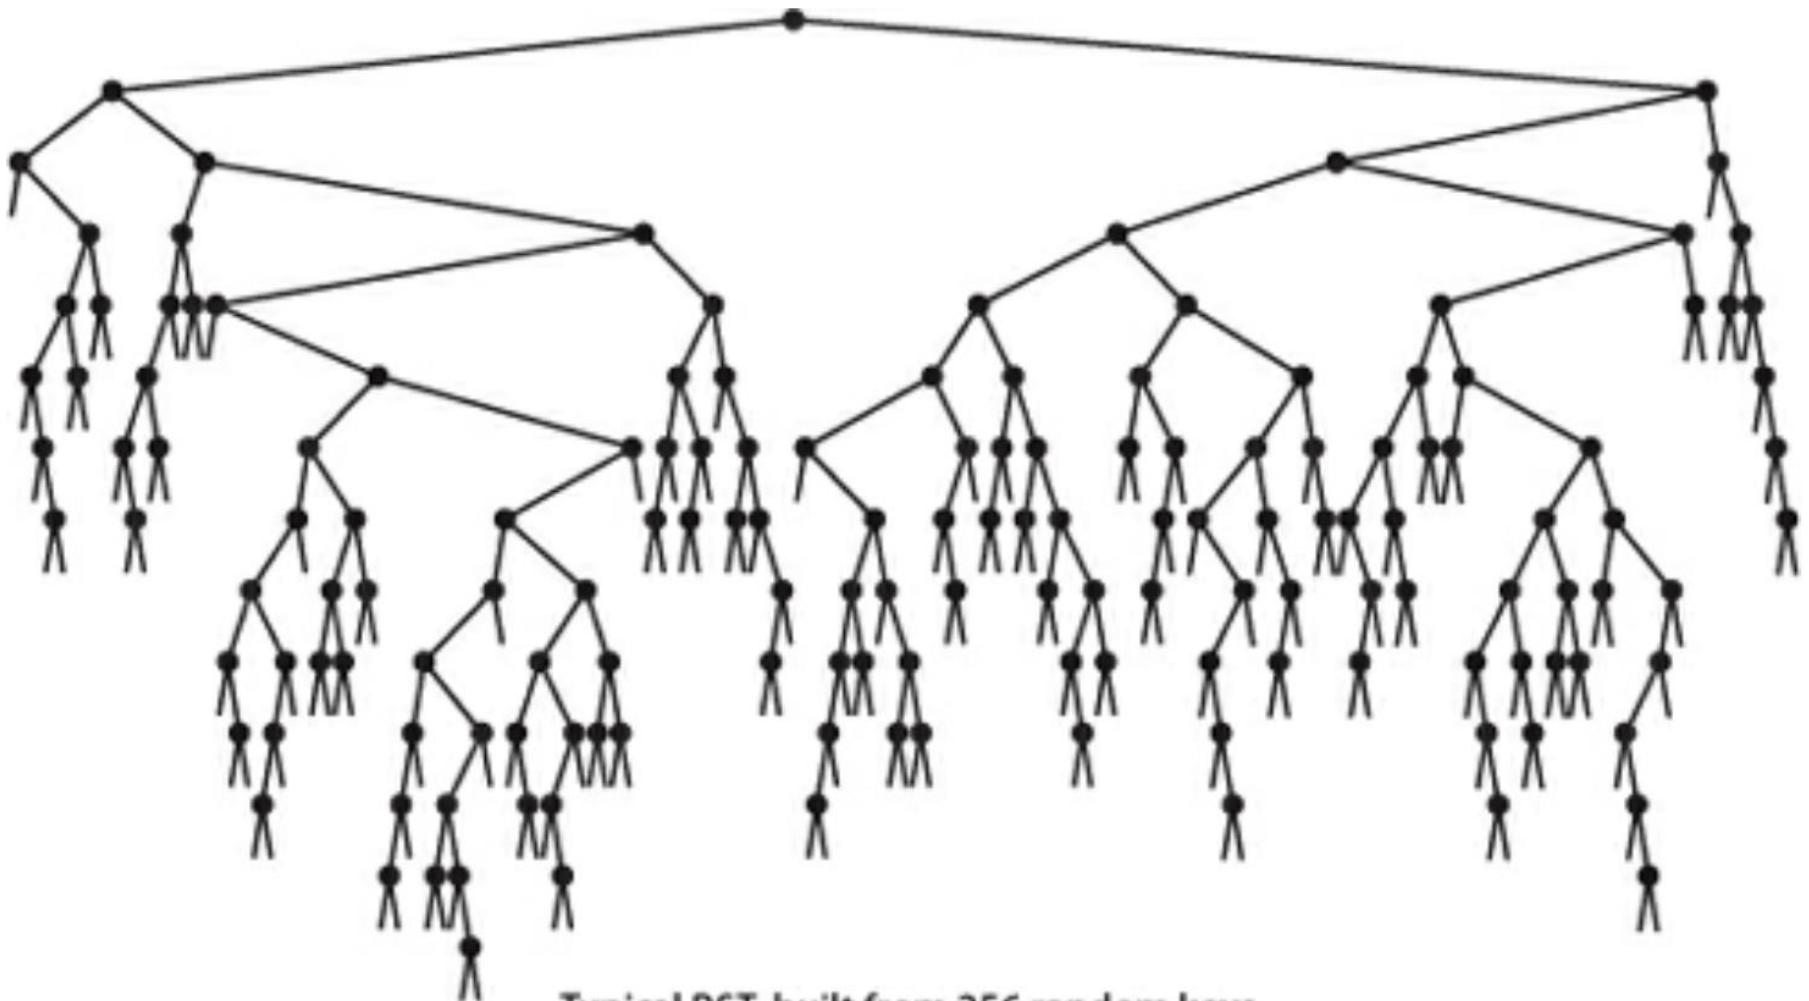
$1.39 \lg N$

$1.39 \lg N$

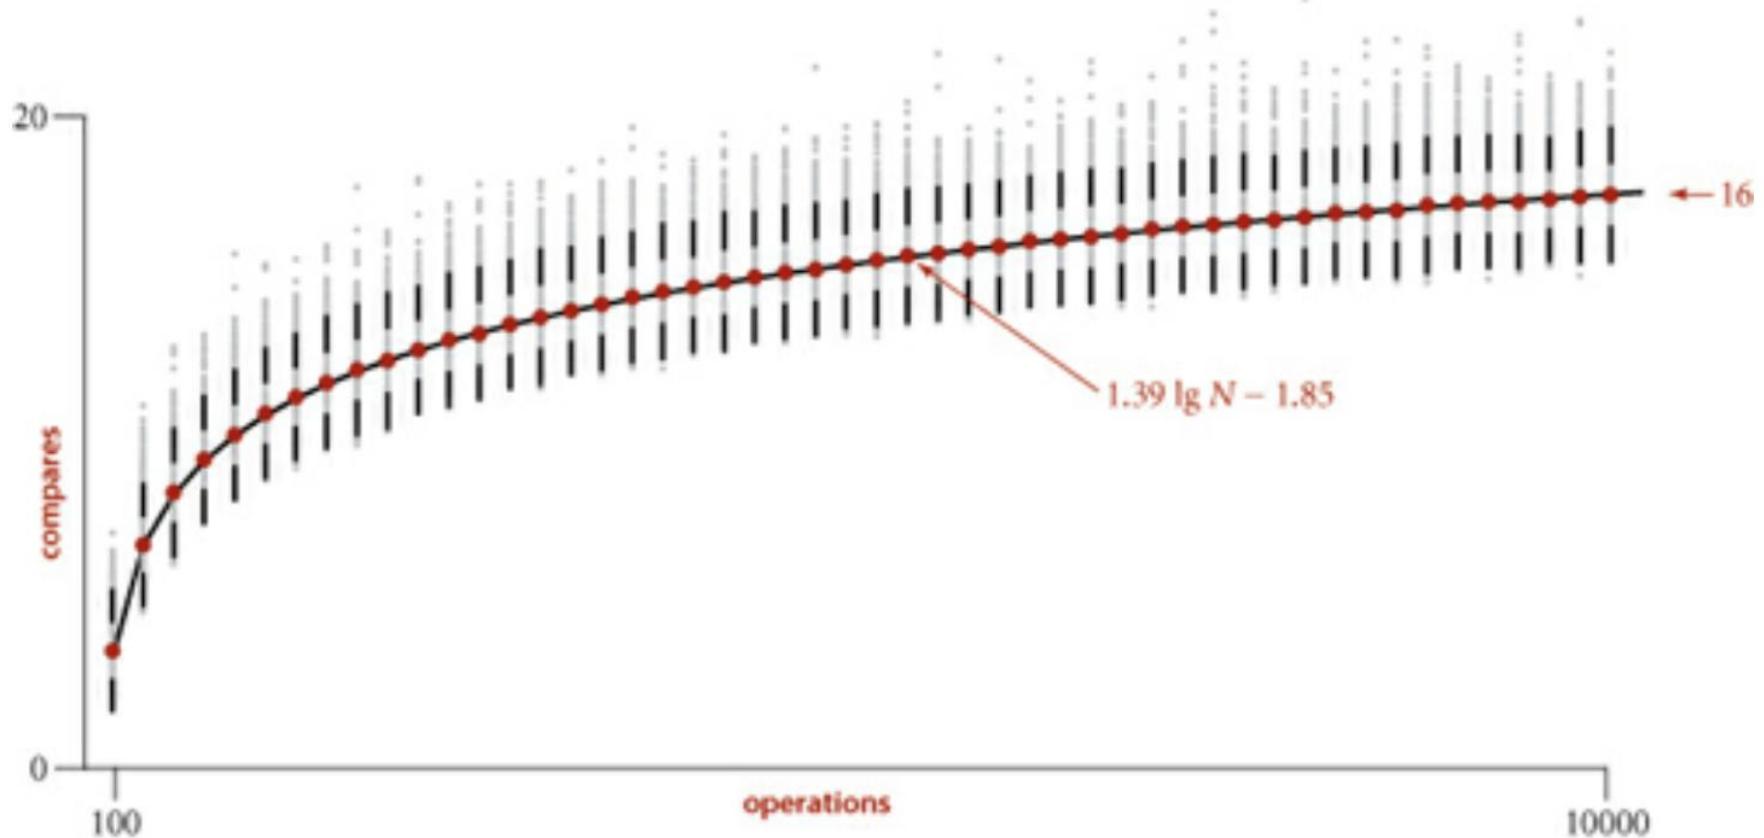
yes

Cost summary for basic symbol-table implementations (updated)

The good news is that the average case (with random data) is closer to the best case than the worst case!



Typical BST, built from 256 random keys





How can we improve upon this?

References

- [1] *Algorithms, Fourth Edition*; Robert Sedgewick and Kevin Wayne (and associated slides)
- [2] Slides from <https://www.cs.princeton.edu/~rs/talks/LLRB/RedBlack.pdf>