

# Processes

---

## **What is a process?**

- “An execution stream in the context of a particular process state.”
- State variables
  - Instruction Pointer
  - Registers
  - Stack Variables
  - Memory
  - etc.

# Processes

---

- A process is similar to a program...
- A **program** is a set of instructions to perform, in a certain order
- A **process** is one thing that is executing that program
- But there can be **multiple processes** running the **same program**

# Processes

---

- Internally, processes **pretend** that they are running on a uni-programming machine
  - Only one CPU, and we run on it 100% of the time
  - Absolute control over memory
  - Never blocked
- Syscalls break this illusion, but we mostly treat them **like function calls**

# Processes

---

- Processes routinely link to **shared libraries**
  - Classic code, often provided as part of the OS. Or can be custom libraries, just for you
- Shared libraries are **part of the process**
  - Even though you didn't write them, they are integrated into your program
  - No special rights or powers (or limits)
  - OS treats library code the same as main

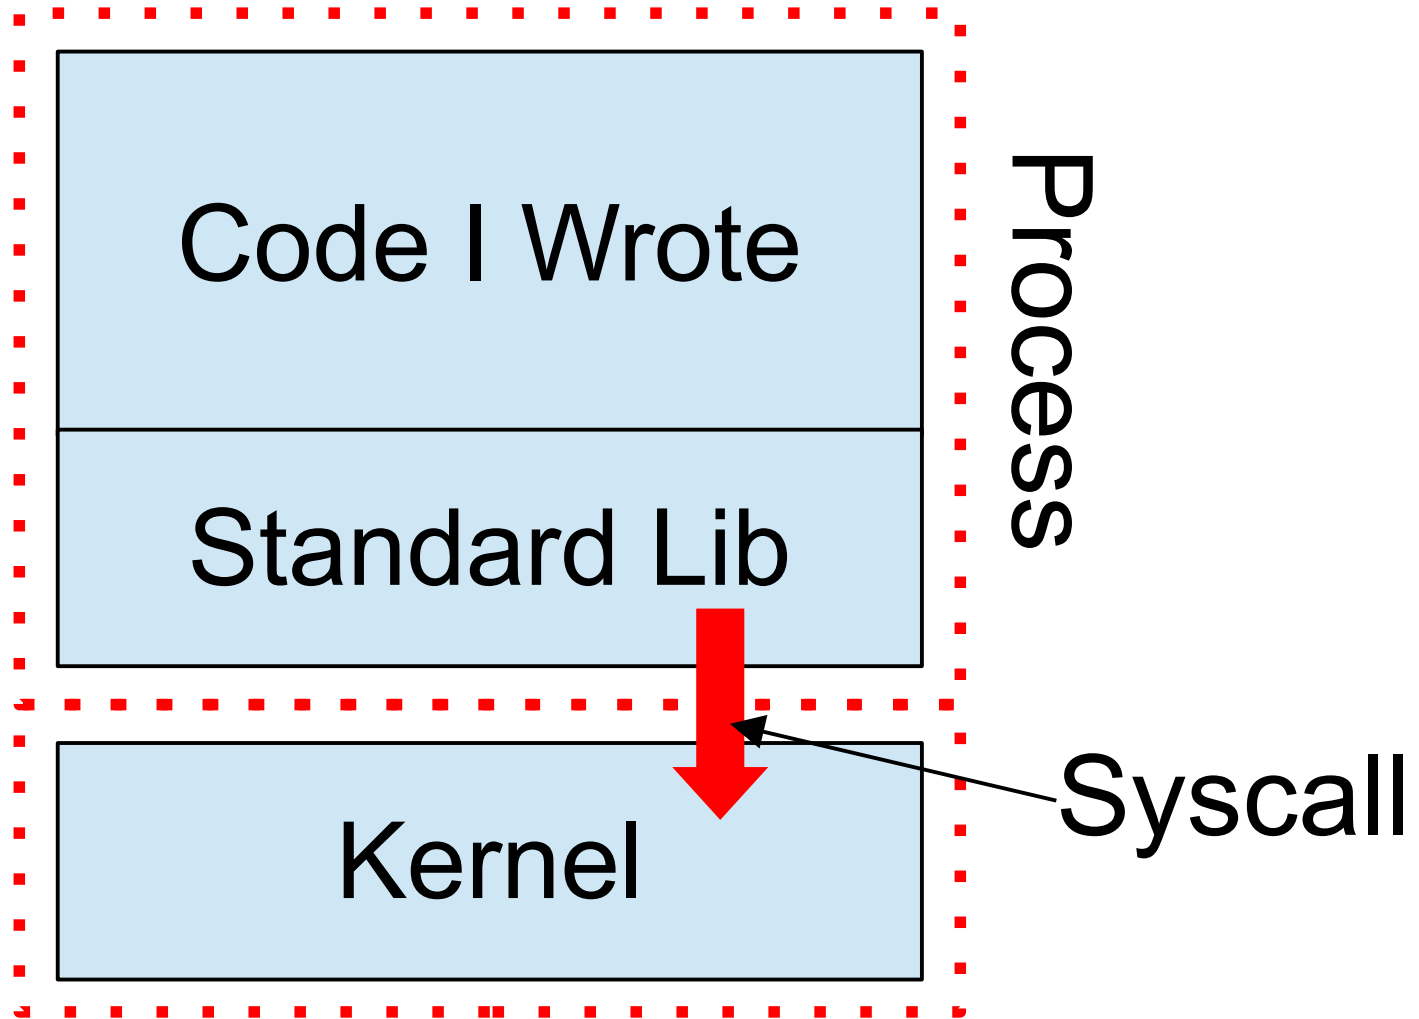
# Syscalls

---

- **syscalls** are special routines that are part of the kernel
  - Often, wrapped with nice functions to make them easy to use
  - But the system call itself is actually known as a “synchronous interrupt”
    - Blocks the CPU, forces you into the kernel
    - But has to happen ***now***, to ***this*** process
      - Not on any other CPU

# Syscalls

---



# Processes

---

## Context Switching


# Multiprocessing

---

- In truth, many processes are on the box at the same time (100s)
  - Have to share CPU
  - Have to (safely!) share memory



# Multiprocessing



# Blocking

---

- Processes will sometimes block for a while
  - Sleeping
  - Waiting for I/O
  - Waiting for communication
    - Inter-process communication
    - Sockets & network
- Other times, they run endlessly
  - Useful or a bug???

# Time Sharing

---

- What happens when multiple processes want the CPU?
- Old, broken idea: hope for the best
  - Can switch processes at any syscall or interrupt. But what if none happen?
- Better: timer interrupt

# Pre-emptive Multitasking

---

- **Timer interrupt** fires 100-1000 times per second
  - Each interrupt is a chance to switch processes, if we wish
  - Processes run for a while (efficiency)
  - But not forever (correctness)
- CPU chooses **scheduling policy**
  - Many options
  - Round-robin common

# Pre-emptive Multitasking

---

- From inside a process it is **impossible** to reliably predict your speed
  - Cache hits & misses
  - Page faults
  - Random interrupts (now including timer!)
  - OS code could run for short or long
  - etc.

# Pre-emptive Multitasking

---

- From inside a process all interrupts feel **random** (even if they are predictable)
- **Any instruction** can be interrupted
- **Impossible** to guarantee “I’ll do this quickly”
  - Could be context-switched out for long time

# Pre-emptive Multitasking

---

## More correct:

- From inside a process, you ***never know*** that interrupts occur
  - Process only cares about its own ***internal, logical state***
  - Individual instructions run fast or slow, and **nobody cares**

# Context Switching

---

- When the OS chooses to remove a process from the CPU, it performs a **context switch**
  - Saves all CPU registers, including program counter
  - Changes virtual memory config
  - Then loads CPU with new info




# Context Switching

---

- CPU perspective
  - Process A was running. Now process B is running
- Process perspective
  - I was running, and then I was frozen in time
  - Later, I was restored to the **exact same state**
- *Truer* process perspective
  - I never stopped running


# Context Switching

Over time, a single CPU does work for two different processes.

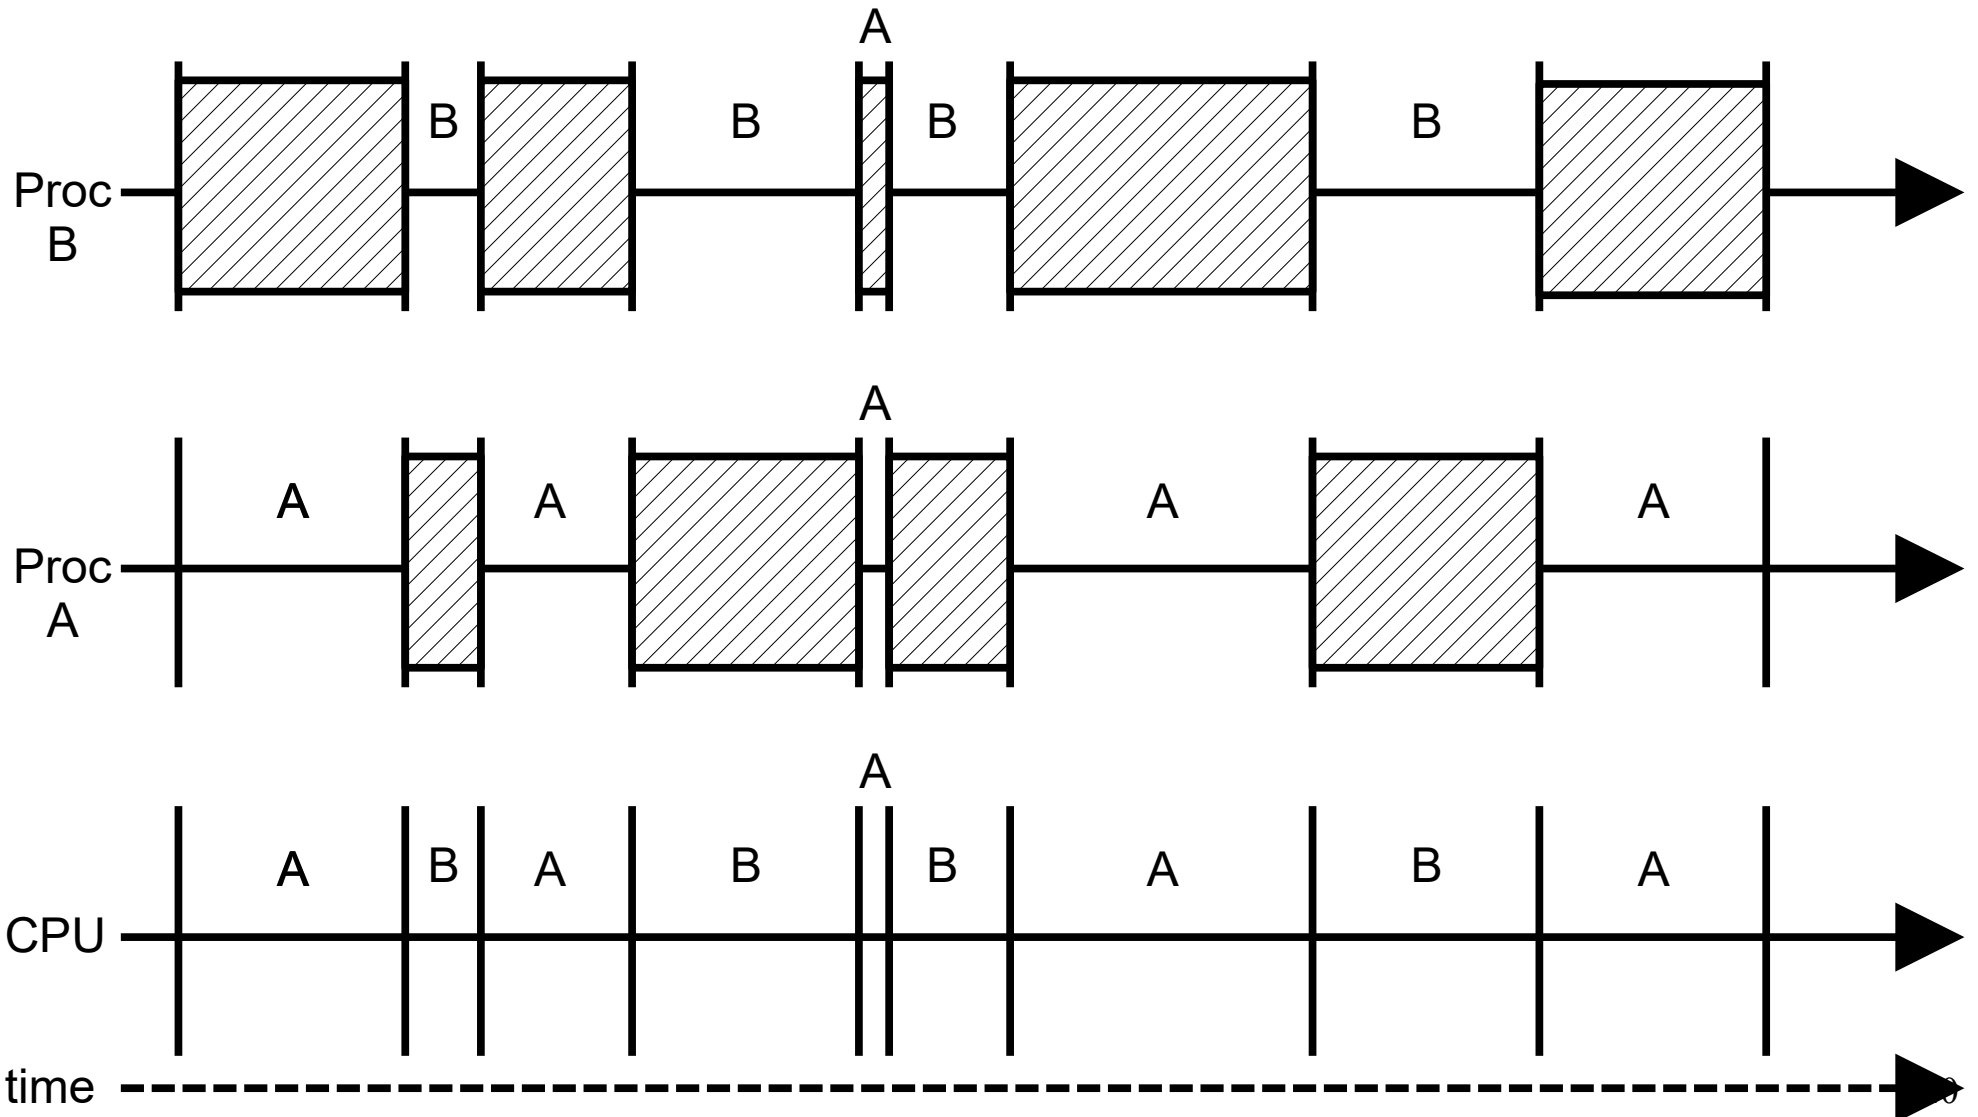


# Context Switching

From the perspective of Process A, it looks like there were long pauses.



# Context Switching



# Context Switching - A Kernel Perspective

---

- Context switching in the user is always **implicit**
  - Ask for an operation to be performed, but it will take time
  - Interrupt occurs unexpectedly
- In the kernel, it has to be **explicit**
  - Some functions will (or might) perform a context switch ***away from your process***

# Context Switching - A Kernel Perspective

---

```
void some_func(...)
{
    USLOSS_Context *oldCtx = ... ;
    USLOSS_Context *newCtx = ... ;
    USLOSS_Console("Before\n");
    USLOSS_ContextSwitch(oldCtx, newCtx) ;
    USLOSS_Console("After\n");
}
```

When does `Before` get printed?  
When does `After` get printed?

**Q:** When does `Before` get printed?

**A:** Immediately!

**Q:** When does `After` get printed?

**A:** A long time from now!

Why is that???

;  
;

```
USLOSS_Console("Before\n");
```

```
USLOSS_ContextSwitch(oldCtx,newCtx);
```

```
USLOSS_Console("After\n");
```

```
}
```

# Processes

---

- When we restore a process after it's been frozen for a while...
  - Restore CPU registers
  - Set up virtual memory
- Process should pick up where it left off
- What does it execute next?
  - The line **after** it performed the context switch!



# Context Switching - A Kernel Perspective

---

```
void some_func(...)
{
    USLOSS_Context *oldCtx = ... ;
    USLOSS_Context *newCtx = ... ;
    USLOSS_Console("Before\n");
    USLOSS_ContextSwitch(oldCtx, newCtx) ;
    USLOSS_Console("After\n");
}
```

**When we come back...**

When did `Before` get printed?

When does `After` get printed?

# When we come back...

**Q:** When did `Before` get printed?

**A:** A long time ago

**Q:** When does `After` get printed?

**A:** Immediately

;  
;

```
USLOSS_Console("Before\n");  
USLOSS_ContextSwitch(oldCtx,newCtx) ;  
USLOSS_Console("After\n");  
}
```

# Processes

---

## Process Control Block

# Process Control Block

---

- A **process control block** stores everything that the kernel knows about a process
- The **process table** is a fixed-length array of PCBs

# Process Control Block

---

- The **process control block** stores:
  - Process ID & name
  - Runnable state
  - Context variables (for saving state)
  - Accounting information
  - Open files & allocated resources
  - ... anything that the kernel knows about the process ...