

Messages

Messages for Communication & Concurrency Control

- IPC
- Our Message Protocol
- Mutexes & Signaling

Messages

IPC

IPC

- **Interprocess communication** (IPC) refers to any of a number of systems which allow processes to talk to each other
 - Messages/mailboxes
 - Signals
 - Pipes, files
 - Semaphores
 - Shared memory

IPC

- IPC usually denotes communication between processes on the same computer
- However, most forms of IPC are also useful for **distributed processes**
 - Processes run on different computers
 - Connected through network
 - Sockets or messaging most common

IPC

- **Mailboxes** allow processes to send **messages** to other processes
 - Atomic messages, small payload
 - Buffering (often)
 - Blocking (often)
 - Structured vs. raw buffers?
 - Non-blocking versions?

IPC

- A **signal** is a software interrupt, sent from one process to another
 - OS must deliver, forces receiver into a **signal handler** function
 - Processes can block signals temporarily (same as disabling interrupts)
 - Various signals, but usually no payload
 - “Default” behavior if no handler
 - Ignore, terminate, core dump, etc.

IPC

- A **pipe** is a file-like mechanism which handles a stream of bytes, but is never stored on disk
 - `stdin, stdout, stderr`
 - `cmd1 | cmd2`
 - Usually only one reader, but many writers OK
 - Block writers when buffer fills

IPC

- **Semaphores** are integers used for counting events, block processes when insufficient resources
 - No payload
 - $P()$ decrements, blocks if insufficient
 - $V()$ increments, releases blocked procs
 - Useful for mutexes, resource control, buffering
 - Slide deck later!

IPC

- **Shared memory** is any time that two processes map the same writable virtual page
 - Probably map to different virtual addresses
 - No linked data structures (lists, trees)
 - Instantaneous data flow, but no way to block
 - Hard to dynamically resize

IPC

- **Non-blocking** functions allow a program to ask if something is available, but refuse to block if not
 - Mostly used for IPC & dist. communication
 - Great for handling multiple simultaneous connections
 - Often, confusing to use
 - Can hurt performance if you don't have some way to sleep (spinning or polling)

Messages

Our Messaging Protocol

Our Messaging Protocol

- In our simulation (Phase 2), we will implement mailboxes and messages
- Mailboxes are created with `MboxCreate()`, and identified by integer IDs
- Each mailbox will have two limits:
 - Max bytes per message (can be 0)
 - Max buffered messages (can be 0)

Our Messaging Protocol

- `MboxSend()` will “buffer” a message when there is no receiver ready to receive it
 - Each mailbox has a max (can be 0)
 - Sender will block if can't buffer
- Buffered messages must be delivered in order
- Blocked senders must be kept in order, too
 - Must buffer in same order they blocked
 - Beware priority problems!

Our Messaging Protocol

- We have a large pool of “message buffers”
 - Global pool, shared by all mailboxes
 - Allocate when you have a message
 - In `MboxSend()`, not `MboxCreate()`
 - Over-allocation is allowed
 - Buffering can fail if no free slots
 - Free in `MboxRecv()`

Our Messaging Protocol

- `MboxRecv()` will receive one message
 - Block if nothing buffered
- Messages must be received in order
- Blocked receivers must be kept in order
 - Beware priority problems here, too!

Our Messaging Protocol

- `MboxCondSend()` , `MboxCondSend()` are **non-blocking** versions
 - Send or recv if possible
 - Return special value if you would have blocked

Messages

Mutexes & Signaling

Mutexes & Signaling

- A mailbox can be used (abused?) as a mutex
- Version A:
 - Initialize with no messages
 - `Send()` to lock, `Recv()` to unlock
 - Class question: What blocks, and why?
- Version B:
 - Class question: Can you reverse it?

Mutexes & Signaling

- Soon, we'll be implementing syscalls
 - Need to block current process until something happens
- Class discussion: how to use mailboxes for sleep & wakeup
 - Nicier interface than calling Phase 1 code
 - How to do it?
 - When to allocate mailboxes?
 - What can be shared, or not?

Mutexes & Signaling

- Sometimes, we need a way to communicate that an event occurred to another process
 - Example: disk interrupt
- Class discussion: how to use mailboxes?
 - What if the receiver is busy?
 - What if nobody is listening?
 - Is it permissible to block in an interrupt handler?

Mutexes & Signaling

- Sometimes, we want to communicate a regular sequence, like a clock – perhaps to many processes
- Class discussion: how to wake up many procs?
 - What if nobody is listening?
 - Do we want to buffer or lose “missed” events?