

# Topic 9: Synchronization via Semaphores

- Requirements for a mutual exclusion mechanism:
  - Must allow only one process into a critical section at a time.
  - If several requests at once, must allow one and only one process to proceed.
  - Processes must be able to go on vacation outside the critical section.
  - Must handle an arbitrary number of processes.
- Desirable properties for a mutual exclusion mechanism:
  - Fair: if several processes waiting, let each in eventually.
  - Efficient: do not use up substantial amounts of resources when waiting. E.g., no busy waiting.
  - Simple: should be easy to use (e.g., just bracket the critical sections).
- Desirable properties of processes using the mechanism:
  - Always lock before manipulating shared data.
  - Always unlock after manipulating shared data.
  - Do not lock again if already locked.
  - Do not unlock if not locked by you (there are a few exceptions to this...).
  - Do not spend large amounts of time in critical section.

- Semaphore
  - Synchronization variable that takes on non-negative integer values.
  - Invented by E. Dijkstra in the mid-60's.
  - Patrick's train story here...



- **P( semaphore ):**
  - Atomic operation that waits for semaphore to become positive, then decrements it by 1.
  - “Proberen” in Dutch (“try”)
- **V( semaphore ):**
  - Atomic operation that increments semaphore by 1.
  - “Verhogen” in Dutch (“increase”)
- Semaphores are simple and elegant. They allow the solution of many interesting problems.
- They can do a lot more than just mutual exclusion.

- Too much milk problem with semaphores:

***Processes A & B***

```
1    P( OKToBuyMilk );  
2    if ( noMilk ) {  
3        BuyMilk();  
4    }  
5    V( OKToBuyMilk );
```

- Notes: **OKToBuyMilk** must initially be set to 1.
  - What happens if it is set to zero? set to two? etc.
- *Binary Semaphores*
  - Have only two possible values: 0 and 1.
  - Used for mutual exclusion.
- *Counting Semaphores*
  - Have any non-negative value (  $\geq 0$  ).
  - Used for scheduling
    - To wait for some event to happen.
    - To wait for some resource to become available.

- Semaphore properties:
  - + Machine independent.
  - + Simple (to use).
  - + Works with many processes (not just two).
  - + Can have many critical sections in the code, each with its own semaphore.
  - + Can acquire many resources simultaneously (multiple P's).

```

/* I want three of these to go, please... */
P( pizza );
P( pizza );
P( pizza );
... using the resources to make 3 pizzas ...
V( pizza );
V( pizza );
V( pizza );

```

- What happens if only 2 pizza resources are available?
- Not provided by hardware.

Semaphores are used in two different ways:

- **Mutual Exclusion:**

- Ensure that only one process is accessing shared information at a time.
- If there are separate groups of data that can be accessed independently, there may be separate semaphores, one for each group of data.
  - These are also binary semaphores.

- **Scheduling:**

- To permit processes to wait for certain things to happen.
- If there are different groups of processes waiting for different things to happen, there will usually be a different semaphore for each group. These semaphores are not necessarily binary semaphores.
- Resource allocation is an important form of scheduling.
  - If there are N items of a resource, then set an associated semaphore to N initially.
  - P to acquire the resource.
  - V to release the resource.
  - Processes block (on P) when the resource runs out and wait for more to be released (V).

### Semaphore Example: Producer — Consumer:

- Suppose one process is creating information that is going to be used by another process.
  - E.g., one process is reading a program from the disk, and another process will compile the program.
  - E.g., one process is building a list of processes, and another process is printing only those processes owned by you.
  - E.g., Unix pipes
- Processes should not have to operate in perfect lock-step: producer should be able to get ahead of consumer.
- Producer: creates copies of a resource.
- Consumer: uses up (destroys) copies of a resource.
- Buffers: used to hold information after producer has created it, but before consumer has used it.
- Synchronization: keeping producer ahead of consumer.
- Define constraints (definition of what is “correct”).
  - Consumer: if all buffers are empty, must wait for producer to fill buffer(s). Scheduling.
  - Producer: if all buffers are full, must wait for consumer to empty buffer(s). Scheduling.
  - Only one process may manipulate buffer pool at a time. Mutual exclusion.
  - A separate semaphore is used for each constraint ( = 3 semaphores).

- Initialization:
  - Put all buffers in a pool of empties.
  - Initialize semaphores (specify the initial value of the semaphore when creating it):

```
semaphore empties = sem_create( NUM_BUFFERS );
semaphore fulls = sem_create( 0 );
semaphore mutex = sem_create( 1 ); /* 'mutex' = mutual exclusion */
```

- The solution:

Producer	Consumer
<pre>P( empties ); P( mutex );   &lt;get empty buffer from pool of empties&gt; V( mutex );   &lt;put data in buffer&gt; P( mutex );   &lt;add full buffer to pool of fulls&gt; V( mutex ); V( fulls );</pre>	<pre>P( fulls ); P( mutex );   &lt;get full buffer from pool of fulls&gt; V( mutex );   &lt;consume data in buffer&gt; P( mutex );   &lt;add empty buffer to pool of empties&gt; V( mutex ); V( empties );</pre>

- Questions:
  - Why does producer **P( empties )** but **v( fulls )**?
  - Why are data in the buffer accessed outside the critical section?
  - Is the order of first two P's important in each case?
  - Is the order of the last two V's important in each case?
  - Could we have separate semaphores for each pool?
  - How would this be extended to have two consumers?
  - How about two producers?



### Semaphore Example: Readers — Writers:

- Shared database with readers and writers.
- It is safe to have several readers access the database simultaneously.
- A writer must have exclusive access.
  - No other writers.
  - No other readers.
- Invariant:  
$$(\text{ActiveWriters} == 1 \ \&\& \ \text{ActiveReaders} == 0) \ || \ (\text{ActiveWriters} == 0 \ \&\& \ \text{ActiveReaders} \geq 0)$$
- Use semaphores to enforce these policies.
  - Note: Writers are actually readers too.
- Constraints:
  - Writers can only proceed if there are no active readers or writers. Use semaphore **OKToWrite**.
  - Readers can only proceed if there are no active writers.
    - Implies that readers get priority.
  - Need a shared variable to keep track of who is reading.
    - Must ensure that only one process manipulates shared variable at once — need a mutex semaphore.

- State variables
  - **ActiveReaders** = number of active readers.
  - **WaitingReaders** = number of waiting readers.
  - **ActiveWriters** = number of active writers.
  - **WaitingWriters** = number of waiting writers.
  - **ActiveWriters** is always 0 or 1.
  - **ActiveReaders** and **ActiveWriters** may not both be non-zero at the same time.
- Initialization:

```
semaphore OKToRead = create_sem(0);
semaphore OKToWrite = create_sem(0);
semaphore Mutex = create_sem(1);

int ActiveReaders, WaitingReaders, ActiveWriters, WaitingWriters;
ActiveReaders = WaitingReaders = ActiveWriters = WaitingWriters = 0;
```
- Scheduling
  - Writers-preference.
  - Note: readers-preference and no-preference solutions are also possible.

Reader Processes	Writer Processes
<pre> P( Mutex ); if ( (ActiveWriters == 0) &amp;&amp;       (WaitingWriters == 0) ) {     V( OKToRead );     ActiveReaders++; } else     WaitingReaders++; V( Mutex );  P( OKToRead ); &lt;read the data...&gt;  P( Mutex ); ActiveReaders--; if ( (ActiveReaders == 0) &amp;&amp;       (WaitingWriters &gt; 0) ) {     /* wake up a writer */     V( OkToWrite );     ActiveWriters++;     WaitingWriters--; } V( Mutex ); </pre>	<pre> P( Mutex ); if ( (ActiveWriters == 0) &amp;&amp;       (ActiveReaders == 0) &amp;&amp;       (WaitingWriters == 0) ) {     V( OKToWrite );     ActiveWriters++; } else     WaitingWriters++; V( Mutex );  P( OKToWrite ); &lt;write data...&gt;  P( Mutex ); ActiveWriters--; if ( WaitingWriters &gt; 0 ) {     V( OKToWrite ); /* wake up a writer */     ActiveWriters++;     WaitingWriters--; } else { /* wake up <u>all</u> the waiting readers */     while ( WaitingReaders &gt; 0 ) {         V( OKToRead );         ActiveReaders++;         WaitingReaders--;     } } V( Mutex ); </pre>

- Consider several examples:
  - Reader enters and leaves system.
  - Writer enters and leaves system.
  - Two readers enter system.
  - Writer enters system (and has to wait).
  - Reader enters system (wait? what? where?).
  - Two readers leave system — writer does what?
- Questions:
  - In case of conflict between readers and writers, who gets priority? What lines of the code help insure this?
  - Is the **waitingWriters** necessary in the writer's first **if**?
  - Can **OKToRead** ever be greater than 1?
  - Can **OKToWrite** ever be greater than 1?
  - Is the first writer to execute **P( Mutex )** guaranteed to be the first writer to access the data?

## Semaphore Implementation.

- No existing hardware implements semaphores directly.
  - All involve some sort of CPU scheduling; it is not clear that scheduling decisions should be made in hardware anyway.
- Must be built up in software using a lower-level synchronization primitive provided by hardware.
- Need a simple way of doing mutual exclusion in order to implement P's and V's.
- Could use atomic reads and writes, as in the too-much-milk problem. Clumsy!
- Uniprocessor solution: disable interrupts.

```
typedef struct {
    int count;
} Semaphore;

P( Semaphore *sem ) {
    while (1) {
        interruptsDisable();
        if ( sem->count > 0 ) {
            sem->count--;
            break;
        }
        interruptsEnable();
    }
    interruptsEnable();
} /* P */
```

```
V( Semaphore *sem ) {
    interruptsDisable();
    sem->count++;
    interruptsEnable();
} /* V */
```

- Step 1:
  - When P fails, put process to sleep.
  - On V just wake up everybody; processes all try P again.
  - Is this “busy-waiting”?
- Step 2:
  - Label each process with the semaphore it is waiting for.
  - Then, just wake up relevant processes.
- Step 3:
  - Just wake up a single process.
- Step 4:
  - Add a queue of waiting processes to the semaphore, wake up FCFS.
  - On failed P, add to queue.
  - On V, remove from queue.

```

typedef struct {
    int count;
    queue q;
} Semaphore;

P( Semaphore *sem ) {
    interruptsDisable();
    if ( sem->count > 0 ) {
        sem->count--;
        interruptsEnable();
    }
    else {
        Move process fm ready list to sem->q;
        interruptsEnable();
        dispatcher();
    }
} /* P */

V( Semaphore ) {
    interruptsDisable();
    if ( sem->q not empty ) {
        Move first process fm sem->q to
        ready list
        dispatcher();
    }
    else
        sem->count++;
    interruptsEnable();
} /* V */

```

- There are several trade-off's implicit here:
  - How many processes are in the system?
  - How much queueing on semaphores, storage requirements, etc.?
  - Most important thing: avoid busy-waiting.

- What do we do in a multiprocessor to implement P's and V's? Cannot just turn off interrupts to get low-level mutual exclusion.
  - Turn off all other processors?
  - Use atomic read and write, as in too-much-milk?
- In a multiprocessor, there will have to be busy-waiting at some level: cannot block without mutual exclusion.
  - Marking the process as blocked and blocking it must be atomic.
- Most machines provide some sort of atomic read-modify-write instruction. Read existing value, store back in one atomic operation.
  - Example: Test-and-set — TAS
    - Implemented initially by IBM, later by many others.
  - The TAS instruction sets the value to one, but returns old value.
    - Use ordinary write to set back to zero.
  - Using test-and-set for mutual exclusion:
    - Like a binary semaphore in reverse (but does not include waiting).
    - 1 = someone else is already using it.
    - 0 = okay to proceed.
    - By definition, two (or more) processes are prevented by TAS from getting a 0-to-1 transition simultaneously.

```

la    $s2, memAddr
lw    $v0, 0($s2)
addi   $s1, $zero, 1
sw    $s1, 0($s2)

```



- Test-and-set is tricky to use:
  - Cannot get at it from HLLs.
  - Typically, must use a routine written in assembler.
- Read-modify-write's may be implemented:
  - Directly in memory hardware (e.g., IBM 360), or
  - In the processor by refusing to release the memory bus (PDP-11).
- Using test-and-set to implement semaphores in a multiprocessor:
  - For each semaphore, keep a test-and-set integer in addition to the semaphore integer and the queue of waiting processes.
  - This integer is often called a *spin-lock* since the processes will busy-wait (or *spin*) on it.

```

typedef struct {
    int count;
    queue q;
    int test;      // set to 0 by SemCreate
} Semaphore;

```

```

P( Semaphore *sem ) {
    while (1) {
        interruptsDisable();
        while ( TAS( &sem->test ) != 0 )
            ; /* do nothing */
        if ( sem->count > 0 ) {
            sem->count--;
            break;
        }
        Enqueue process on sem->q
        sem->test = 0; /* ordinary assignment */
        interruptsEnable();
        dispatcher(); /* cannot return! */
    }
    sem->test = 0;
    interruptsEnable(); /* return okay */
} /* P */

```

```

V( Semaphore *sem ) {
    interruptsDisable();
    while ( TAS( &sem->test ) != 0 )
        ; /* do nothing */
    sem->count++;
    if ( sem->q is not empty ) {
        Dequeue first process from sem->q;
        Wake it up;
    }
    sem->test = 0; /* ordinary assignment */
    interruptsEnable();
} /* V */

```

- Important point: Implement some mechanism once, very carefully. Then, always write programs that use that mechanism.
  - Importance of *layering*.
- Why do we still have to disable interrupts in addition to using test-and-set?
  - Semaphores synchronize *processes*.
  - An interrupt handler is not a separate process; it is a procedure run by the interrupted process.
  - What if the interrupted process was in a critical section and the interrupt handler tries to enter the same section? Neither will make progress: deadlock!
- Semaphores cannot be used to provide mutual exclusion between interrupt handlers and processes.
  - On a uniprocessor, you must disable interrupts.
  - On a multiprocessor, you must disable interrupts and busy-wait (spin-lock).
  - Different forms of mutual exclusion complicate the code.
- Interrupt handlers can only V semaphores (unless you code very carefully).
  - Provides a limited form of scheduling