

An Introduction to Virtual Memory

Russell Lewis

A Virtual Computer


Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1



Memory

0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

A Virtual Computer

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

17

Memory

0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

A Virtual Computer

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1



Memory

0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

A Virtual Computer


Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1



Memory

0	17
1	18
2	0
3	0
4	0
5	0
6	0
7	0

A Virtual Computer

Memory	
0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

Process B

```
load r1, mem(0)  
mul  r1, r1, r1  
store r1, mem(1)
```

r1

0

A Virtual Computer

Memory

0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1

17

A Virtual Computer

Memory	
0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1

289

A Virtual Computer

Memory

0	17
1	289
2	0
3	0
4	0
5	0
6	0
7	0

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1


289

A Virtual Computer

Process A

```
load r1, mem(0)
inc  r1
store r1, mem(1)
```

r1




Memory

0	17
1	0
2	0
3	0
4	0
5	0
6	0
7	0

Process B

```
load r1, mem(0)
mul  r1, r1, r1
store r1, mem(1)
```

r1



A Virtual Computer

Process A

```
load r1, mem(0)
inc  r1
store r1, mem(1)
```

r1

18

Memory

0	17
1	????
2	0
3	0
4	0
5	0
6	0
7	0

Process B

```
load r1, mem(0)
mul  r1, r1, r1
store r1, mem(1)
```

r1

289

Isolation

- How to keep processes from corrupting each other's memory?

Page Tables

Process A

load r1, mem(0)
inc r1
store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	17

Process B

load r1, mem(0)
mul r1, r1, r1
store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	17

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	0
5	0
6	0
	17

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)
inc r1
store r1, mem(1)

r1 **17**

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	0
5	0
6	0
	17

Process B

load r1, mem(0)
mul r1, r1, r1
store r1, mem(1)

r1 0

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)
inc r1
store r1, mem(1)

r1 17

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	17

Process B

load r1, mem(0)
mul r1, r1, r1
store r1, mem(1)

r1 17

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

18

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	17

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1

289

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

18

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	18
5	0
6	0
7	17

Process B

load r1, mem(0)

mul r1, r1, r1

store r1, mem(1)

r1

289

Page Table

virt	phys
0	7
1	6

Page Tables

Process A

load r1, mem(0)
inc r1
store r1, mem(1)

r1

18

Page Table

virt	phys
0	7
1	4

Memory

0	0
1	0
2	0
3	0
4	18
5	0
6	289
7	17

Process B

load r1, mem(0)
mul r1, r1, r1
store r1, mem(1)

r1

289

Page Table

virt	phys
0	7
1	6

Virtual Address Space

- What memory can each process see?

Virtual Address Space

Virtual Memory

Process A

load r1, mem(0)
inc r1
store r1, mem(1)

r1

18

0

1

17

18

17

289

Process B

load r1, mem(0)
mul r1, r1, r1
store r1, mem(1)

r1

289

The Flat Address Model

- Every process has its own isolated memory
- Every address is usable
 - No holes
 - No limits
- App/OS decide what addresses are in use
 - Things can go **anywhere!**

Implementation

- Must be fast!
- Must be flexible!

Implementation

- Must be fast!
- Must be flexible!
- Implemented by CPU, configured by OS
(page tables)
- CPU asks OS what to do when it's confused
(page fault)

The Mechanics of Translation

Process A

```
load r1, mem(0x1234)  
inc r1  
store r1, mem(0x1235)
```

Virtual Address

0x1234

Virtual Page #

0x123

0x4

Page Offset

Physical Page #

0xabc

0x4

Page Offset

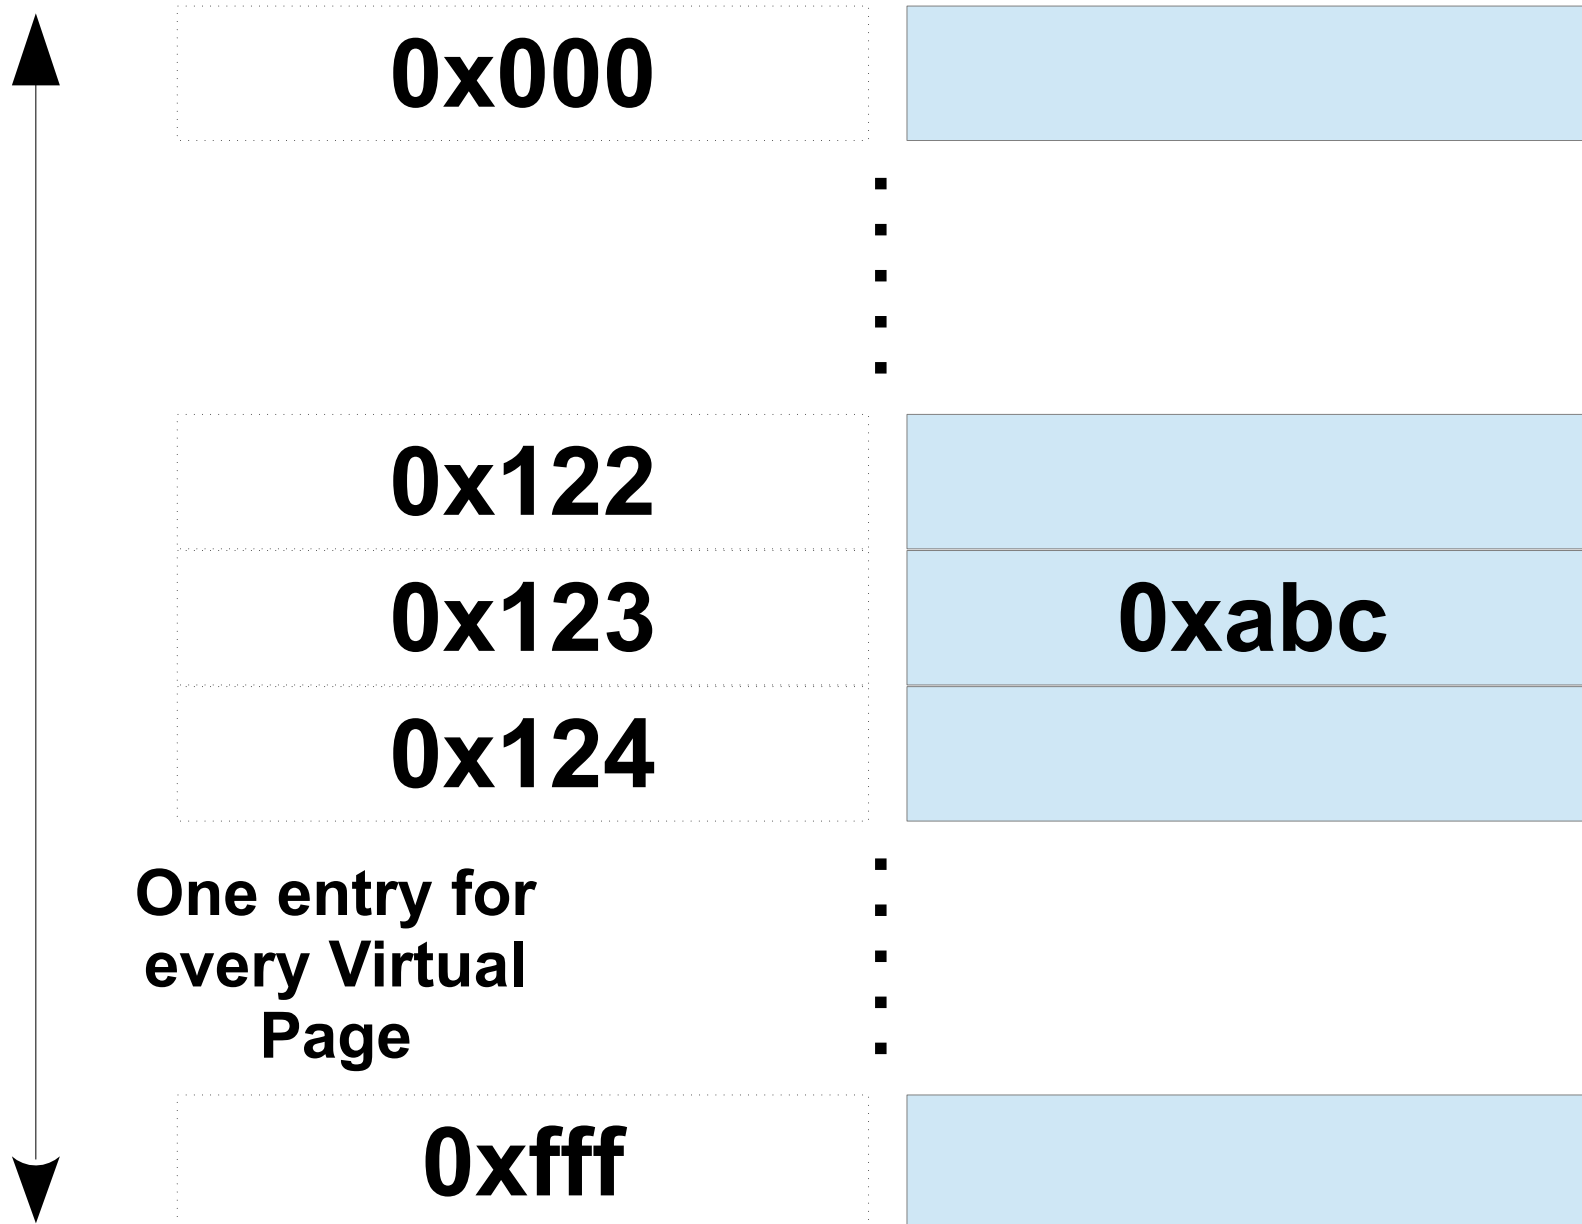
Physical Address

0xabc4


The Page Table

0x000	
⋮	
0x122	
0x123	0xabc
0x124	
⋮	
0xfff	

The Page Table



The Page Table



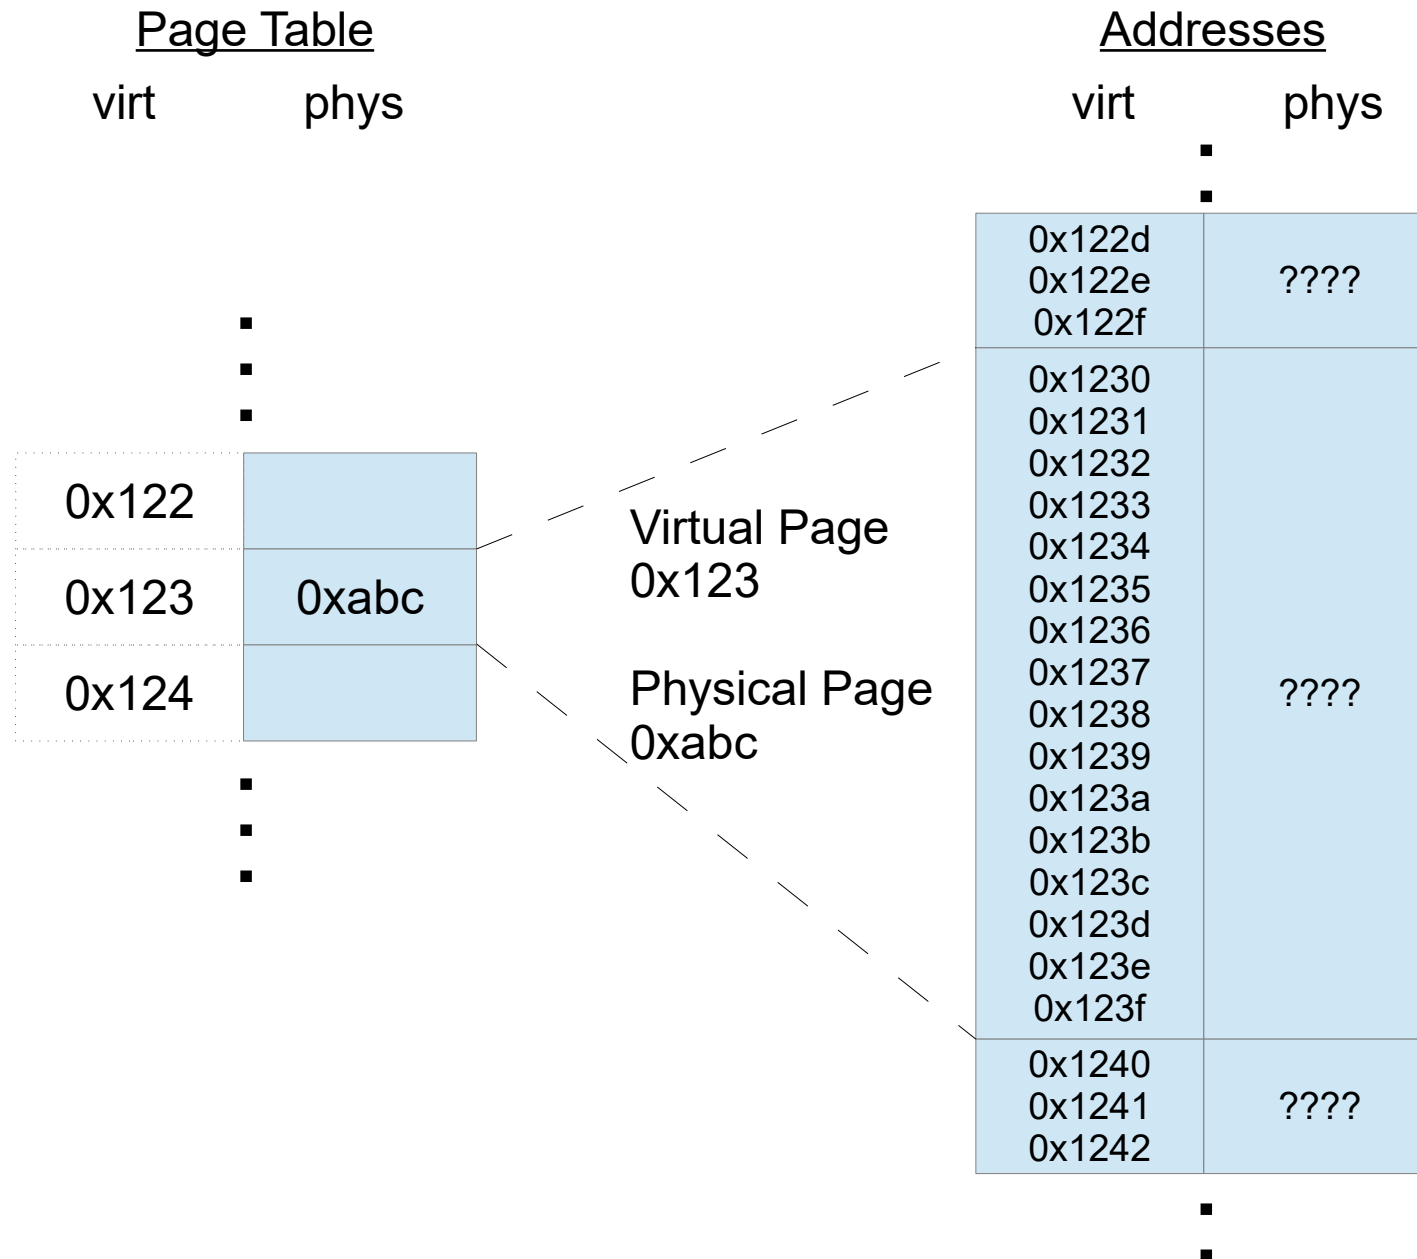
The Mechanics of Translation

Page Table

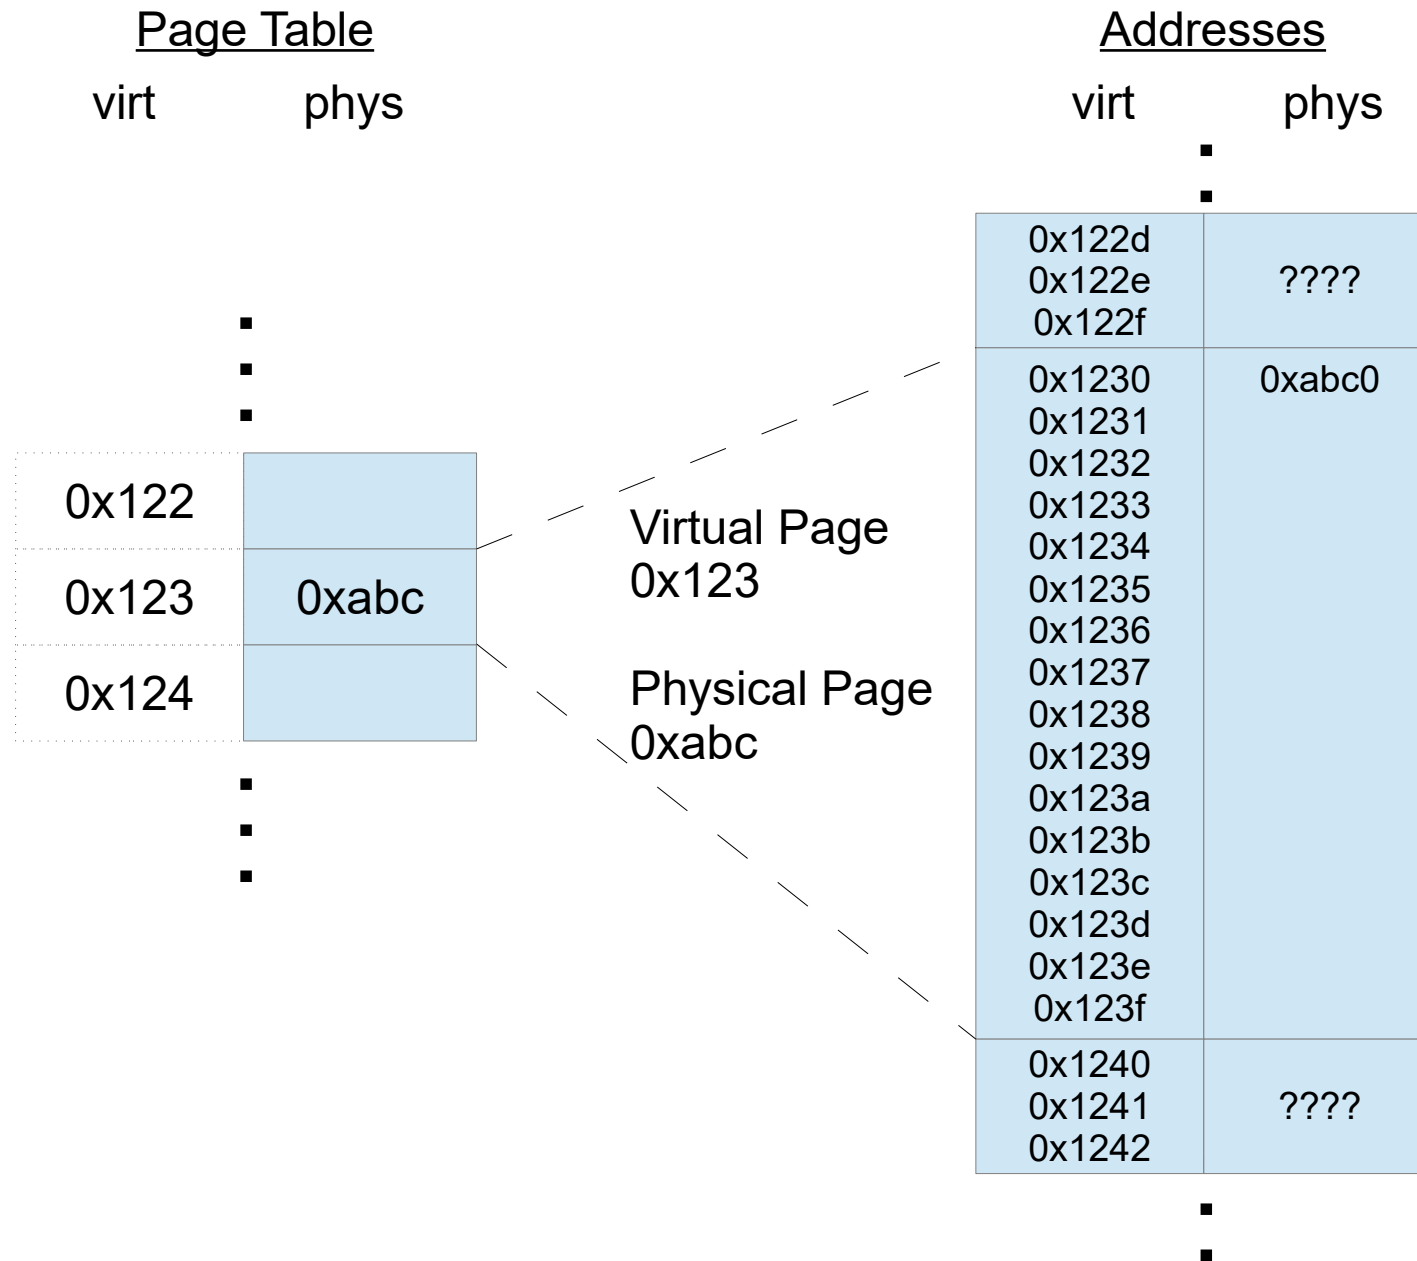
virt phys

	▪
	▪
	▪
0x122	
0x123	0xabc
0x124	
	▪
	▪
	▪

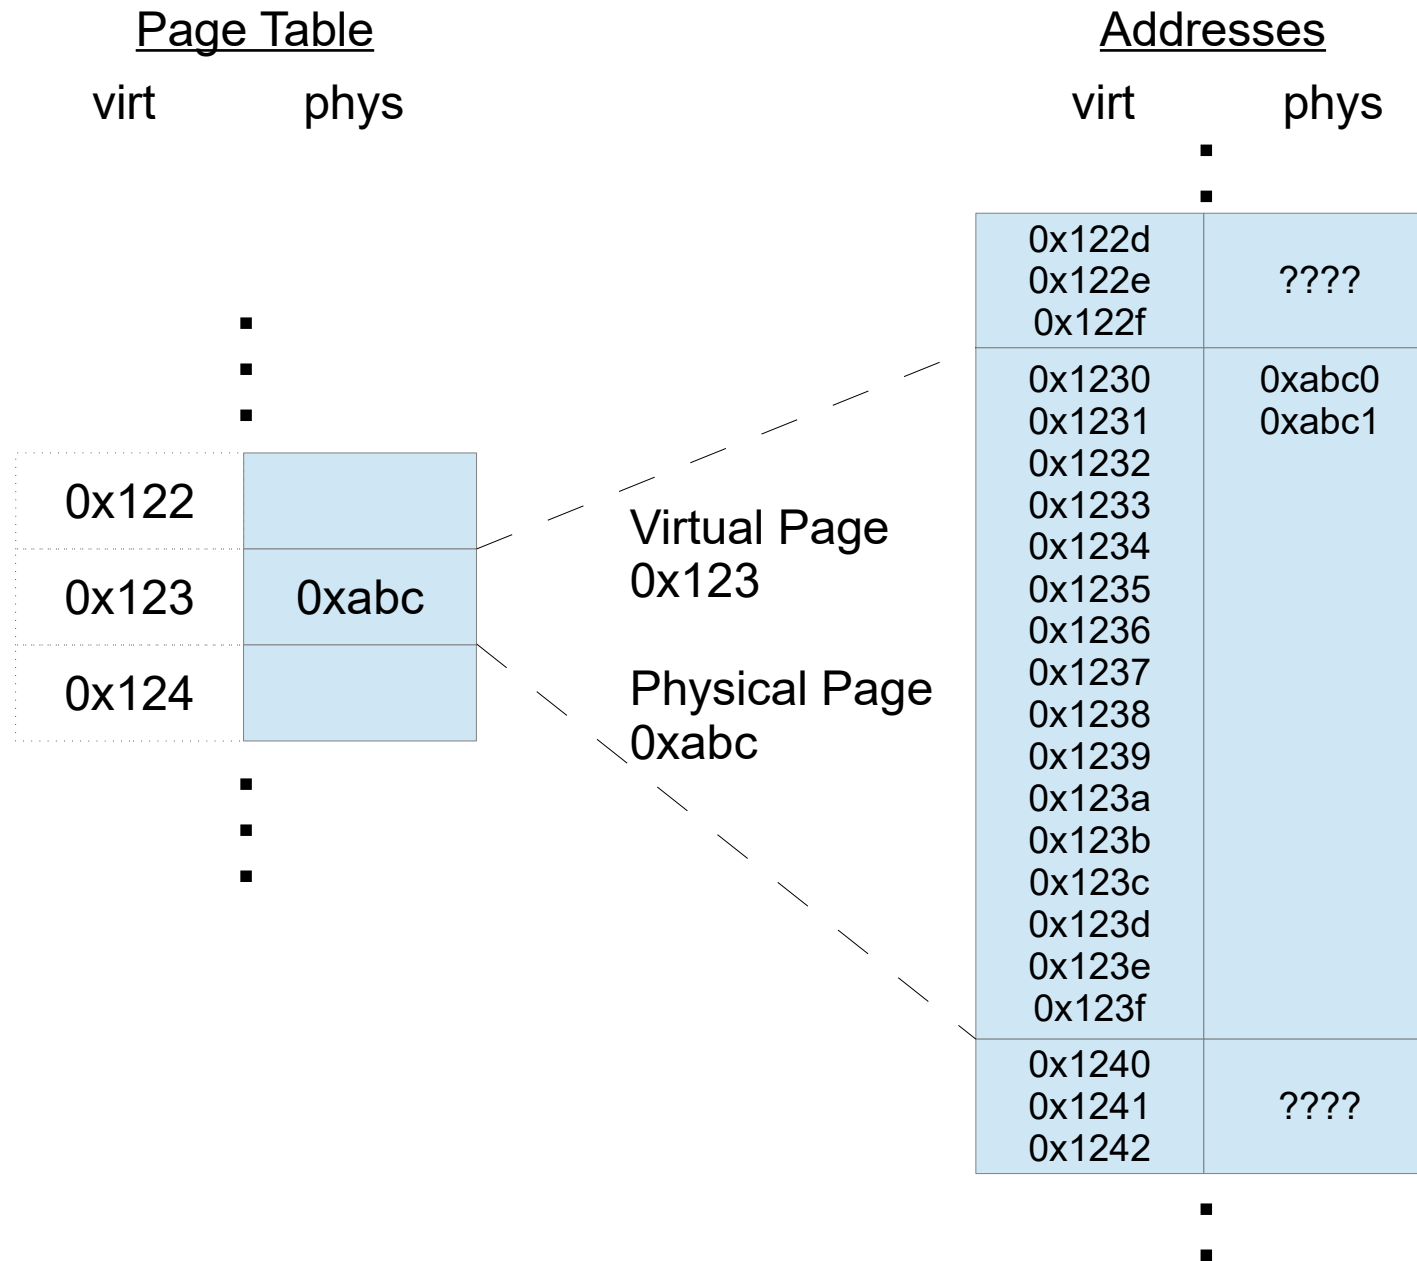
The Mechanics of Translation



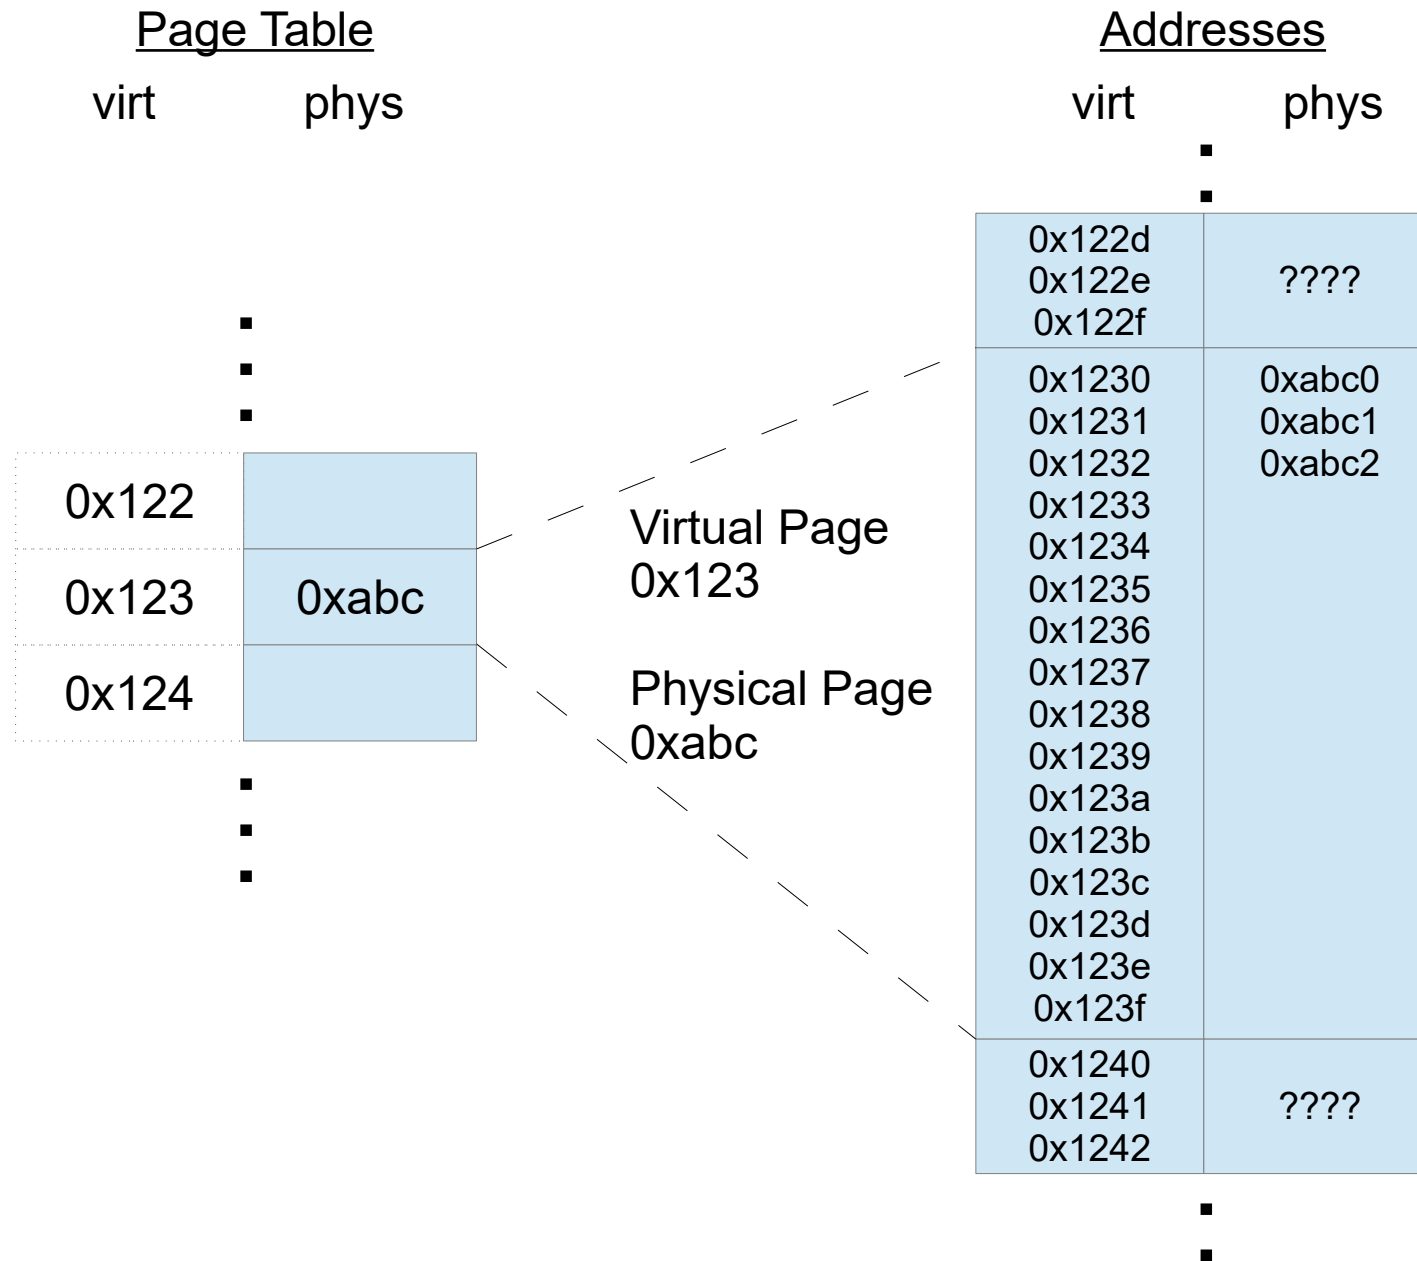
The Mechanics of Translation



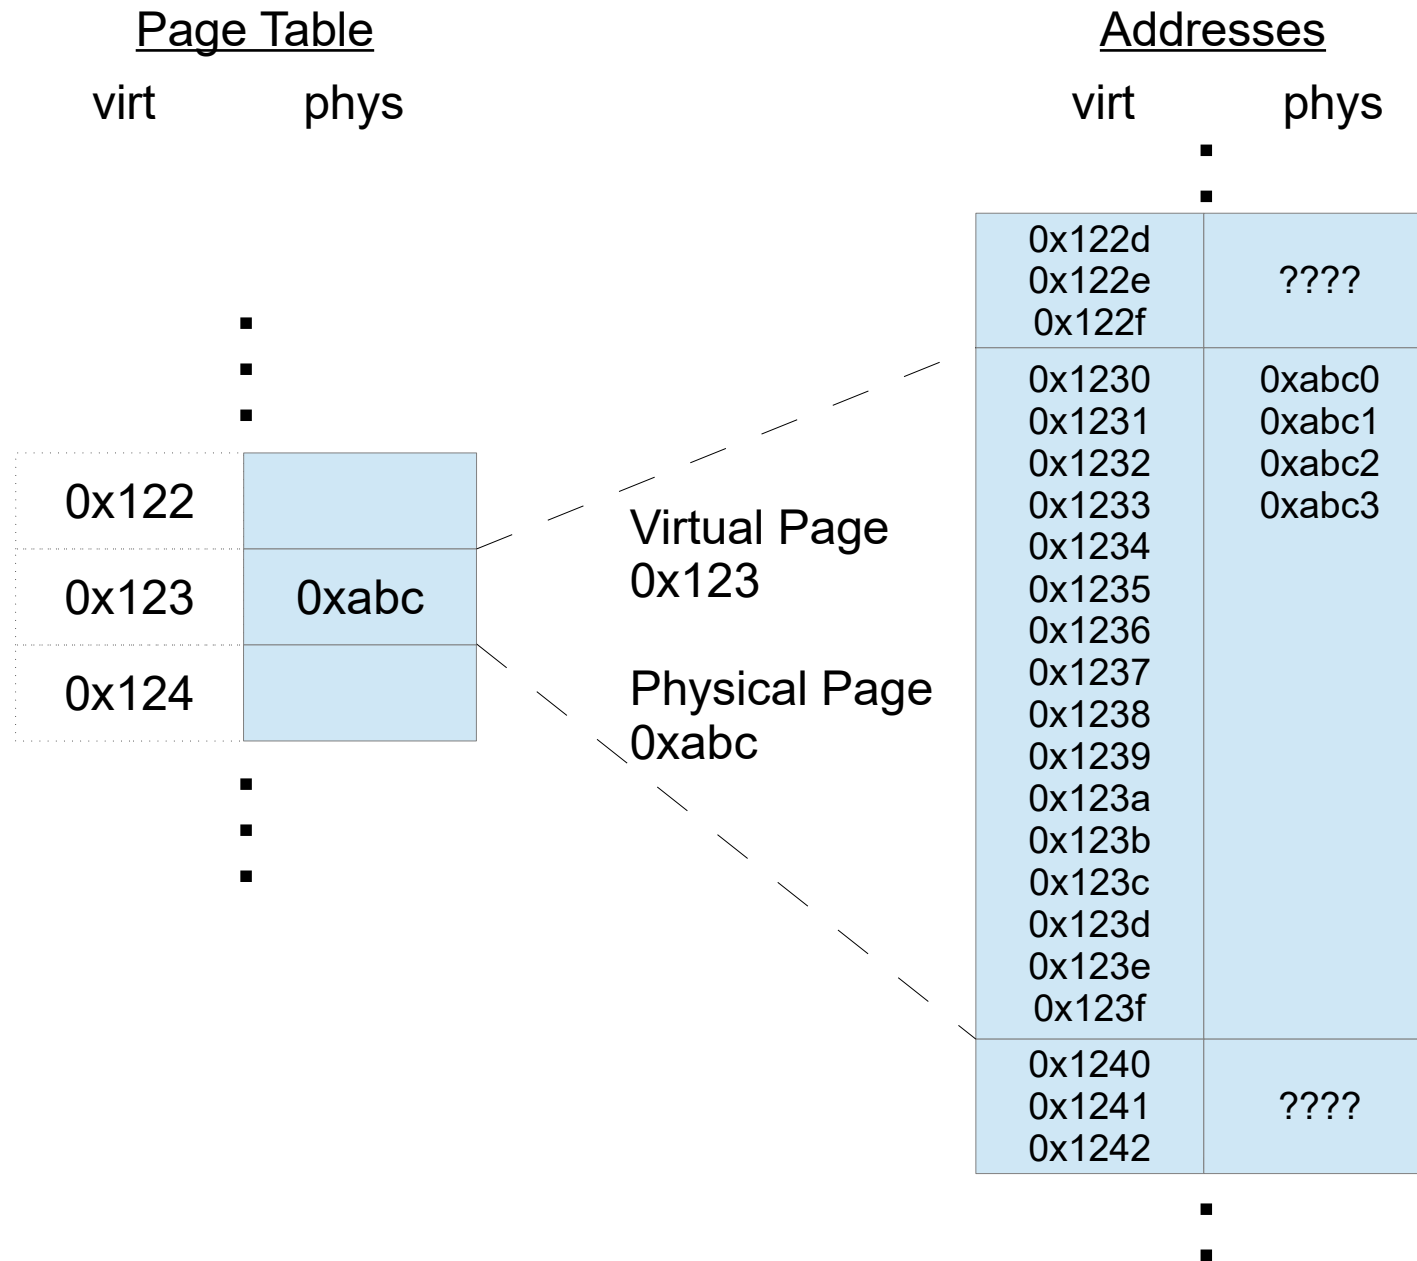
The Mechanics of Translation



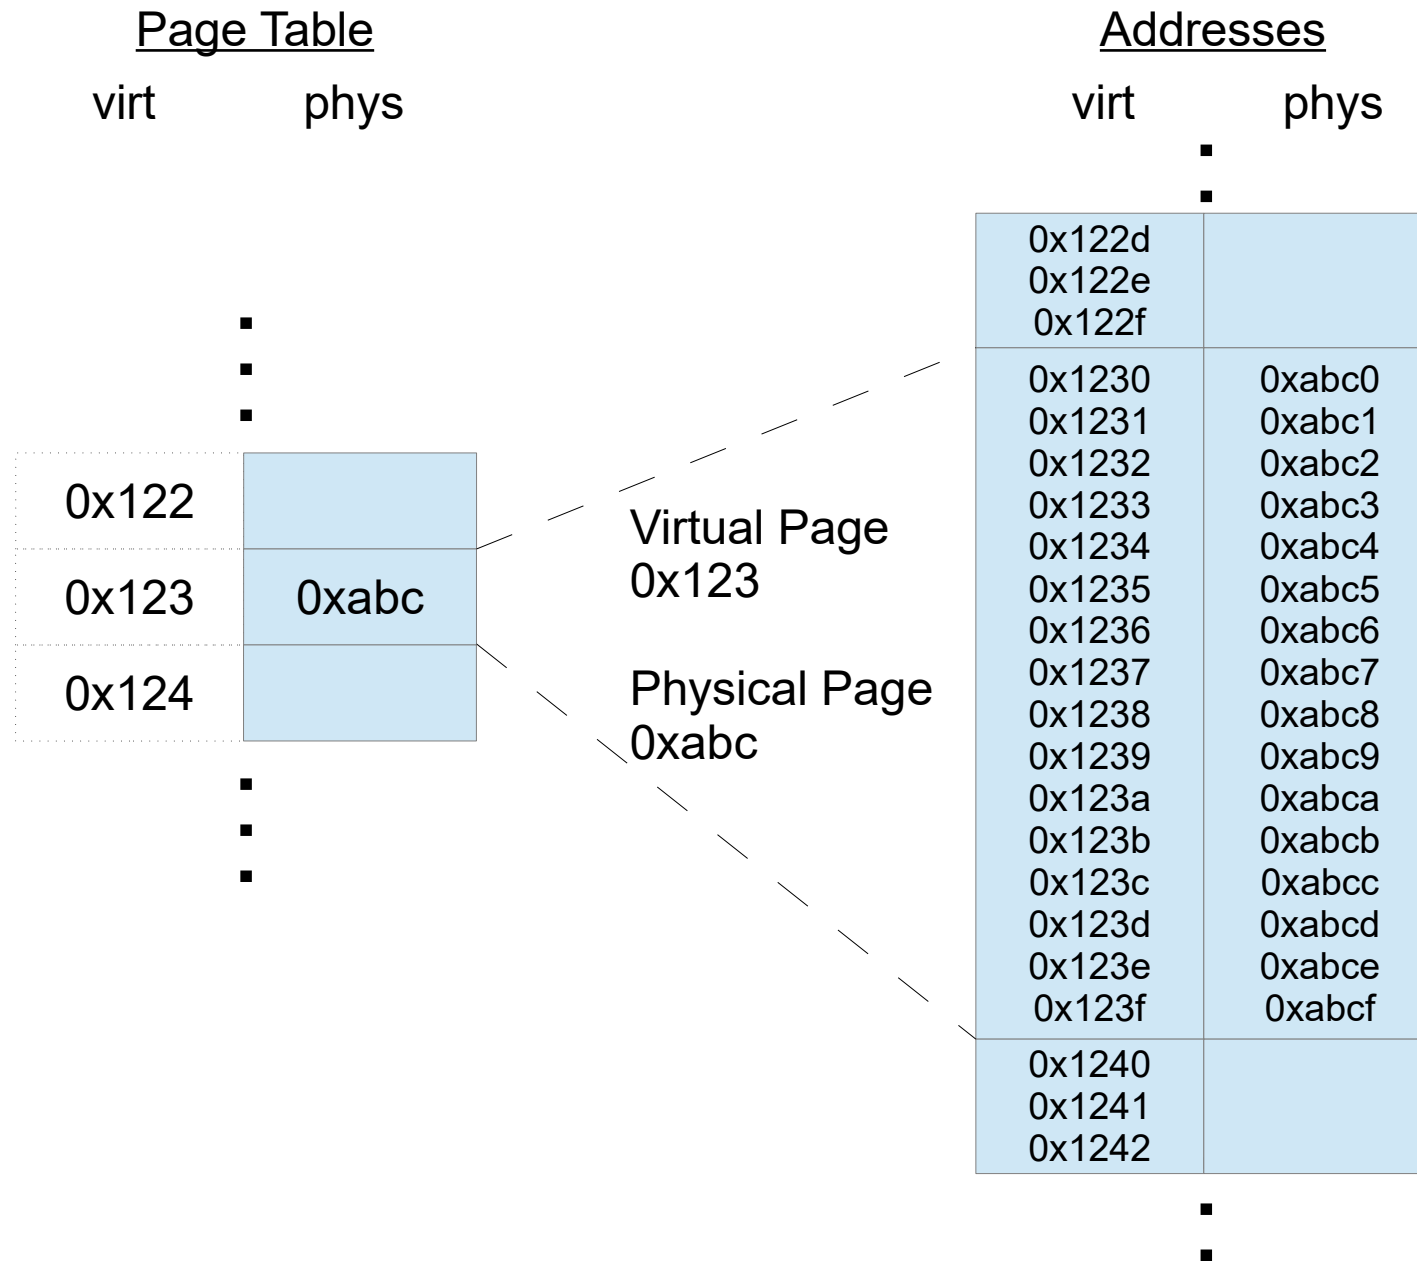
The Mechanics of Translation



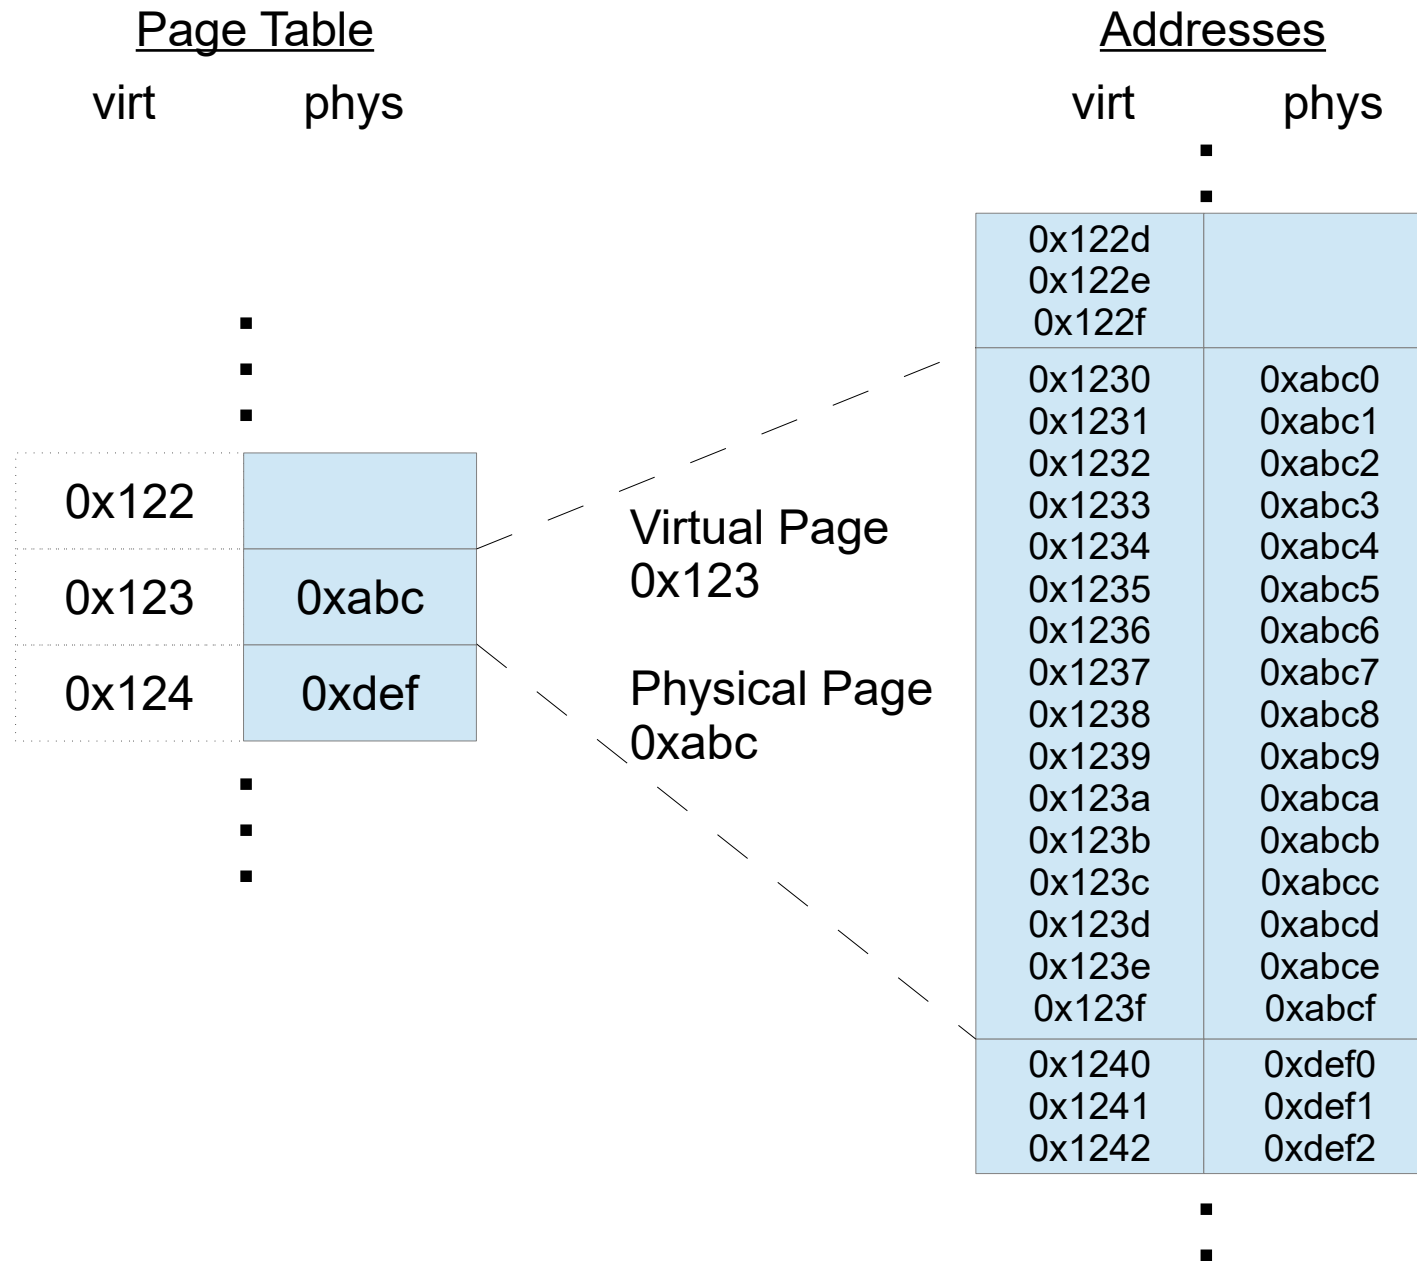
The Mechanics of Translation



The Mechanics of Translation



The Mechanics of Translation



The Mechanics of Translation

<u>Page Table</u>	
virt	phys
▪	
▪	
▪	
0x122	
0x123	0xabc
0x124	0xdef
▪	
▪	
▪	

<u>Addresses</u>	
virt	phys
▪	
▪	
0x122d	
0x122e	
0x122f	
0x1230	0xabc0
0x1231	0xabc1
0x1232	0xabc2
0x1233	0xabc3
0x1234	0xabc4
0x1235	0xabc5
0x1236	0xabc6
0x1237	0xabc7
0x1238	0xabc8
0x1239	0xabc9
0x123a	0xabca
0x123b	0xabc b
0x123c	0xabcc
0x123d	0xabcd
0x123e	0xabce
0x123f	0xabc f
0x1240	0xdef0
0x1241	0xdef1
0x1242	0xdef2
▪	
▪	

H
e
l
l
o

What is a Page?

- Fixed-size, aligned section of virtual memory
- All bytes in the same virtual page **MUST** be in the same physical page
- Always power of 2 in size
 - Power of 2 = simple hardware = fast

How Big is a Page?

- 4K is common (Intel, PowerPC)
- 8K is rare (SPARC)
- “Huge Pages” sometimes also supported
 - Larger pages = smaller page table
 - Can be mixed with ordinary 4K pages

The Page Table

0x000	
⋮	
0x122	
0x123	0xabc
0x124	
⋮	
0xfff	

A Page Table Entry

1	0xabc	rw
----------	--------------	-----------

Valid Bit

Physical Page #

Permissions

Page Fault

- Program accesses invalid virtual memory
 - Page table entry not valid, or permissions don't allow operation
- CPU interrupts program, forces into OS
- OS decides what to do

Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

0

Page Table

virt	phys
0	
1	

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	(free)
5	(free)
6	(free)
7	17

Page Fault: Demand Paging

Process A

```
load r1, mem(0)
inc  r1
store r1, mem(1)
```

r1

0

Page Table

virt	phys
0	
1	

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	(free)
5	(free)
6	(free)
7	17

**No valid page
table entries (yet)**

Page Fault: Demand Paging

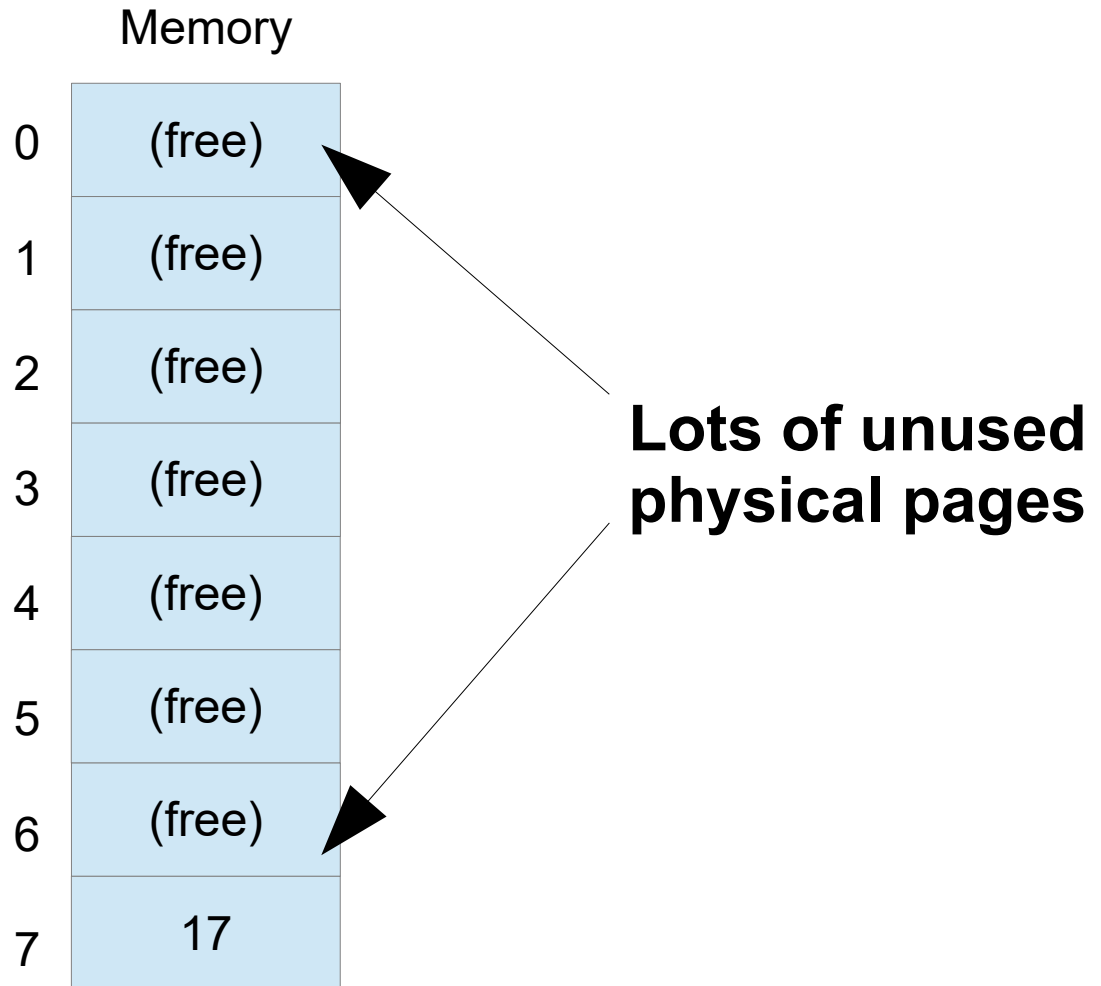
Process A

```
load r1, mem(0)
inc  r1
store r1, mem(1)
```

r1 0

Page Table

virt	phys
0	
1	



Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

0

Page Table

virt

phys

0

1

**Page
Fault!**

Memory

0

(free)

1

(free)

2

(free)

3

(free)

4

(free)

5

(free)

6

(free)

7

17

Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

0

Page Table

virt	phys
0	7
1	

**Add
Entry**



Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	(free)
5	(free)
6	(free)
7	17

Page Fault: Demand Paging

Retry

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

0

Page Table

virt

phys

0

7

1

Memory

0

(free)

1

(free)

2

(free)

3

(free)

4

(free)

5

(free)

6

(free)

7

17

Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)



Success!

Page Table

virt	phys
0	7
1	

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	(free)
5	(free)
6	(free)
7	17

Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

18

Page Table

virt

phys

0

7

1

**Page
Fault!**

Memory

0

(free)

1

(free)

2

(free)

3

(free)

4

(free)

5

(free)

6

(free)

7

17

Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

18


Page Table

virt	phys
0	7
1	4

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	0
5	(free)
6	(free)
7	17

**Allocate
memory**



Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

18

Page Table

virt

phys

0

7

1

4

Retry

Memory

0

(free)

1

(free)

2

(free)

3

(free)

4

0

5

(free)

6

(free)

7

17

Page Fault: Demand Paging

Process A

load r1, mem(0)

inc r1

store r1, mem(1)

r1

18

Page Table

virt	phys
0	7
1	4

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	(free)
6	(free)
7	17

Success!

Demand Paging: Summary

- Valid virtual address, but page table entry is not filled in
- Page fault on first access to the page
- OS verifies that address is valid, then fills in page table entry
 - May need to allocate & init a page
- Retry operation

Virtual Address Space

Virtual Memory

Process A

load r1, mem(0)
inc r1
store r1, mem(1)

r1

18

0

1

17

18

17

289

Process B

load r1, mem(0)
mul r1, r1, r1
store r1, mem(1)

r1

289

Page Fault: Copy-On-Write (COW)

Process A
...code...

Page Table

virt	phys
0	7 (ro)
1	4

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	(free)
6	289
7	17

Process B
set r1, 19
store r1, mem(0)

r1	289
----	-----

Page Table

virt	phys
0	7 (ro)
1	6

**Shared physical page.
Both marked read-only.**

Page Fault: Copy-On-Write (COW)

Process A
...code...

**Going to
overwrite
page 0**

Page Table

virt	phys
0	7 (ro)
1	4

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	(free)
6	289
7	17

Process B
set r1, 19
store r1, mem(0)

r1	289
----	-----

Page Table

virt	phys
0	7 (ro)
1	6

Page Fault: Copy-On-Write (COW)

Process A
...code...

<u>Page Table</u>	
virt	phys
0	7 (ro)
1	4

Memory	
0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	(free)
6	289
7	17

Process B
set r1, 19

store r1, mem(0)

r1	19
----	----

Page Fault!

<u>Page Table</u>	
virt	phys
0	7 (ro)
1	6

Page Fault: Copy-On-Write (COW)

Process A
...code...

**Allocate page,
copy data**

Page Table

virt	phys
0	7 (ro)
1	4

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	17
6	289
7	17

Process B
set r1, 19
store r1, mem(0)

r1	19
----	----

Page Table

virt	phys
0	7 (ro)
1	6

Page Fault: Copy-On-Write (COW)

Process A
...code...

**Update
page table**

<u>Page Table</u>	
virt	phys
0	7 (ro)
1	4

Memory

0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	17
6	289
7	17

Process B
set r1, 19
store r1, mem(0)

r1	19
----	----

<u>Page Table</u>	
virt	phys
0	5
1	6

Page Fault: Copy-On-Write (COW)

Process A
...code...

<u>Page Table</u>	
virt	phys
0	7 (ro)
1	4

Memory	
0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	17
6	289
7	17


Process B
set r1, 19

store r1, mem(0)

r1	19
----	----

<u>Page Table</u>	
virt	phys
0	5
1	6

Retry



Page Fault: Copy-On-Write (COW)

Process A
...code...

<u>Page Table</u>	
virt	phys
0	7 (ro)
1	4

Success!

Memory	
0	(free)
1	(free)
2	(free)
3	(free)
4	18
5	19
6	289
7	17

Process B
set r1, 19
store r1, mem(0)

r1	19
----	----

<u>Page Table</u>	
virt	phys
0	5
1	6

COW: Summary

- Single physical page
- Lots of page table entries point to it
 - All read-only
- Page fault on first write
 - Allocate new page
 - Copy data
 - Update page table
- Retry operation

Other things to think about:

- How to implement `fork()` ?
- How to implement a swap file?
- How to implement memory-mapped files?
- How to implement shared memory?
- How to implement pages filled with zeroes?

Other things to think about:

- Policy questions:
 - How much physical memory should a process be allowed to consume?
 - How much virtual memory should a process be allowed to allocate?

Other things to think about:

- How do we reduce the size of the page table?
- How do we cache page table entries?
 - And when does the cache need to be flushed?
- When, if ever, can two processes share the same page table?
 - And how do you handle cache flushing if they do?