

CSC 525:
Principles of Computer Networks

On Distributed Communications Networks


- by Paul Baran at RAND in early 60's
- one of the first proposals for packet switching and dynamic routing
- To understand the fundamental reasons for packet switching and dynamic routing

Side Topic: Reading Papers

- First Pass:
 - What is the goal? What is the problem to solve?
 - What is the problem context?
 - Background, previous work, assumptions
 - What is the main idea of the work?
 - To accomplish this:
 - Carefully read abstract, intro, main design, and conclusions
 - Skim analysis/evaluation section
- Second Pass
 - Start with the assumption that the main point is wrong
 - (even if you agree with it :)
 - Focus on analysis/evaluation sections and the details of the design.
 - How do the analysis and results prove otherwise?
 - What are the weak points?

Telephone Networks

- Circuit switched, time division or frequency division.
- Centralized switching control
- System reliability relies on components' reliability
 - individual nodes are not expected to fail
- Dumb terminals, complex network, only app is voice.



IP Networks

- Packet switched, statistical multiplexing.
- Dynamic Routing
 - adapt to component failures and other changes in networks.
- Complex hosts, dumb network, built on redundant unreliable components to achieve system reliability.
- Supporting a wide range of applications
- All envisioned in Baran's paper

Building **Robust** Communication Systems

- To achieve high system availability in the presence of component failures
 - Redundancy in components (links and nodes)
 - Absence of centralized control
 - Cut data into packets, each packet carries its own destination address.
 - Dynamically route packets around failed parts
- Provide a common all-digital service for a wide range of applications.

Types of Networks




Fig. 1—(a) Centralized. (b) Decentralized. (c) Distributed networks.

Common Telco Designs

Proposed Mesh

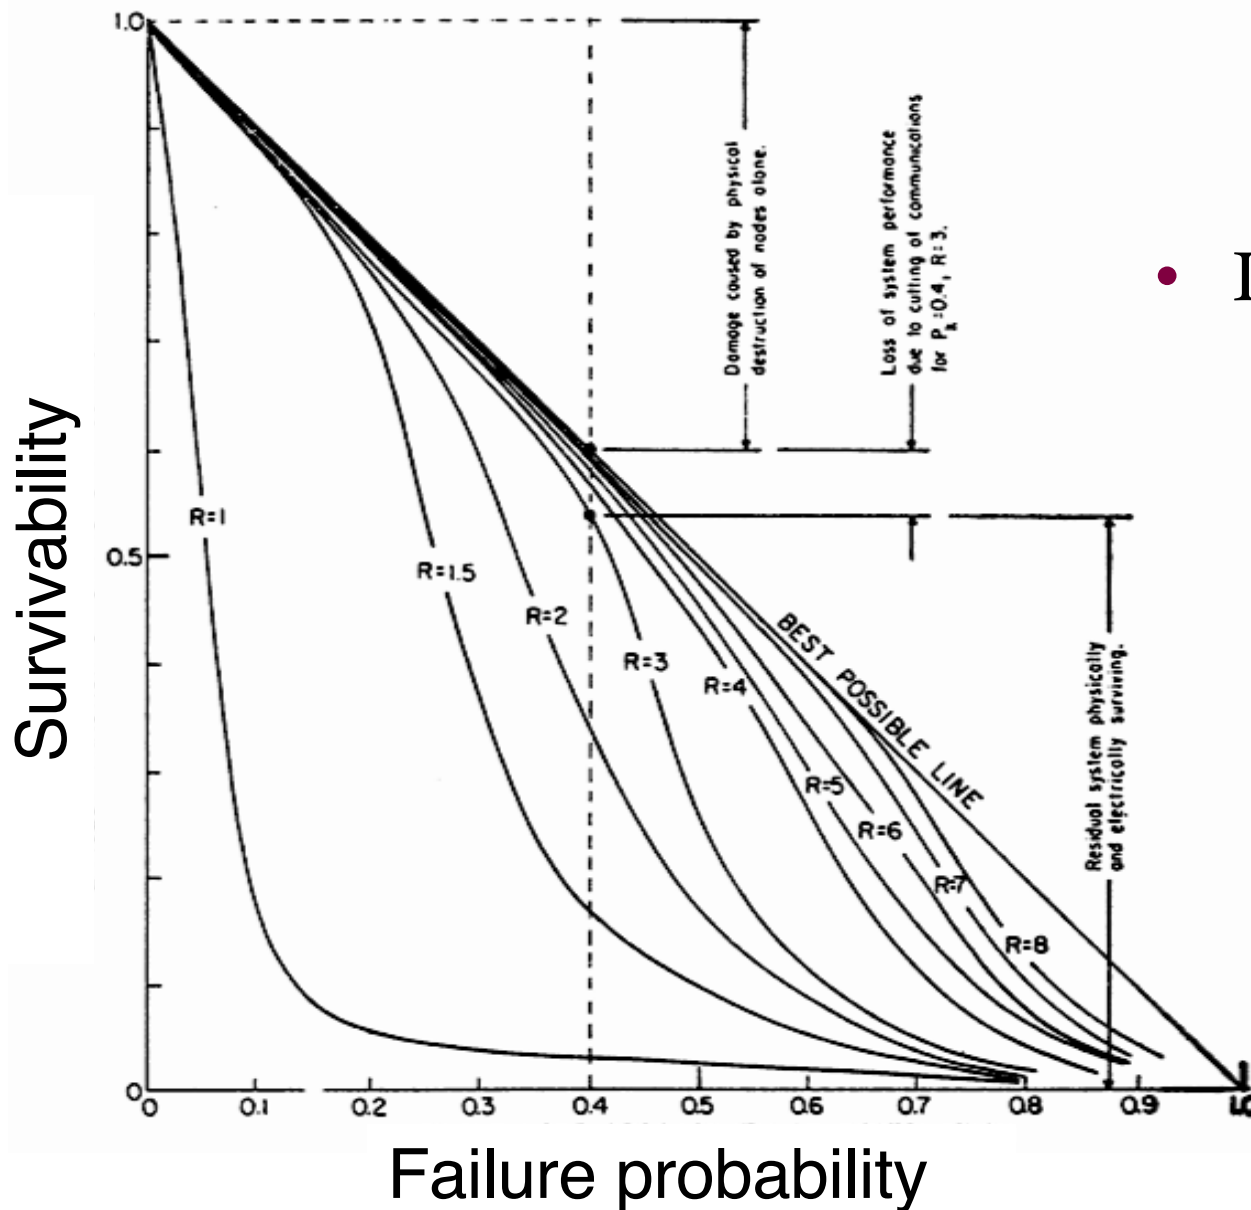
Network Redundancy

- Define redundancy level = $\#links/\#nodes$
- E.g., centralized star topology has redundancy level of 1.
 - A star topology has n nodes and $n-1$ edges.
- Best case is a full mesh (clique)
 - Every node connects to every other node.
 - But clique is expensive in number of links
 - Does not *scale*

Side Topic: Scalability

- Being scalable is one of the most fundamental challenges in Internet protocol design.
 - How do you build a system that can grow to an arbitrary large size and still maintain its reliability and efficiency?
 - Many problems are easy in small scale, but hard in large scale.
- See how the design in each paper handles scalability issue.
- Another fundamental challenge is to handle dynamics/changes, especially unexpected ones.

Survivability and Redundancy

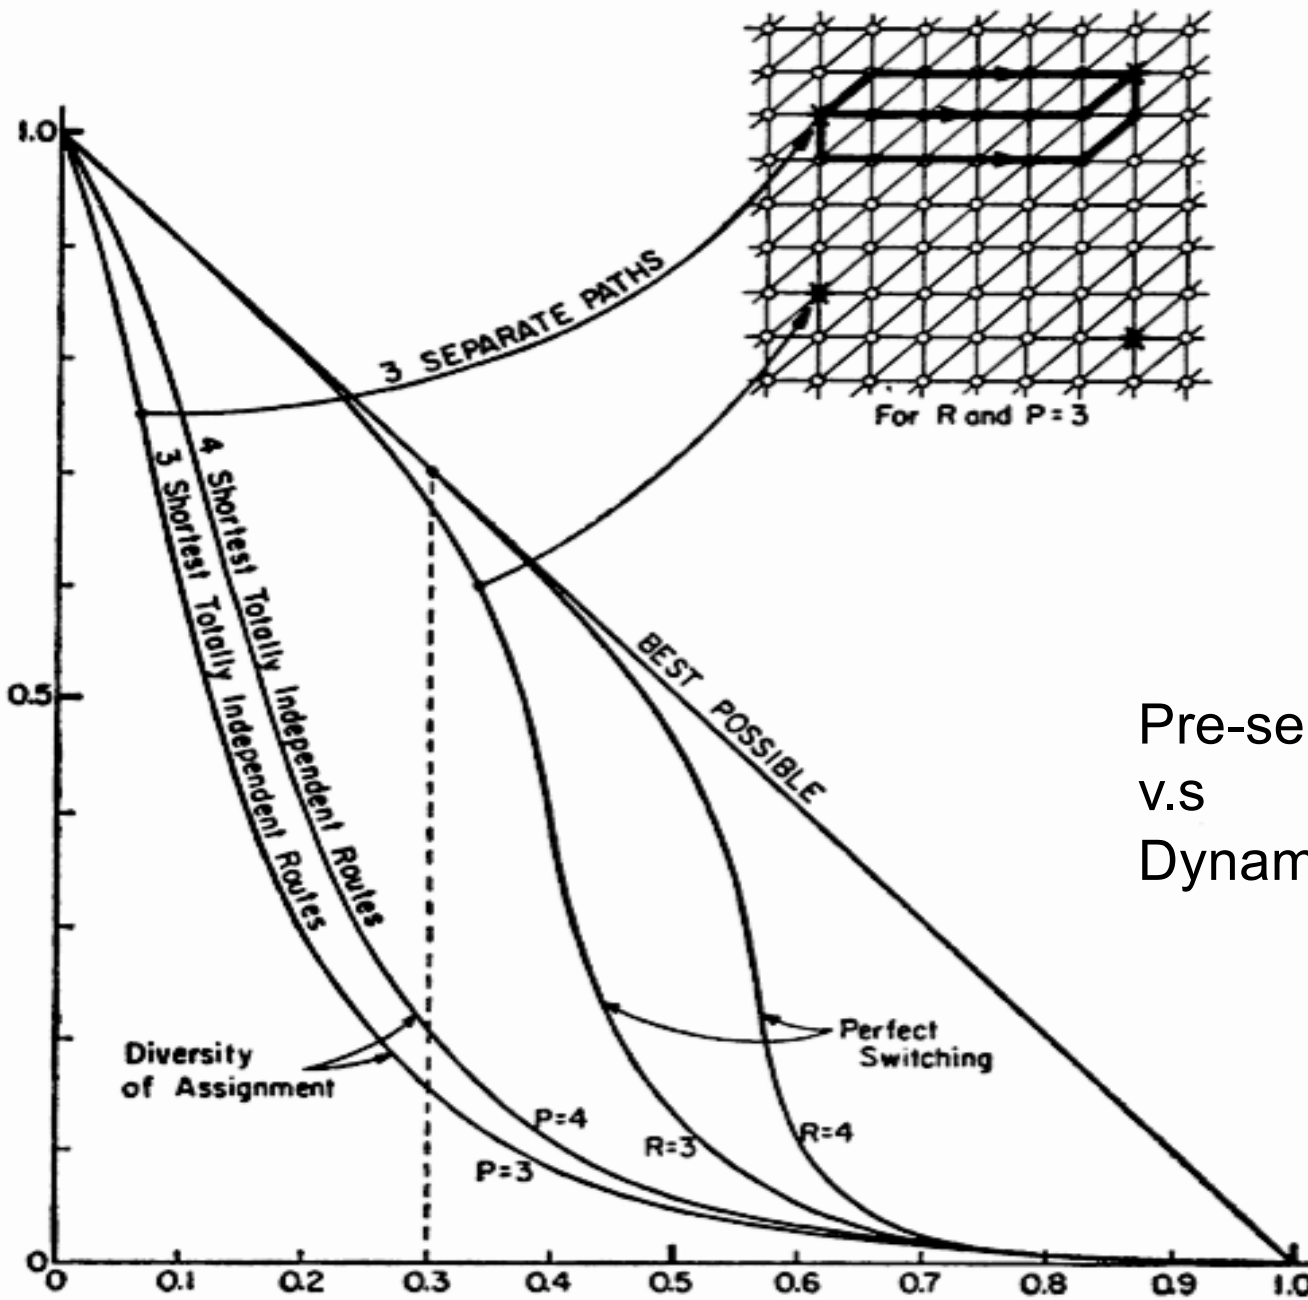


- Do we need a full mesh?
 - If yes, solution is not realistic.
 - Result is that redundancy of 3-4 may suffice.
 - Further increase of redundancy doesn't help very much.

How to Select Routes

- Goal: To take advantage of link redundancy in case of failures.
- Options:
 - Pre-select a single route, clearly bad.
 - Pre-select several routes (diversity of assignment)
 - So no need for runtime path computation
 - Dynamic routing (perfect switching)
 - Will need to find alternative path upon failures

Survivability




Pre-select Routes
v.s
Dynamic Routing

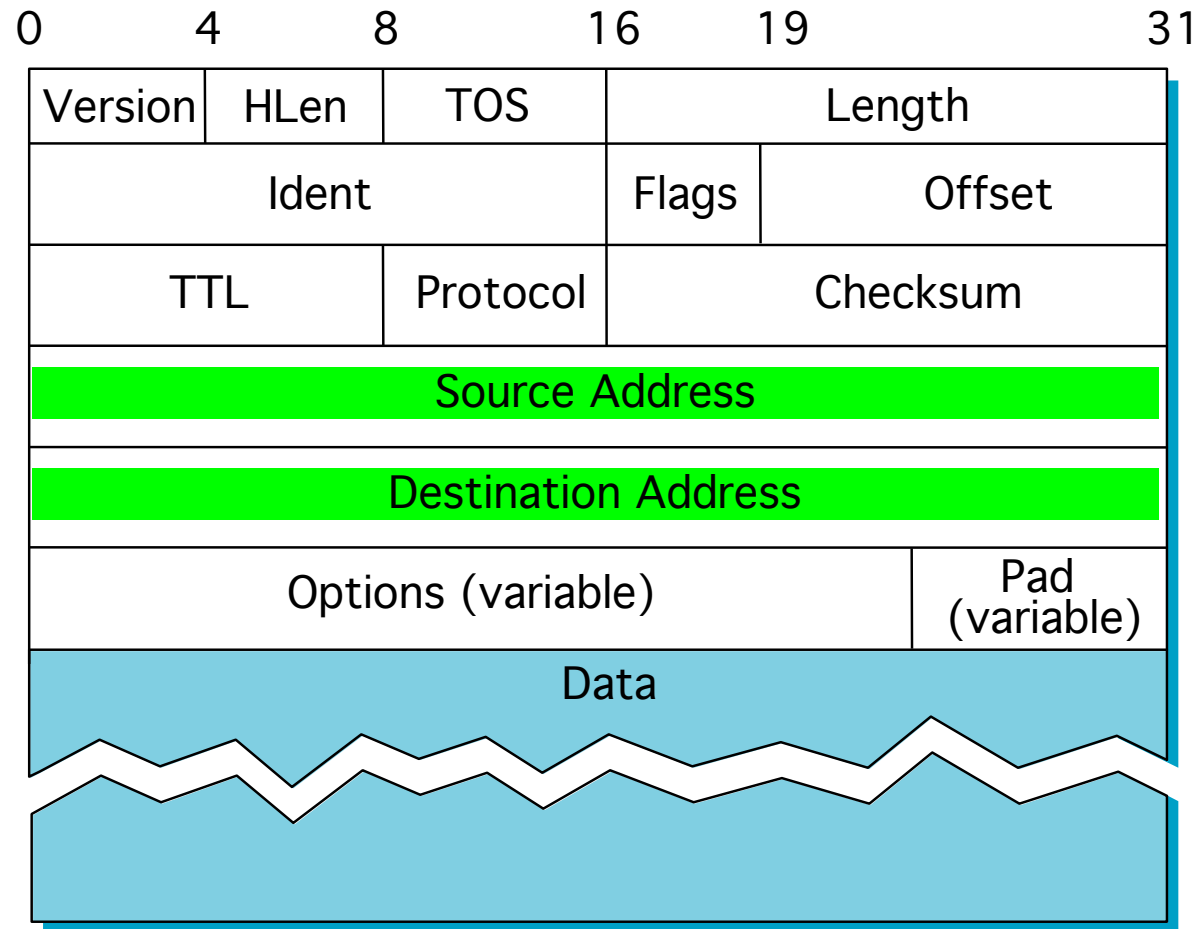
Failure probability

Packets

- A standard packet header format
 - source & dest. addresses, hop-count
- Store-and-Forward
 - concept of statistical multiplexing
 - different users/applications have different throughput requirements
 - send whenever packets are ready, no connection setup overhead.

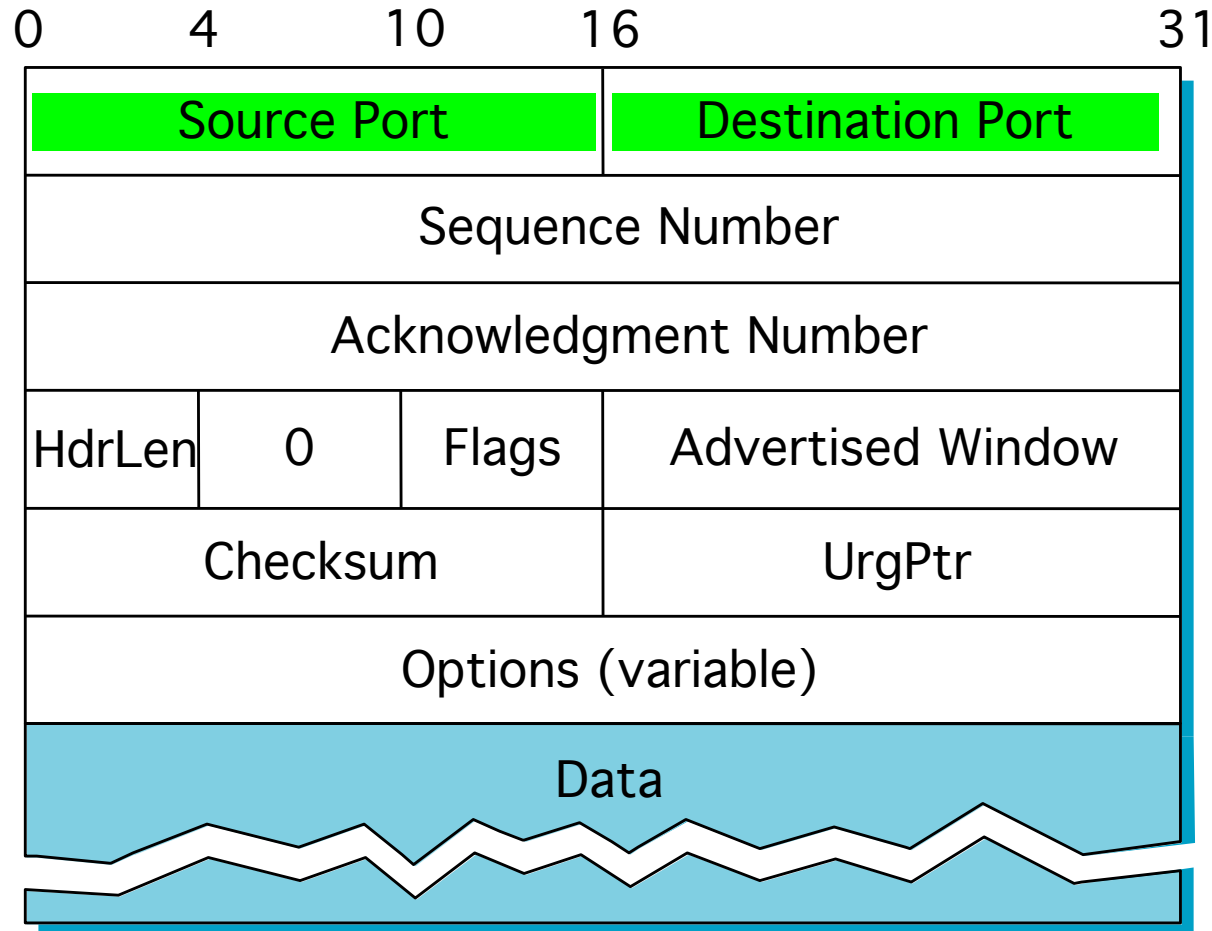


IP Header



Best effort packet delivery

TCP Header



End-to-end reliable byte stream

Dynamic Routing

- Interesting concept
 - Built a richly connected distributed network
 - Built packets with source and destination addresses
 - Claim we need dynamic routes
 - in order to get around failed links
- But how would you do it?
 - What about time, memory, storage, buffer, etc?
 - Today we have a number of routing protocols (RIP, OSPF, BGP etc.)

Hot Potato Routing

- Two equal objectives:
 - Send along shortest working path to destination.
 - Forward the packet as soon as you get it.
- Conflict: many packets may want to go out on the same interface for shortest path.
- **Trade-off:** send along the shortest *available* path

Routing Table and Forwarding Table

	LINK NUMBER							
	1	2	3	4	5	6	7	8
	HANDOVER NUMBER ENTRIES							
A	22	∞	12	10	9	9	8	13
B	5	3	2	2	4	5	12	2
C	7	8	13	9	22	10	7	8
D	21	23	19	21	12	10	12	13
E	7	10	12	14	12	13	13	15
F	7	10	12	13	14	21		
G	6	4	10					

Z	15	20	7	3	10	8	5	10
---	----	----	---	---	----	---	---	----

	BEST CHOICE				
	1st	2nd	3rd	4th	5th
	LINK NUMBER for DECISION CHOICE				
	7	5	6	4	3
	3	4	8	2	5
	1	7	2	8	6
	6	5	7	8	3
	1	2	3	5	6
	1	2	3	4	5
	5	2	1	6	3

4	7	3	6	5
---	---	---	---	---

Routing Table:
For each destination, maintain outgoing interfaces and distances

Forwarding Table:
For each destination, maintain the best outgoing interface.

Routing Lookup

Building the Forwarding Table

- Listen for incoming traffic, use hop count to estimate distance **back to the source**
 - Hop count is recorded in the packet header.
- When a router receives a packet
 - if the router has not seen the source address before, add it to the routing table
 - if the source already in routing table, and this packet has smaller hop count than the one in the table, update the table.
 - Otherwise do nothing.
 - Use info carried in data packets to build the routing table.
- Today's routing protocols use special packets, i.e., routing messages, to communicate and build the routing table.
 - Separate control and data.
 - But Ethernet uses a learning algorithm similar to this paper's.

Be Adaptive

- Need to be adaptive to changes
 - Component failures, malicious attacks, etc.
- Use exponential average to take present and past into consideration
 - $H(\text{next}) = a * H(\text{current}) + (1-a) * H(\text{measurement})$
 - Tune it by adjusting the value of parameter a .
- Assign link weight to reflect the cost or distance of the link.
- These techniques are still being used today.

How to build a robust system

- Learn the changing environment by observations
- Automatically adjust to changes
 - Treat changes as norms, not exceptions
- With adequate redundancy and adaptivity, one can build a highly robust system out of unreliable parts
 - Vulnerable systems fail if any component fails
 - Robust systems fail only if all the parts have failed
 - A wide range in between
- One of the fundamental design principles in the Internet architecture

Summary

- The paper provides some context and history for network systems.
 - Suggests a move to dynamic packet switching networks rather than centralized circuits.
- Introduces some basic building blocks for Internet design
 - Distributed networks without centralized control
 - Standard packet header
 - Dynamic routing algorithm
- Keep these in mind in reading more recent papers and think how these are reflected in actual protocol design, and whether they fit well or not with emerging needs.