

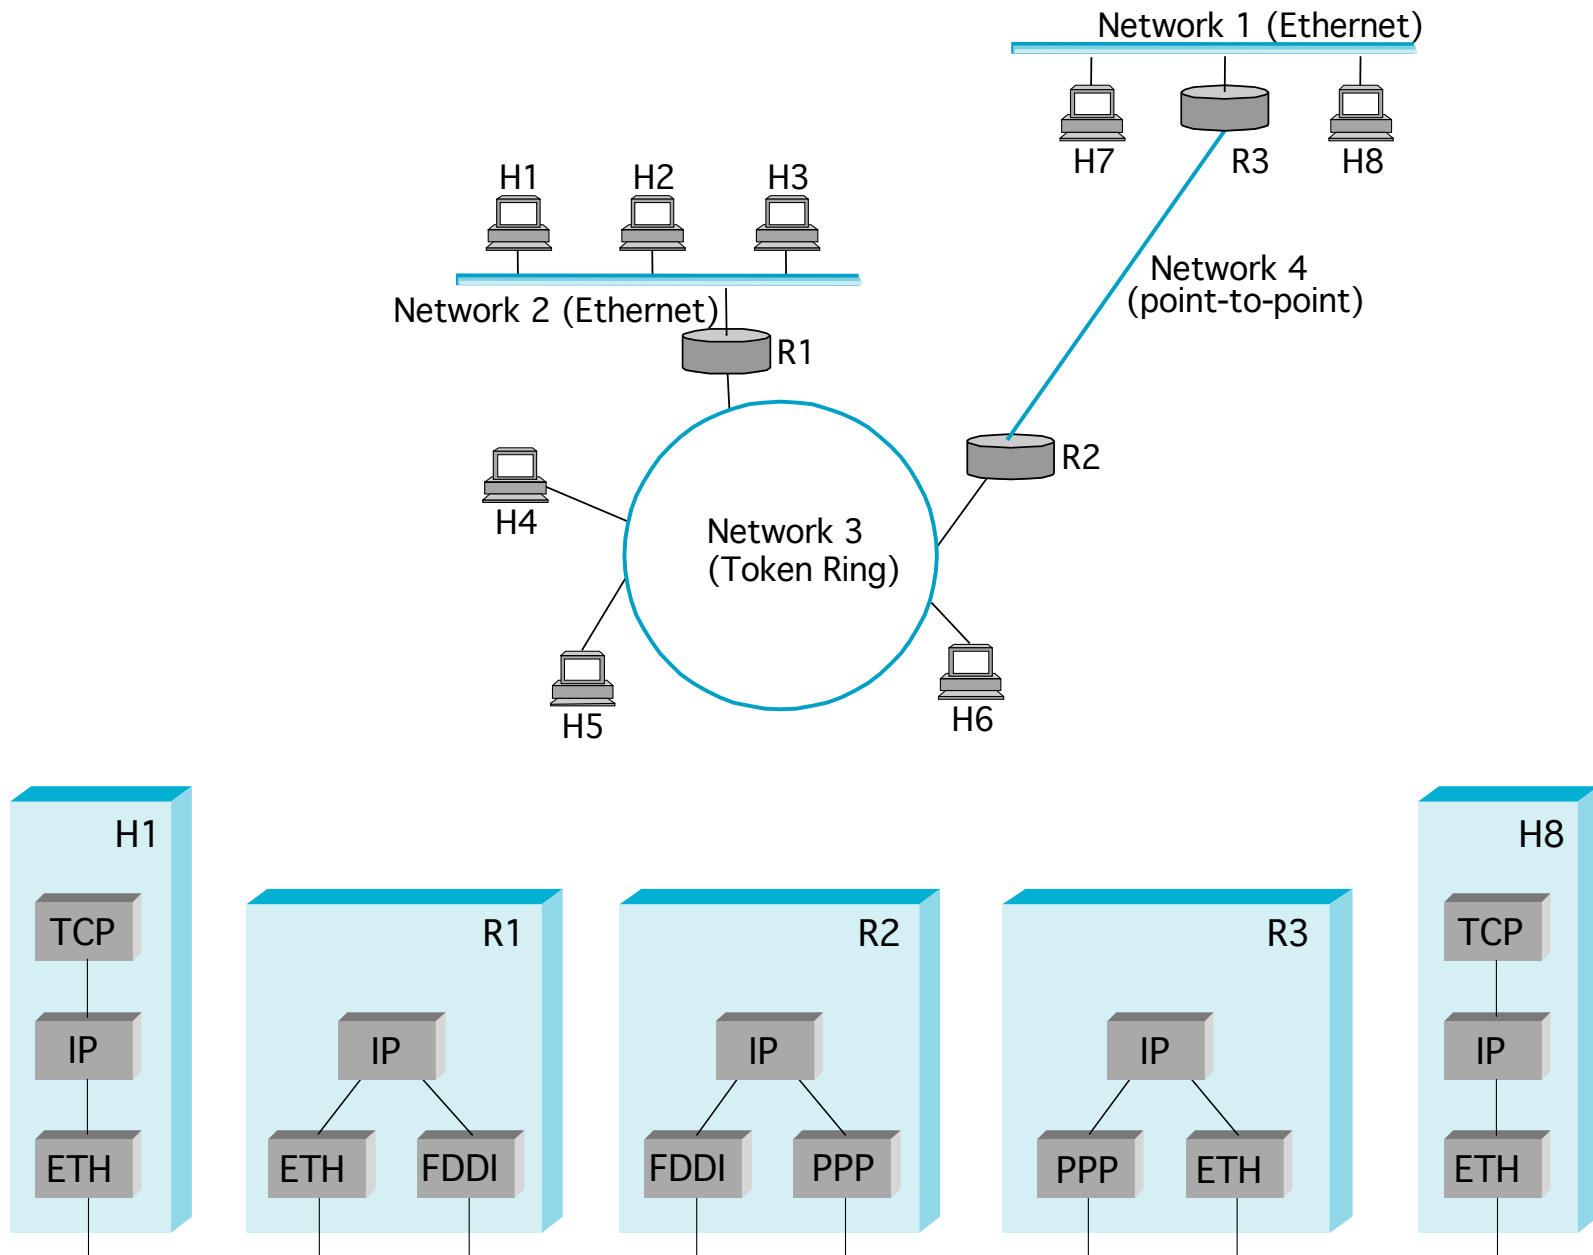
CSC 525: Computer Networks

“Design Principles of the DARPA Internet Protocols”

- Why IP
 - A number of different network technologies developed in early 70's
 - Different media: copper, radio, satellite
 - Different protocol designs
 - Under different administrative control
 - The need for inter-connecting different nets.

Fundamental Goal

- To develop an effective technique for multiplexed utilization of all existing networks
 - no centralized control
 - no changes to individual networks
- Result
 - IP-based packet switched network consisting of distinct networks with store-and-forward gateways (routers) between them.
 - TCP/IP



Second Level Design Goals

In the order of importance

- continued operation despite partial failures
- support multiple types of communication services
- accommodate a variety of networks
- distributed management of resources
- cost effective, performance
- low cost attachment
- resource accountability

What else would you add?

Trade-offs

- But not all the goals can be satisfied equally
 - Which goals win and which lose in the trade-off?
- The **order of the goals is essential**
 - Very strong focus on the first three
 - Survive network and router failures
 - Provide different types of services
 - Accommodate a variety of networks
 - A different order would produce a different design.
 - E.g., accounting barely works in the Internet.

Survivability

- Links and Routers may fail and stop working
 - (side note: design did not anticipate misbehavior)
- Two ends can continue communication
 - Despite faults at any intermediate point
 - Mask any transient failures (e.g. route changes)
- A network with only one failure mode:
complete partition

Achieving Survivability

- Implications of storing network states
 - Any state stored at intermediate nodes must be replicated (since node may fail)
 - Difficult (at best) to design/implement.
- Solution: States are only stored at edges (hosts)
 - Stateless packet switches/routers (middle of network)
 - Consider TCP states: seq#, ack#, window size, etc are all stored at the hosts. No TCP information is stored in the intermediate routers.
 - Fate-Sharing: the fate of the communication is only shared with the fate of the two ends.

Soft states vs. Hard states.

- What if we do have to maintain some states in the network?
 - E.g., to provide Quality of Service guarantees.
- Soft states: periodic refresh regardless of whether there has been any change.
- Hard states: only update when there's a change.
- Another alternative: carry states in packets.
- Pros and cons?

Heterogeneity: Above IP

- To provide different types of services
 - TCP: reliable delivery
 - Debugging: low complexity, no reliability
 - Voice: delay bound, reliability sometimes hurts
- Leave the network layer simple
 - Split TCP (transport) and IP (network)
 - Connectionless datagram as basic building block
 - No assumption on subnet QoS capability
 - Build multiple services (transport protocols) at the hosts

Heterogeneity: Below IP

- To accommodate a variety of networks
- Least assumption on what subnets can do
 - Transport a packet
 - Reasonable packet size
 - Reasonable (not perfect) reliability
- Not assumed
 - Reliable or sequential delivery
 - Network broadcast or multicast
 - Priority or services
 - Failures, speeds, or delays
- Implement functionalities at the hosts

Other Goals

- Distributed Management
 - Success in allowing multiple domains and diverse routing policies.
 - Not very successful as the networks become large and complex.
- Performance
 - Header size, end-to-end retransmissions, routing table lookup.
- Adding Hosts and Nets
 - Speaking IP
 - Host software is complex
 - Relies on correct implementations/behaviors at the host
- Accounting
 - Challenging in the datagram model

Summary

- Identified and **Prioritized** Goals
 - Top three goals very successful
 - Bottom goals less successful
- Building Block: Datagram
 - Very effective for top goals (survivability, type of services, variety of networks)
 - Suggests that “flows” might be better for different priorities of goals
 - Suggests periodic messages: **Soft-State**


“End-to-End Arguments in System Design”

- Fact: layered network architecture
 - why layers: divide and conquer
- Downside of layering: potential duplication of functionalities across layers
 - example: how many layers are concerned with data transmission reliability?
- What functionality to put into which layer?

Basic Observation

- Some applications have end-to-end performance requirements
 - Reliability, delay bound, security etc.
- Implementing these in the network is very hard
 - Every step along the way must be fail-proof
 - Sometimes network doesn't have enough information
- The hosts
 - Can satisfy the requirements without the network
 - Cannot depend on the network

Example: Reliable File Transfer



- Errors can occur at many levels
 - from sender disk to memory
 - from sender OS to sender line card
 - across network
 - from receiver line card to receiver OS
 - from receiver memory to disk

Reliable File Transfer (II)

- Make each step reliable and concatenate them
 - What happens if any router fails or misbehaves?
 - The receiver has to do the check anyway!
- End-to-end check and retry
 - Full functionality can be implemented entirely by the applications with no need for reliability from lower layers.

Conclusion

- Implementing perfectly reliable network
 - Doesn't reduce host implementation complexity
 - Does increase network complexity
 - Probably imposes overhead on all applications, even if they don't need it, sometimes can hurt other applications (e.g. VoIP)
 - Provides false sense of security
- However, in some cases improving network reliability can enhance performance
 - e.g., over very lossy links

What does it mean

- Think twice before putting functionality into the network
- If hosts can implement it correctly, do it at a lower layer **only** as a necessary performance enhancement.
 - Consider marginal gain at lower layer
 - Consider the impact on other applications and services.

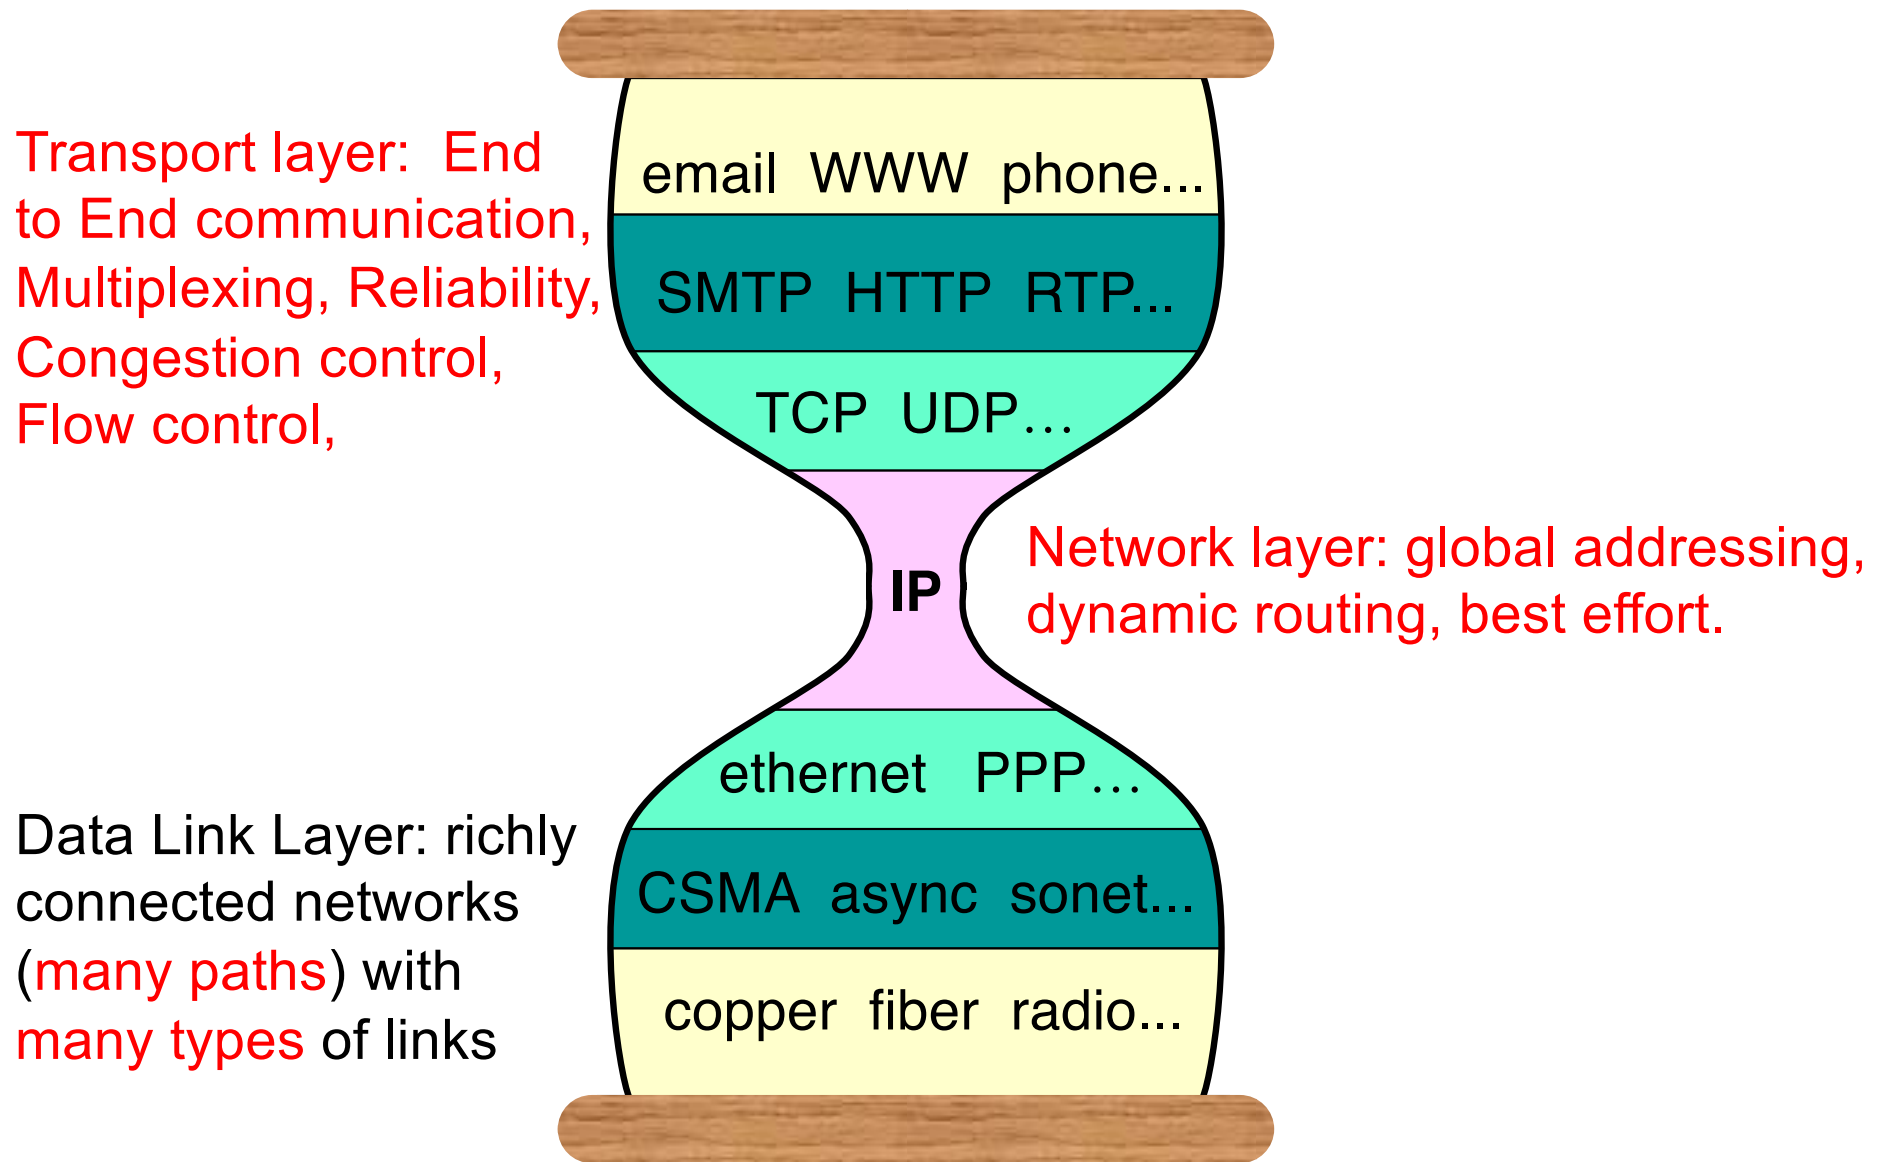
Extended Version

- Don't put application semantics in network
 - Leads to loss of flexibility
 - Cannot change old applications easily
 - Cannot introduce new applications easily
- Short-term application performance vs. long-term architectural flexibility
 - More functionality imposes more restrictions, but benefit known apps more.
 - Sometimes it's a tough call

The Erosion of the Architecture

- Layering and End-to-End principle are regularly violated by middle boxes in today's Internet:
 - Firewalls, NATs, ALG (application layer gateways)...
- Battle between architectural purity, commercial interests, and political pressures
 - ISPs
 - Vendors
 - Governments

The Result from the Design Principles



So much for the principles

We reject kings, presidents and voting.

We believe in rough consensus and running code.

- Dave Clark