


# CSC 525: Computer Networks

# Where we are



# Duties of Transport Layer

- Demultiplex:
  - IP gets packets to a host, i.e., dest address.
    - Its header points to a transport protocol
  - Transport gets packets to the app, i.e., dest port.
    - Identified by a port number on the host.
    - Well-known ports vs. ephemeral ports
- Provide different services for certain apps:
  - Different apps may want different services on top of IP's best-effort delivery
    - UDP: no added feature; just whatever datagrams received.
    - TCP: reliable, in-order, byte-stream, with flow and congestion control.
    - There're many other transport protocols in between.

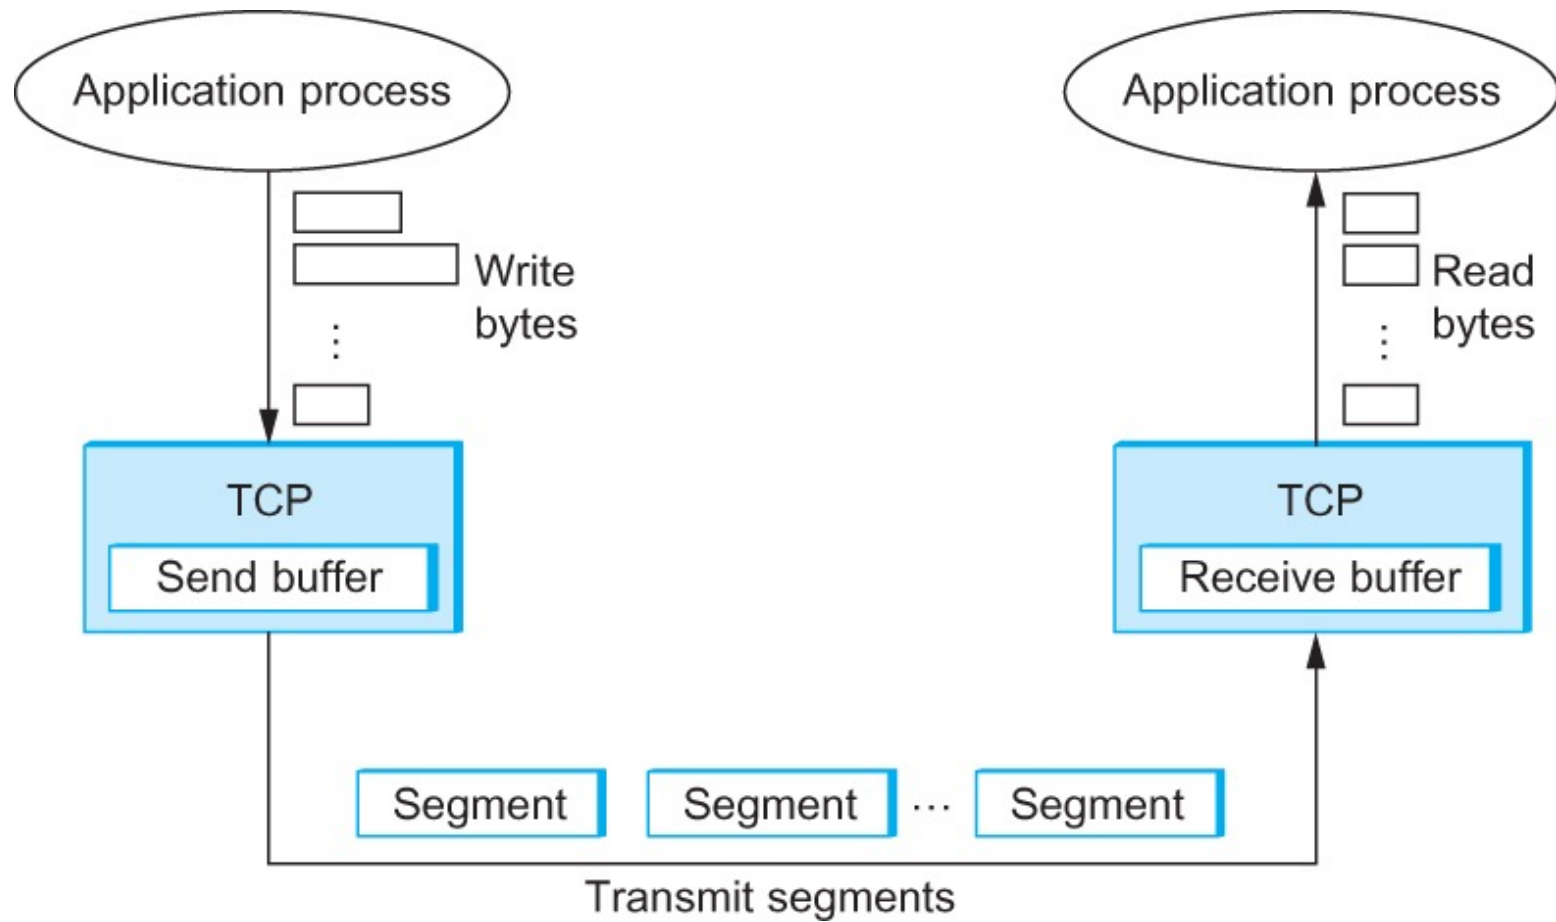
# TCP Overview

- End-to-End unicast: one sender, one receiver
- Connection-oriented:
  - exchange control msgs first to initialize sender & receiver states
- Full duplex data delivery:
  - bi-directional data flow over the same connection
- **Reliable, in-order, byte-stream delivery**
  - no “message boundaries”
  - sender & receiver must buffer data before delivering to apps.
- Flow controlled
  - Prevent sender from overrunning receiver buffer.
- Congestion controlled
  - Prevent sender from overrunning router buffers.

# TCP Segment

- TCP provides a byte-stream service to applications
  - Applications read from and write to a byte-stream.
- But TCP itself does not transmit individual bytes over the network.
  - On the source host TCP buffers enough bytes from the sending process to fill a reasonably sized packet.
  - On the destination host TCP empties the contents of the packet into a receive buffer, and the receiving process reads from this buffer at its leisure.
- The packets exchanged between TCP peers are called *segments*.


# TCP Segment




How TCP manages a byte stream.

# Reliable Transmission


- How to achieve reliable, in-order delivery with controlled rate?
- The basic idea is called **Automatic Repeat reQuest (ARQ)**
  - Sender sets a timer and waits for ACK.
  - Receiver sends ACK.
  - If timed out, sender retransmits.




# Different Scenarios




(a)



(c)




(b)



(d)




# Stop-and-Wait



- At most one outstanding (un-acked) data segment at any time.
- Problem: the pipe is empty most of the time.
  - $\text{Throughput} = \text{one pkt size} / \text{RTT}$

# Sliding Window

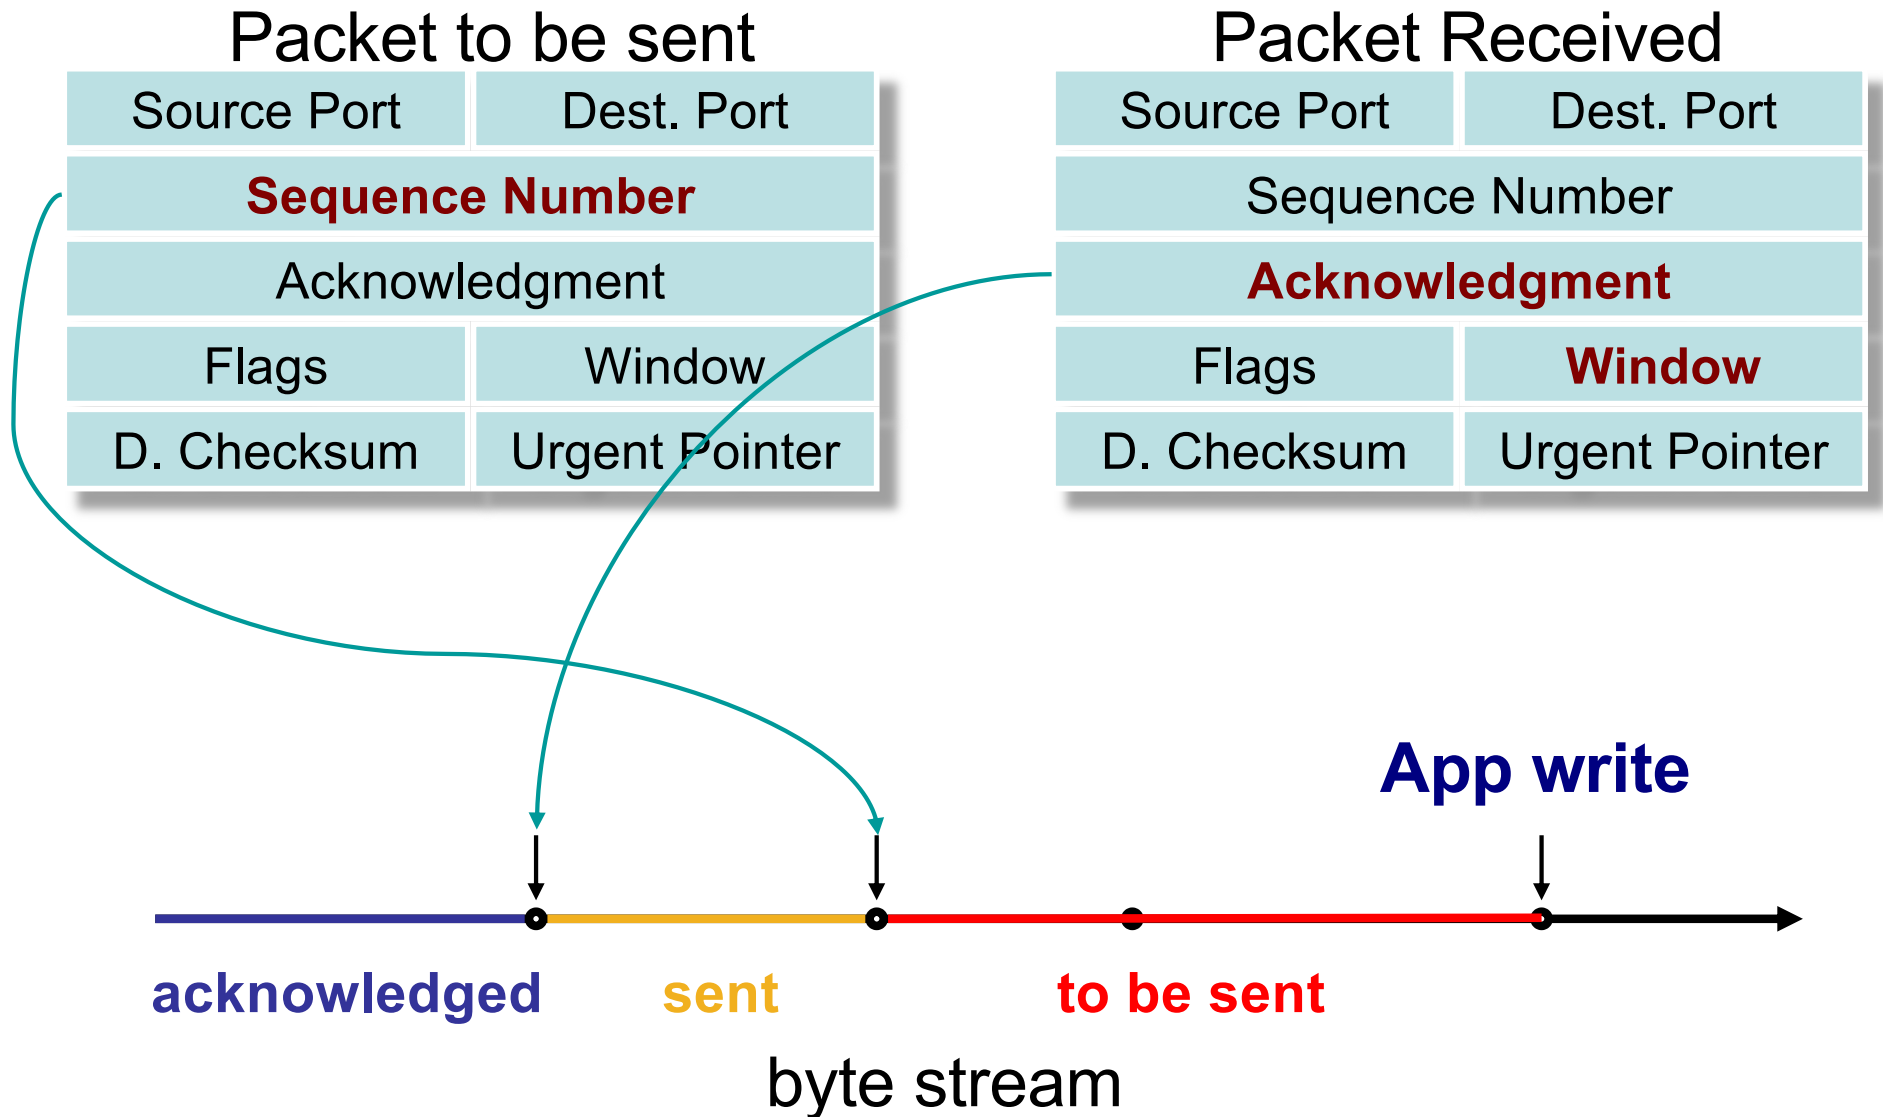
- Pipelining: send multiple outstanding segments (up to a limit), advance the window as the ACKs arrive.
- The max number of un-ACKed segments allowed is called *window size*.



# Acknowledgement

- Receiver sends acknowledgement back to sender.
  - Can be a separate packet, but more often piggybacked on data packet (and mark the A flag bit in TCP header).
- Cumulative ACK: ACK number X means
  - The receiver has got all data preceding X.
  - The next data the receiver is expecting is X.
  - E.g., if received 1, 2, 3, 5, 6, then ACK = 4.
- Other schemes exist, such as selective ACK.

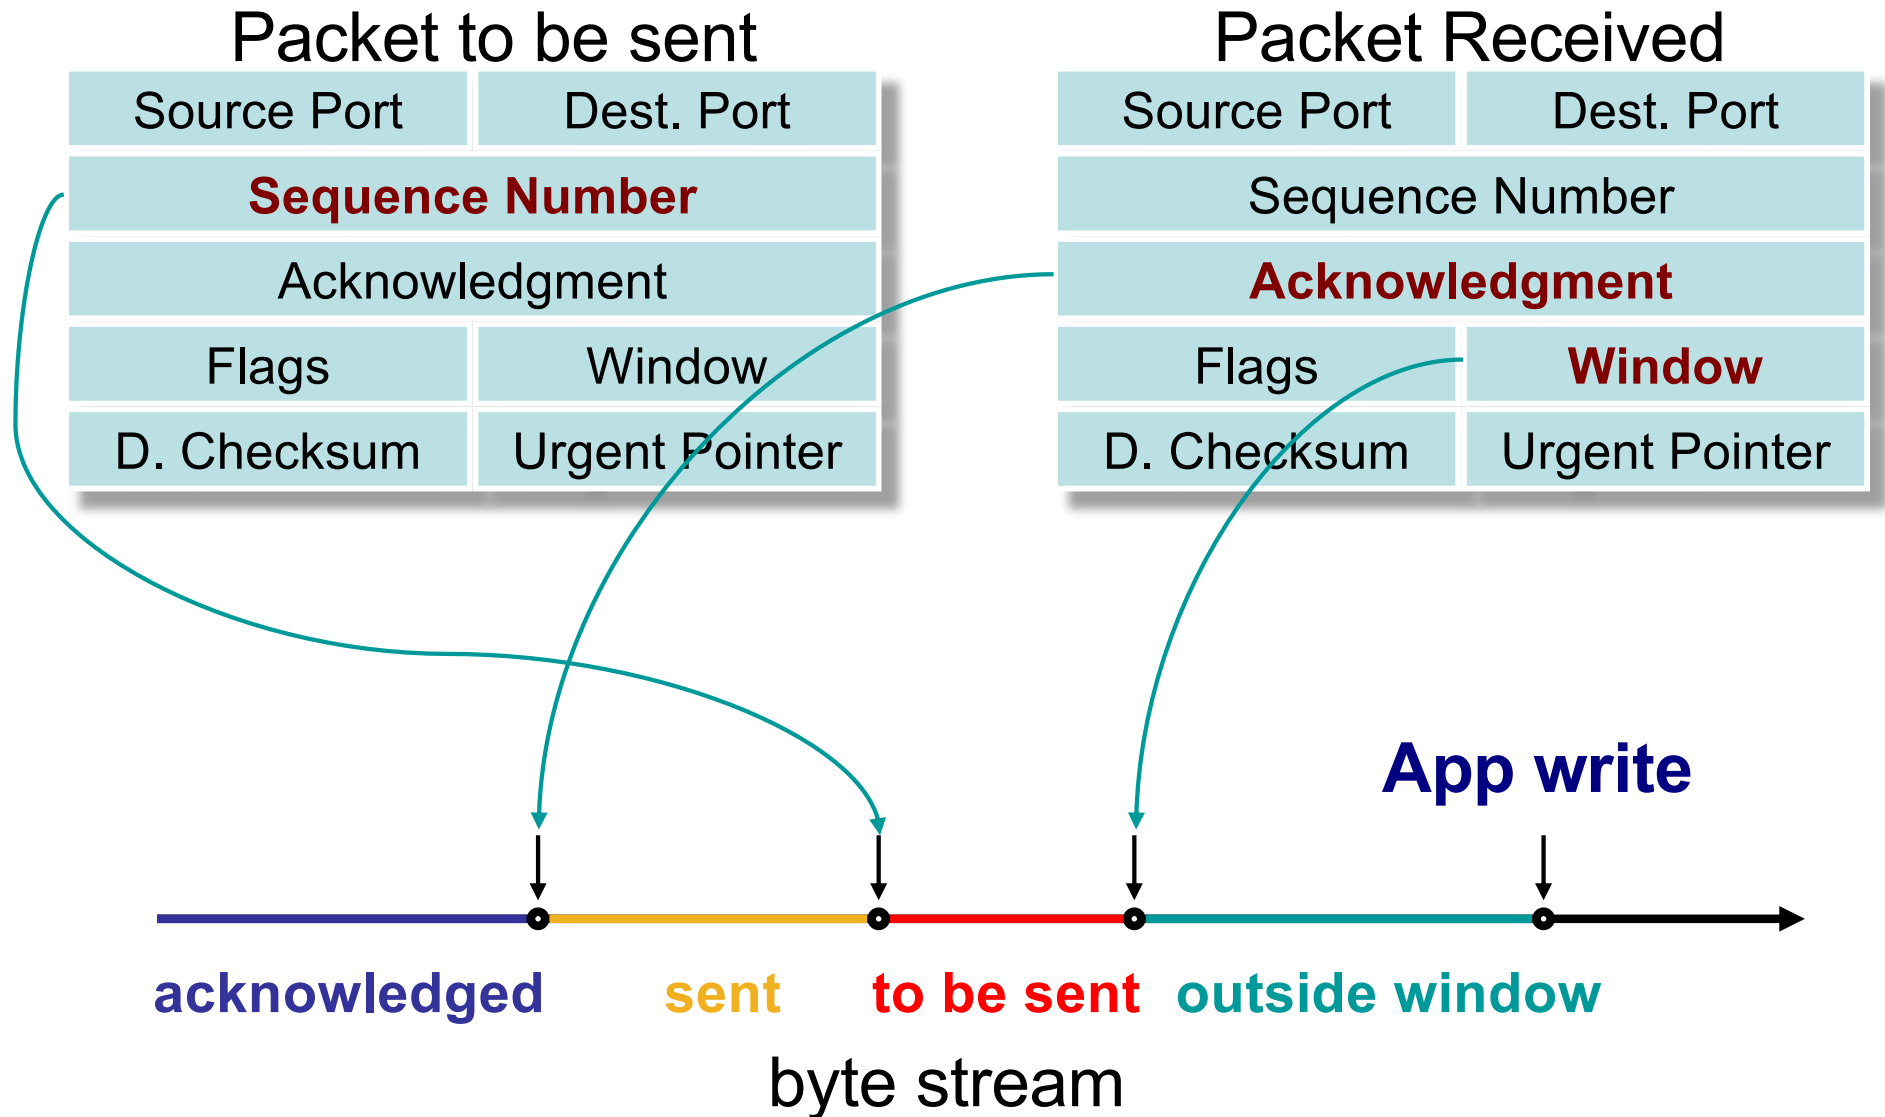
# Seq# and Ack#



# TCP Flow Control

- Make sure receiving end can keep up with incoming data.
- The receiver puts its *available* buffer size in the TCP header as *AdvertisedWindow*.
- The sender should not send more data than what *AdvertisedWindow* allows.

# TCP Flow Control: Sender Side




# TCP Congestion Control

- Can the network keep up with the data rate?
  - Finite link bandwidth
  - Finite buffer size at routers
- End-to-end solution: TCP makes guesses about the available network capacity and adapts. Only end-hosts are involved, not routers.

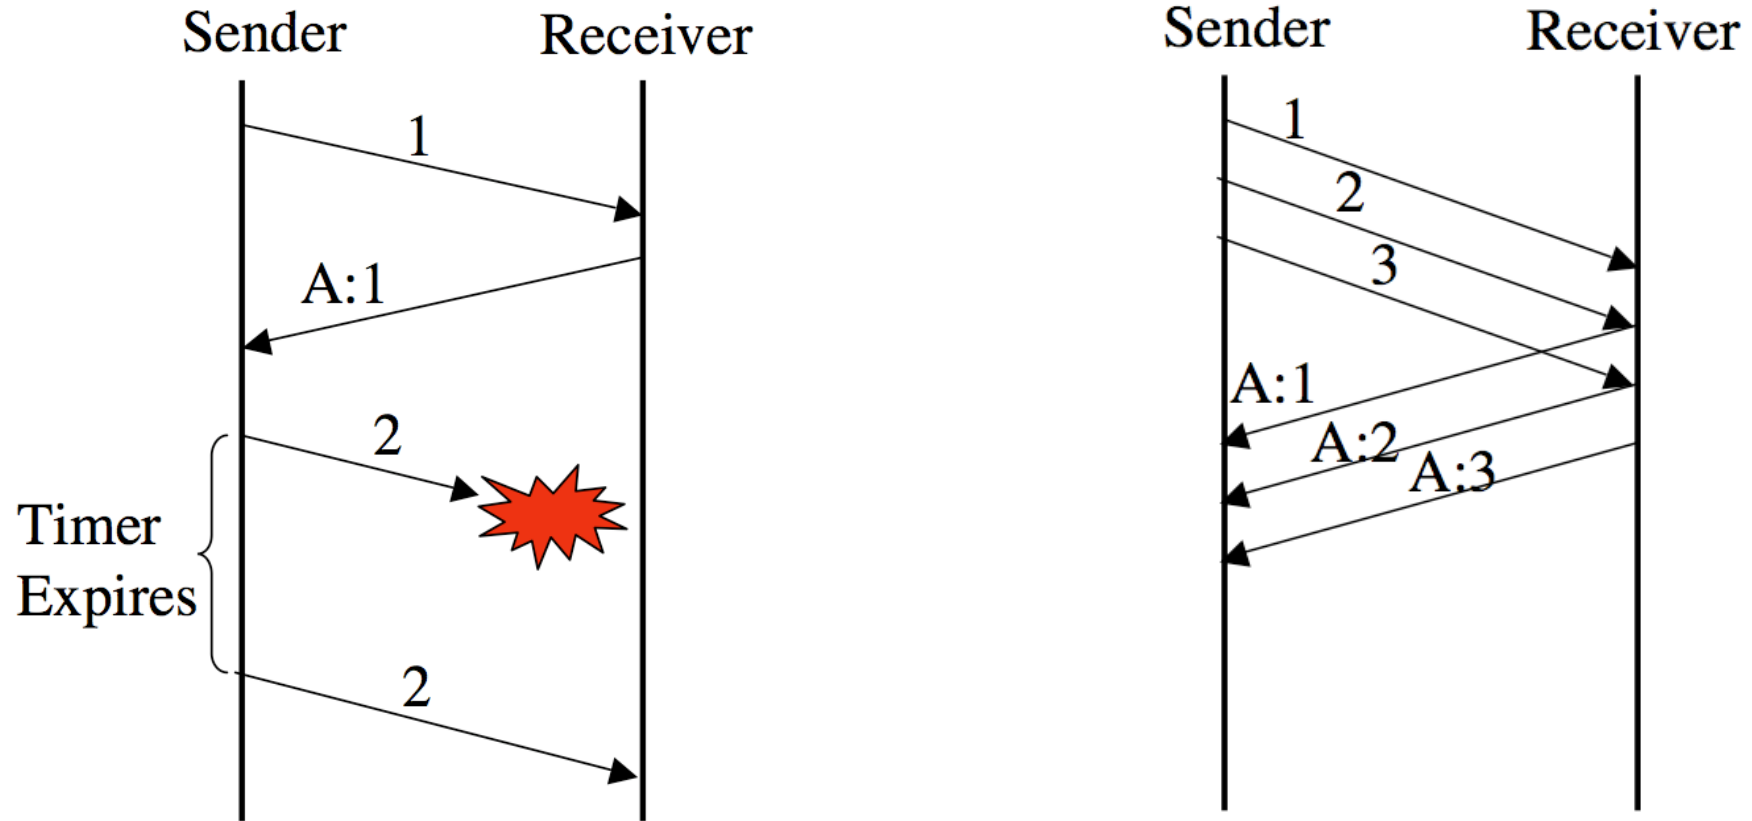
# The danger of congestion

- Knee – point after which
  - Throughput increases very slow
  - Delay increases fast
- Cliff – point after which
  - Throughput starts to decrease very fast to zero (congestion collapse)
  - Delay approaches infinity





# Performance of Sliding Window



- $\text{Throughput} = \text{window size} * \text{pkt size} / \text{RTT}$ 
  - How to set the retransmission timer?
  - How to set window size?

# Setting Retransmission Timer

- Observe and adapt: use exponential moving average based on measurements.

$$\begin{aligned}A(n) &= b \cdot A(n-1) + (1-b)T(n) \\D(n) &= b \cdot D(n-1) + (1-b)(T(n) - A(n)) \\ \text{Timeout}(n) &= A(n) + 4D(n)\end{aligned}$$

*Question: Why not set timeout to average delay?*


$T(n)$  is current measured RTT,  $A(n)$  is estimated average RTT,  $D(n)$  is estimated average deviation.

Karn's algorithm for updating timer in case of retransmission

1. Don't take sample of  $T(n)$  if retransmitted.
2. Double Timeout after timeout ...
3. Reset Timeout for new packet when receiving ACK

# RTT from Experiments

RTT: [gaia.cs.umass.edu](http://gaia.cs.umass.edu) to [fantasia.eurecom.fr](http://fantasia.eurecom.fr)




Blue: sampleRTT, Pink: SRTT

# How to adjust window size?

- Maintain three variables:
  - cwnd: congestion window
  - flow\_win: flow window, i.e., receiver's advertised window
  - ssthresh: threshold size (used to update cwnd)For sending, use:  $\text{win} = \min(\text{flow\_win}, \text{cwnd})$
- Test the limit of the network
  - when  $\text{cwnd} < \text{ssthresh}$ , increase  $\text{cwnd}$  exponentially.
  - when  $\text{cwnd} \geq \text{ssthresh}$ , increase  $\text{cwnd}$  linearly (AI)
- Detect congestion by timeout:
  - If packet lost, have gone too far
  - reduce rate drastically (MD)
- Additive Increase, Multiplicative Decrease (AIMD)

# Additive Increase, Multiplicative Decrease

- Two TCP flows share a single bottleneck link.
- In an ideal design, they should get equal share of the bottleneck bandwidth, and total traffic is within the bottleneck capacity.



# Slow Start Phase

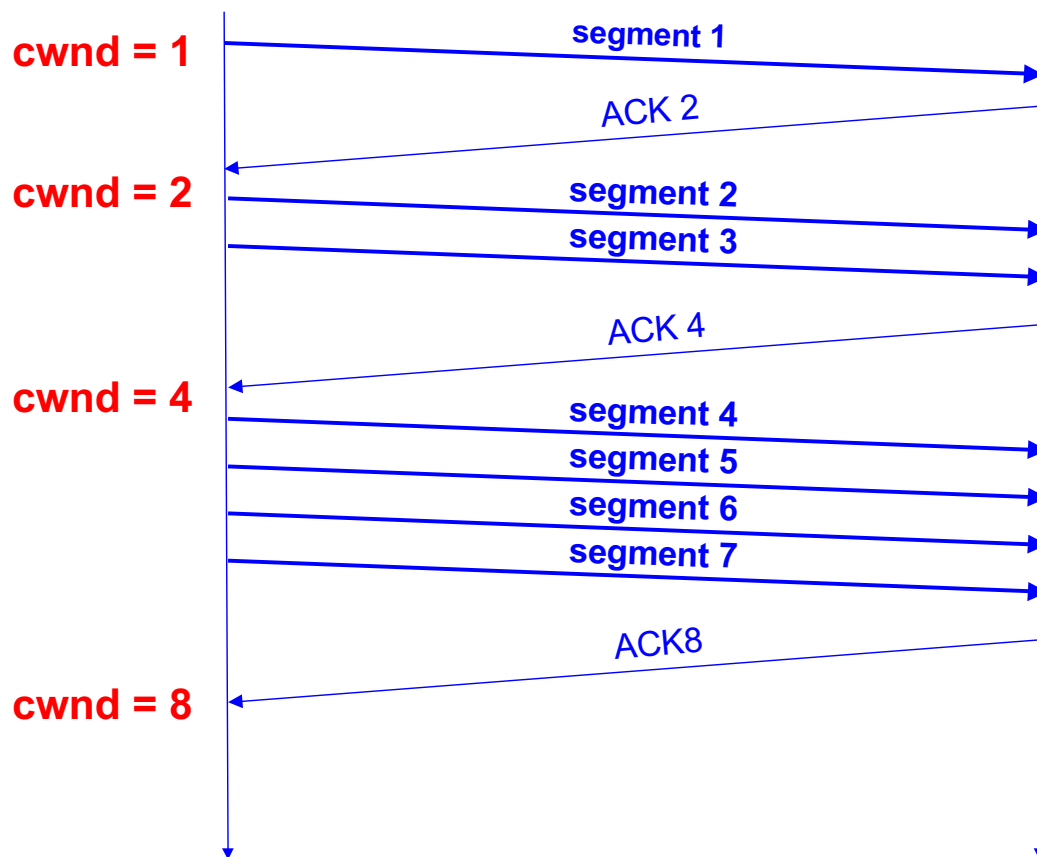
- Goal: reach the knee point quickly
- Upon (re)starting a TCP connection:
  - Set  $cwnd = 1$
  - Each time a segment is acknowledged increment  $cwnd$  by one,  
 $cwnd += 1$
- Slow Start is not actually slow
  - After all segments in the window have been acked,  $cwnd$  has been doubled over one RTT. That is,  $cwnd$  increases exponentially over time.
  - But slower than sending all at once.

# Slow Start Example

- The congestion window size grows rapidly

- TCP will slow down the increase of *cwnd* when

*cwnd*  $\geq$  *ssthresh*




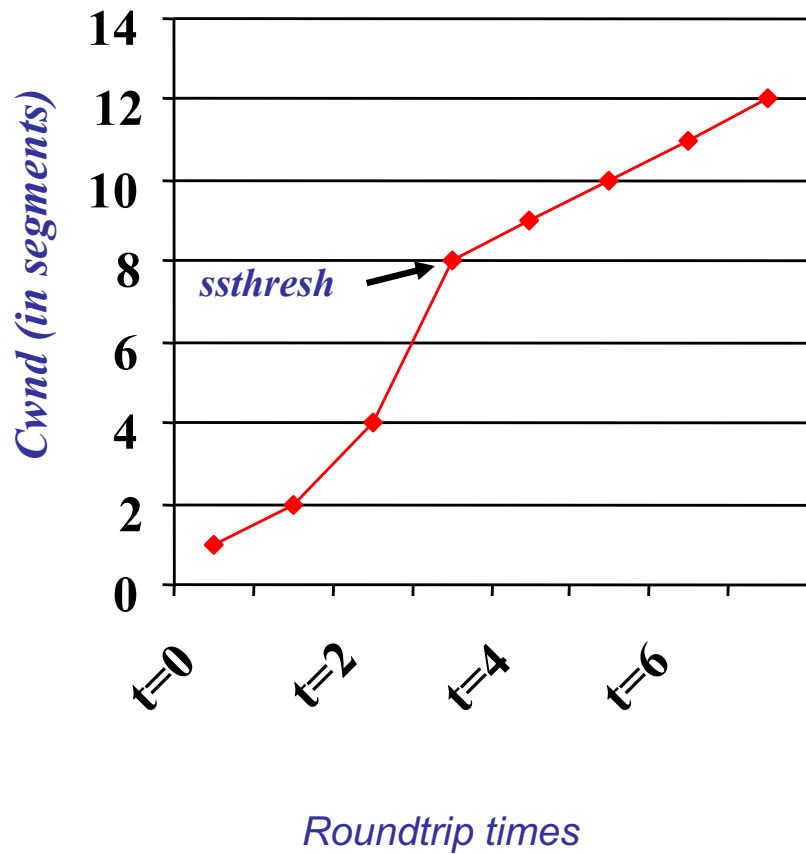
# Congestion Avoidance Phase

- Slow down “Slow Start”
- *ssthresh* is lower-bound guess about location of knee
- **If  $cwnd \geq ssthresh$  then**  
    each time a segment is acknowledged  
    increment *cwnd* by  $1/cwnd$ , i.e,  $cwnd += 1/cwnd$
- The effect is that, after *all* segments in the window have been acknowledged (i.e., one RTT), *cwnd* has been incremented by 1.



# Example

- Assume that  $ssthresh = 8$



# Retransmission Timeout (RTO)

- Upon a retransmission timeout
  - Assume packet loss due to congestion
  - Cut window size drastically

$ssthresh = cwnd/2;$   
 $cwnd = 1;$

# Putting Everything Together

## Initially:

```
cwnd = 1;  
ssthresh = rcvWin;
```

## New ack received:

```
if (cwnd < ssthresh)  
    /* Slow Start */  
    cwnd = cwnd + 1;  
else  
    /* Congestion Avoidance */  
    cwnd = cwnd + 1/cwnd;
```

## Timeout:


```
/* Multiplicative decrease */  
ssthresh = cwnd/2;  
cwnd = 1;
```

```
while (next < unack + win)  
    transmit next packet;
```

```
where win = min(cwnd, rcvWin);
```

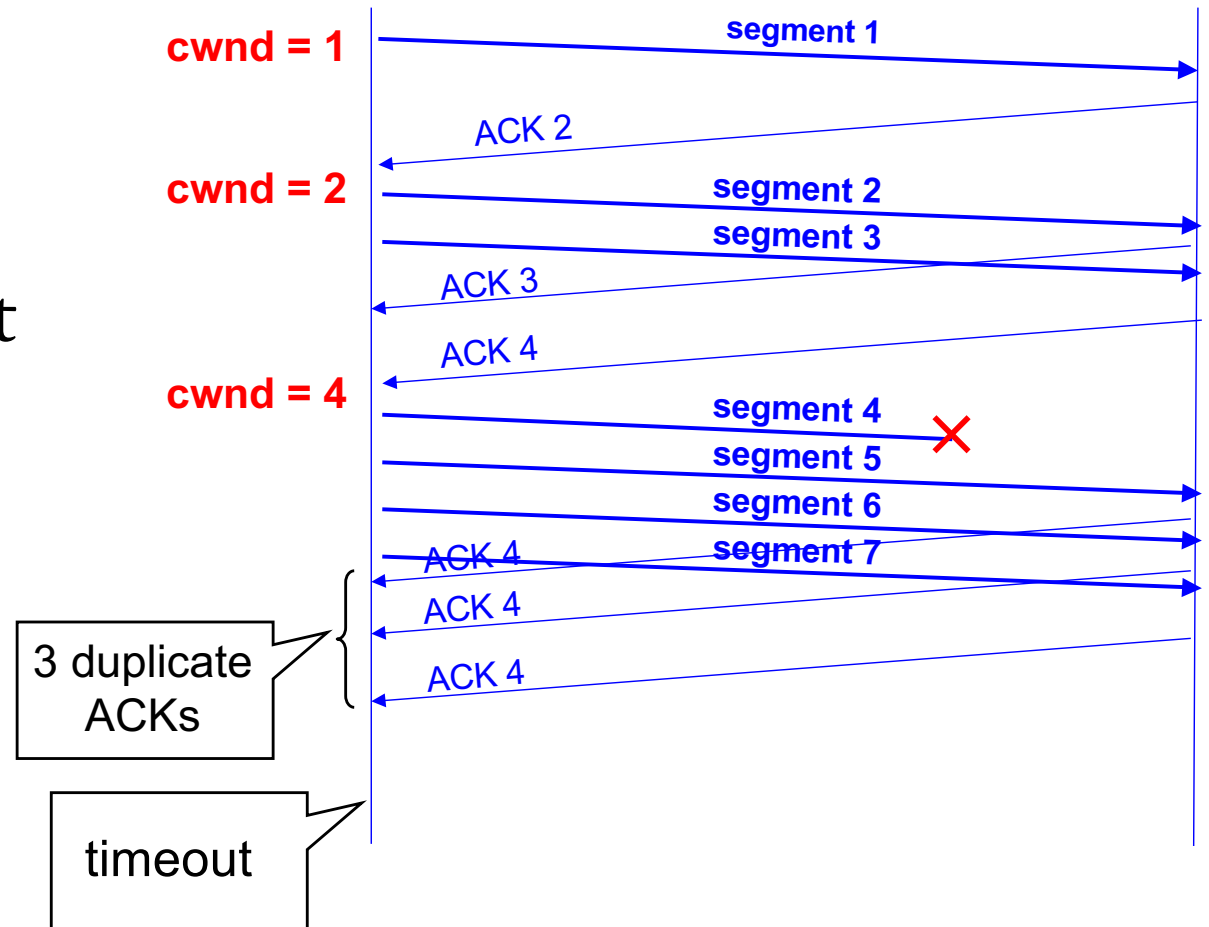


# TCP Reno behavior



# Fast Retransmission

- Don't wait for window to drain
- Resend a segment after 3 duplicate ACKs



# Fast Recovery

- After a fast retransmission, set  $cwnd$  to  $cwnd/2$ 
  - i.e., don't reset  $cwnd$  to 1
- But when RTO expires still do  $cwnd = 1$
- Fast Retransmit and Fast Recovery
  - Implemented by TCP Reno
  - widely used version of TCP today
- Lesson: avoid RTOs at all costs!

# Fast Retx and Fast Recovery

**Initially:**

```
cwnd = 1;  
ssthresh = recvWin;
```


**New ack received:**

```
if (cwnd < ssthresh)      /* slow start*/  
    cwnd = cwnd + 1;  
else                      /* congestion avoidance */  
    cwnd = cwnd + 1/cwnd;
```

**Loss detected:**

```
ssthresh = cwnd/2;        /* multiplicative decrease */  
if (3 duplicated ACKs)    /* fast retransmission */  
    cwnd = ssthresh;      /* fast recovery*/  
else  
    cwnd = 1;
```

# TCP NewReno Behavior (Fast Retransmit and Fast Recovery)



- Retransmit after 3 duplicated acks
  - prevent expensive timeouts
- No need to slow start again
- At steady state, *cwnd* oscillates around the optimal window size.



# Future Development

- Many more TCP congestion control schemes have been proposed and implemented.
  - Only end-hosts are involved, not routers
- There're also designs that require router support
  - Will be able to detect and react to congestion more accurately
  - But need deployment and support from routers
  - Usually only happens in special networks.