

CSC 525: Computer Networks

Sharing Information on the Internet

- The Internet contains a vast collection of information: documents, web pages, media, etc.
- One goal of the Internet is to make it easy to share this information.
- There are many different ways to do this.

Client/Server Model


- In the beginning, there was FTP.
 - Need to know the server address and the file path.
- Then Web was invented
 - Follow hyperlinks
- Now we use Google
 - Search
 - The trend is that users don't need to know where the content is; they can focus on defining, identifying, and consuming the content.

Peer-to-Peer Model

- Not all contents are put on dedicated servers
 - Require technical expertise to set up and maintain
 - Need powerful machines as servers
 - Exposed to the world, attract traffic, ...
- People want to share some contents, but don't want to commit to setting up and maintain a dedicated server.
- Or, we want all the participants to share the load (i.e., pool individual resource together to make the entire system works better).
- Peer-to-peer systems
 - A node can be a server and a client at the same time.


How Did It Start?

- A killer application: Napster: Free music over the Internet
- Key idea: share the content, storage and bandwidth of individual users
- Each node stores a subset of files
- Each node can access/download files from any other node in the system



Main Challenge

- Where to find a particular file?
 - Always a tradeoff between possible locations and search difficulty



Other Challenges

- Scaling: up to hundreds of thousands or millions of machines
- Adaptability: machines can come and go at any time
- Security and privacy
- ...

Napster

- A centralized directory system that maps files (songs) to machines that store the files.
 - Napster app uploads information of local files to Napster server.
- How to find a file (song)
 - Query the directory system → return the address of a machine that has the file
 - Ideally this is the closest/least-loaded machine
 - ftp the file
- Advantages:
 - Simple network operations
 - Easy to implement sophisticated search engines on top of the centralized index system
- Disadvantages:
 - Robustness, scalability, legal liability

Gnutella

- Distributed file location, no single server that knows all files.
- Idea: flood the search request
- How to find a file:
 - Send request to all neighbors
 - Neighbors recursively send the request to their neighbors
 - Eventually a machine that has the file receives the request, and it sends back the answer
- Which neighbors to connect (at overlay):
 - Initially randomly, may improve later on.
- Advantages:
 - Totally decentralized, highly robust
- Disadvantages:
 - Not scalable; the entire network can be swamped by request messages (to alleviate this problem, each request has a TTL)
 - Especially hard on nodes with slow links.

Kazza

- Modify Gnutella protocol into two-level hierarchy
 - Hybrid of Gnutella and Napster
- Supernodes
 - Nodes that have better connection to Internet
 - Act as temporary indexing servers for standard nodes
 - Help improve the stability of the network
- Standard nodes
 - Connect to supernodes and report their lists of files
 - Allow slower nodes to participate
- Search
 - Broadcast (Gnutella-style) search messages across supernodes

Distributed Hash Table (DHT)

- Abstraction: a hash-table data structure implemented over a distributed system.
 - insert(key, value);
 - value = query(key) or lookup(key)
 - key can be a file name, the title of a song, etc.
 - value can be anything: a data object, file, IP address of the node that has the file, etc.
 - Only exact match, very limited expressiveness of queries.
Applications are expected to build more functionality on top of DHT.
- Proposals
 - CAN, *Chord*, Pastry, Tapestry, and many others.

DHT Design Goals


- (key, value) pairs are stored distributedly at many hosts. Given a key, the DHT needs to locate which host stores the (key, value) pair.
- Make sure that an item is always found if it's in the DHT.
- Scales
 - Search takes $O(\log N)$ steps
- Handles rapid arrival and failure of nodes

Chord


- Use a one-dimensional id space $0..2^m-1$
- Use the same hash function to hash keys to ids.
- Use the same hash function to hash nodes' IP addresses to ids.
- Organize the nodes into a structure to facilitate fast search of an id/key.
- Design decision
 - Decouple correctness from efficiency
- Properties
 - Routing table size $O(\log N)$, where N is the total number of nodes
 - Guarantees that a file is found in $O(\log N)$ steps

The Ring

- Model the ID space as a ring.
 - Each node has an id = $\text{hash}(\text{ip})$.
- A node maintains its
 - Successor
 - Predecessor
 - Finger table
- *A key is stored at the node with next higher ID.*
 - E.g., a filename can be converted to a key = $\text{hash}(\text{filename})$
 - Key 24 and 30 are stored in node 32.




Basic Lookup




- Linear search
 - Search the ring clockwise one node at a step.
 - $O(N)$ hops, with only one routing entry (successor) per node.
- Direct search
 - If every node maintains a list of every other nodes, the search can be done in 1 step, but the routing table size is N .
- Anything in the middle?

Finger Table



- Some hints to speed up the lookup
- $\text{Finger}[k]$: successor to $(n + 2^{(k-1)}) \bmod 2^m$

Scalable Lookup



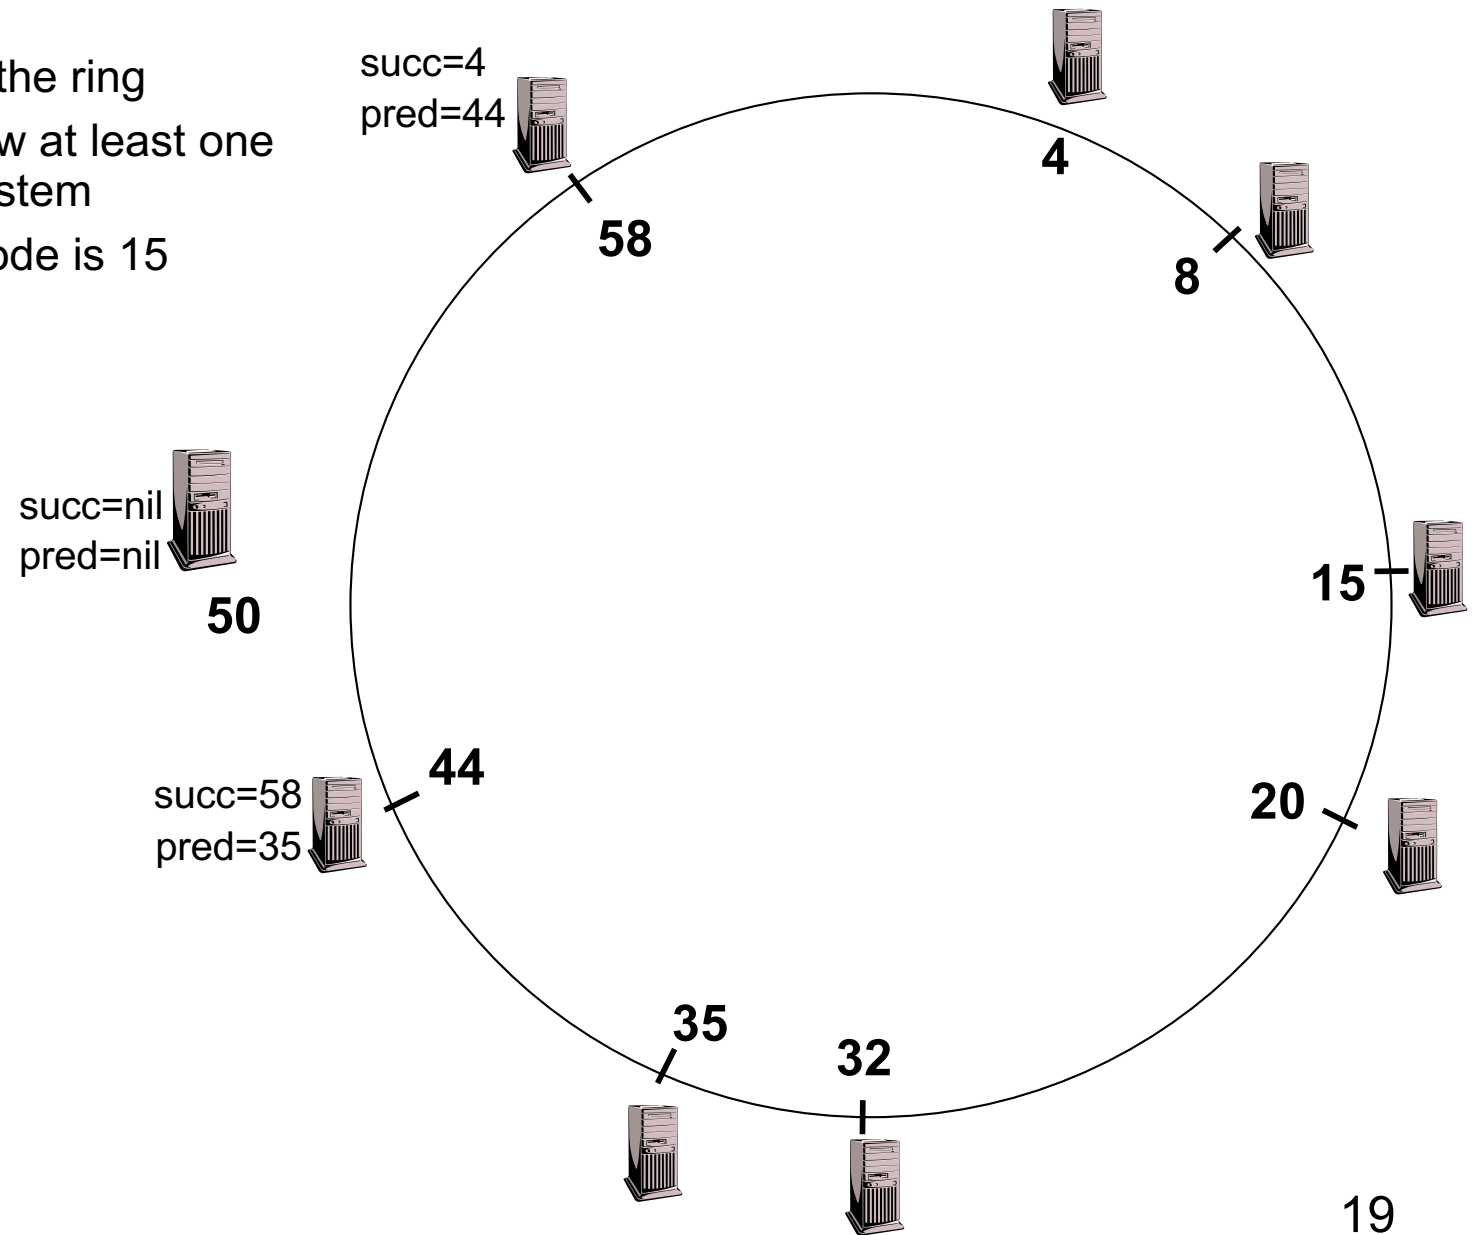
- Binary search using
 - Route to the node according to the finger table.
 - Each hop cuts the distance to target by half.
 - $O(\log N)$ hops with $O(\log N)$ routing table entries.

Join Operation

- Each node A periodically sends a **stabilize()** message to its successor B
- Upon receiving a **stabilize()** message, node B
 - returns its predecessor $B' = \text{pred}(B)$ to A by sending a **notify(B')** message
- Upon receiving **notify(B')** from B,
 - if B' is between A and B, A updates its successor to B'
 - Otherwise ($B' \neq A$), A doesn't do anything.

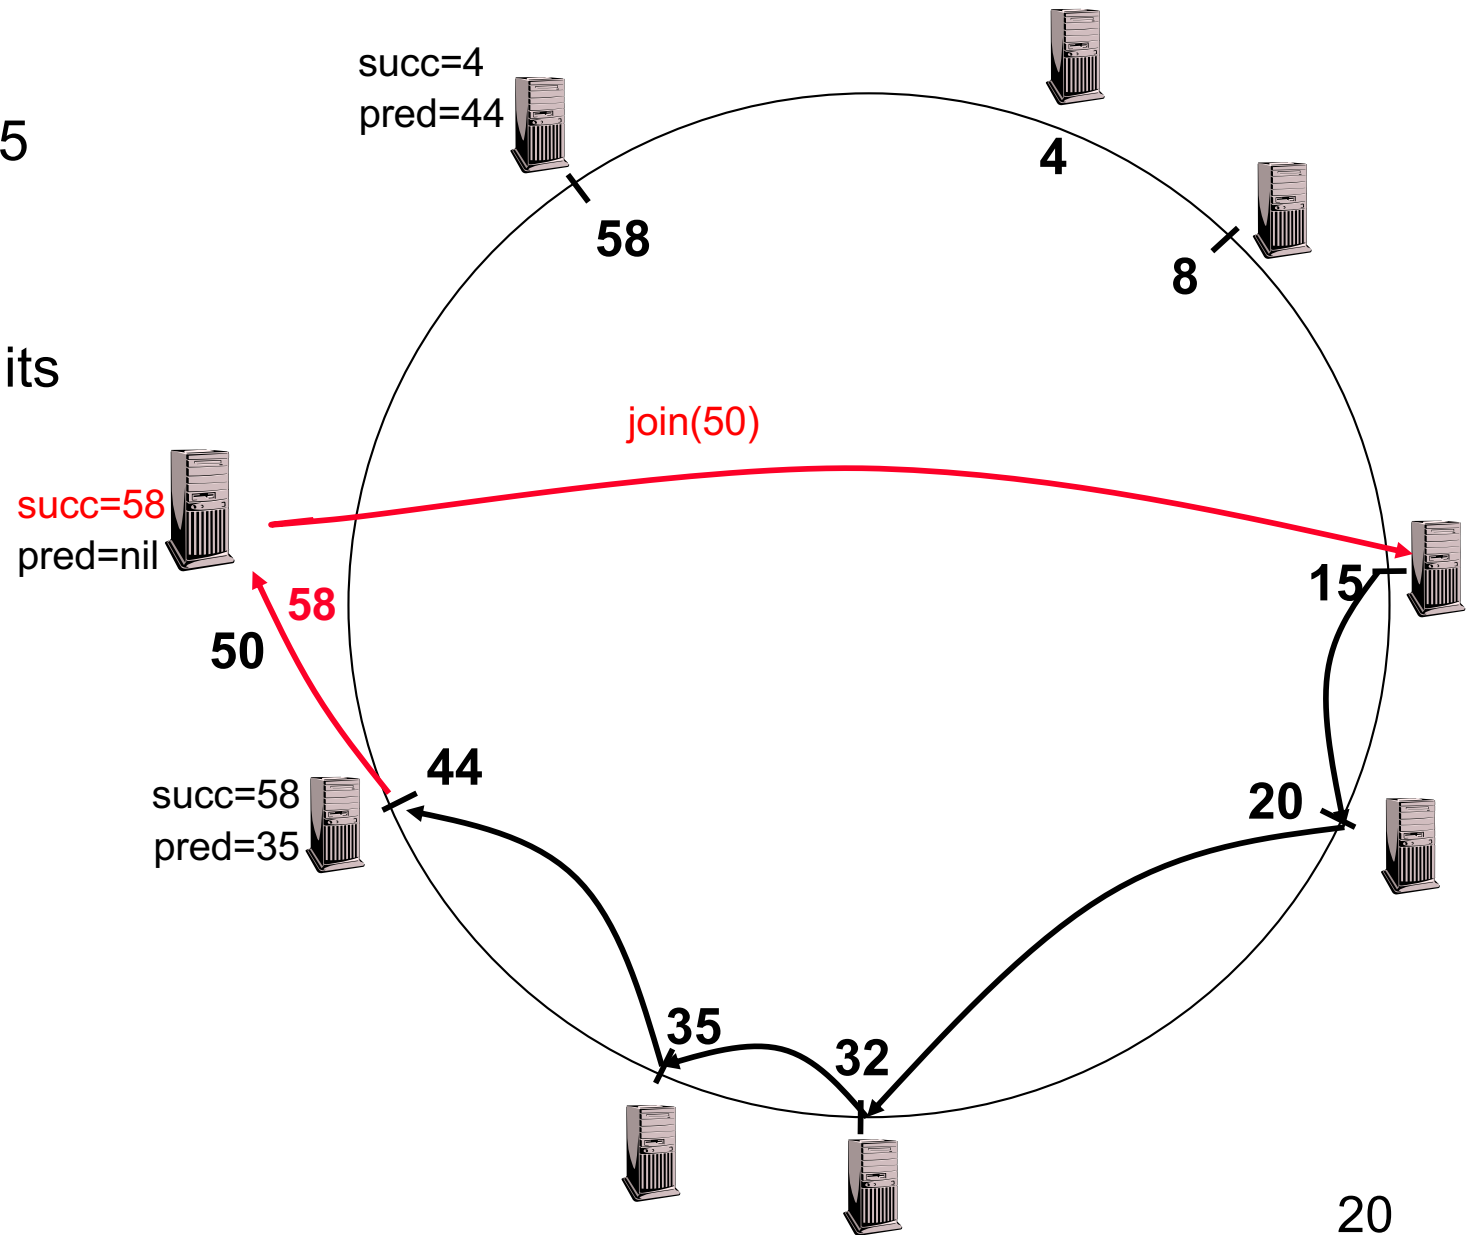
Joining Operation

- Node with id=50 joins the ring
- Node 50 needs to know at least one node already in the system
 - Assume known node is 15



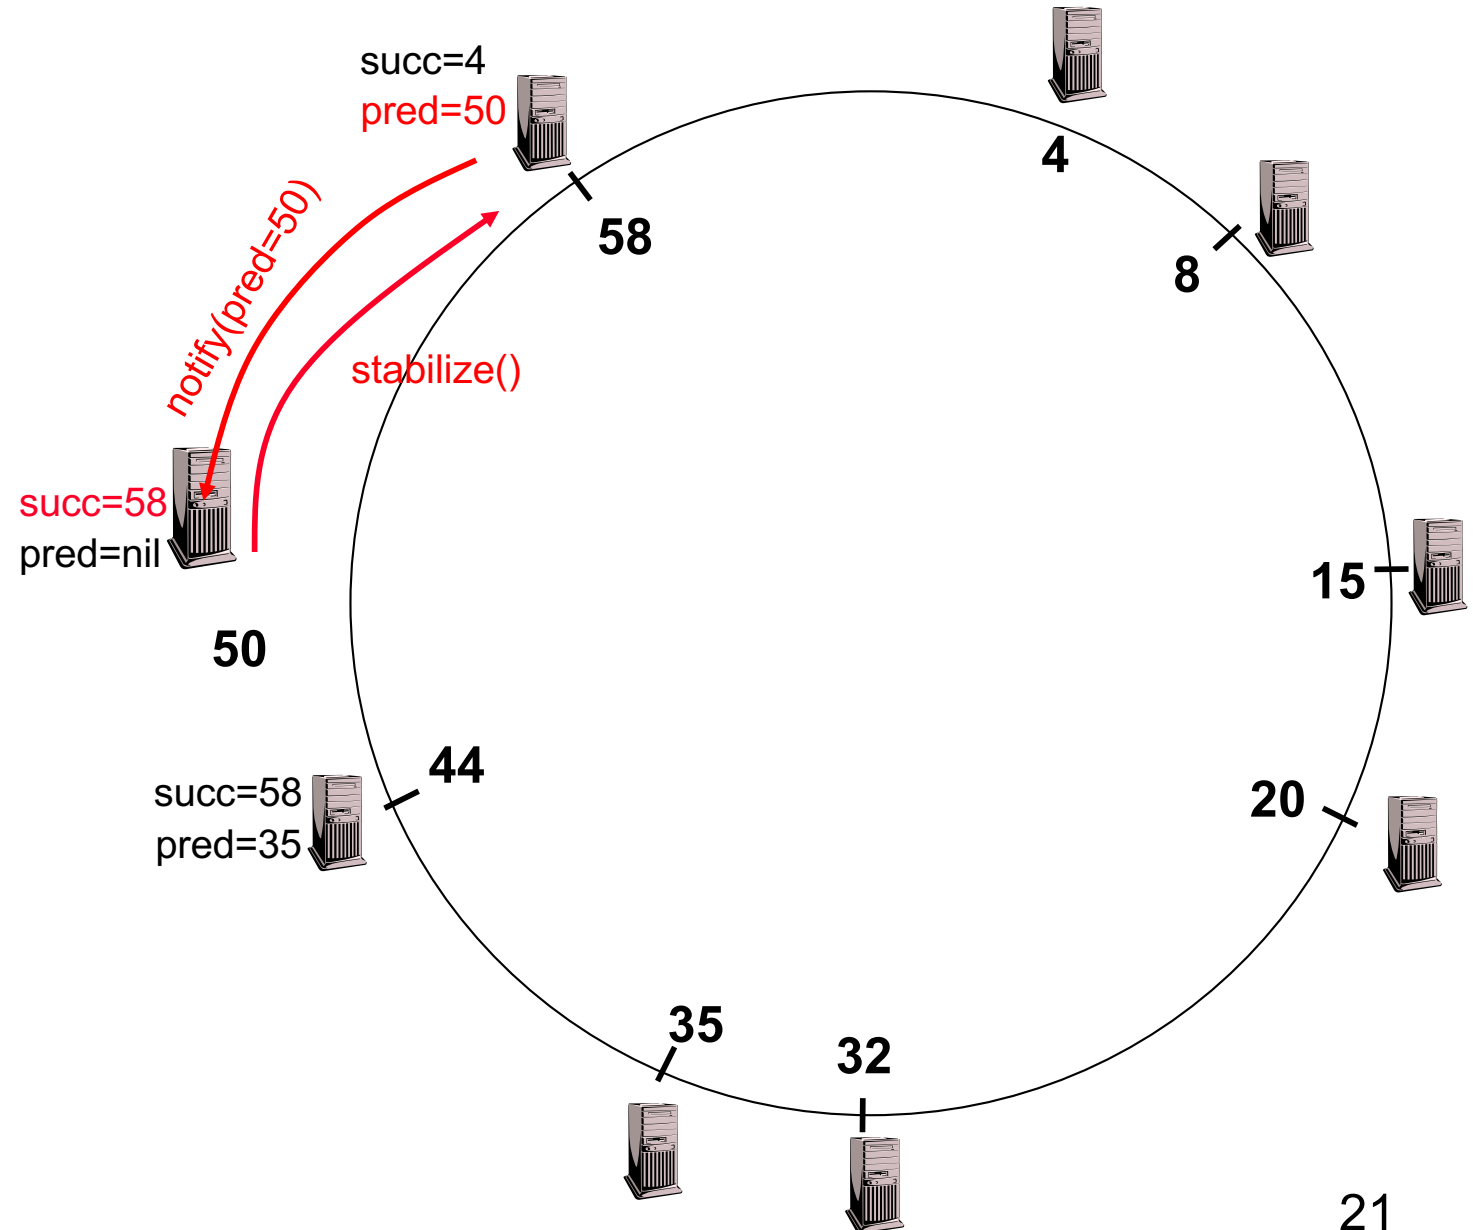
Joining Operation

- Node 50: send join(50) to node 15
- Node 44: returns node 58
- Node 50 updates its successor to 58



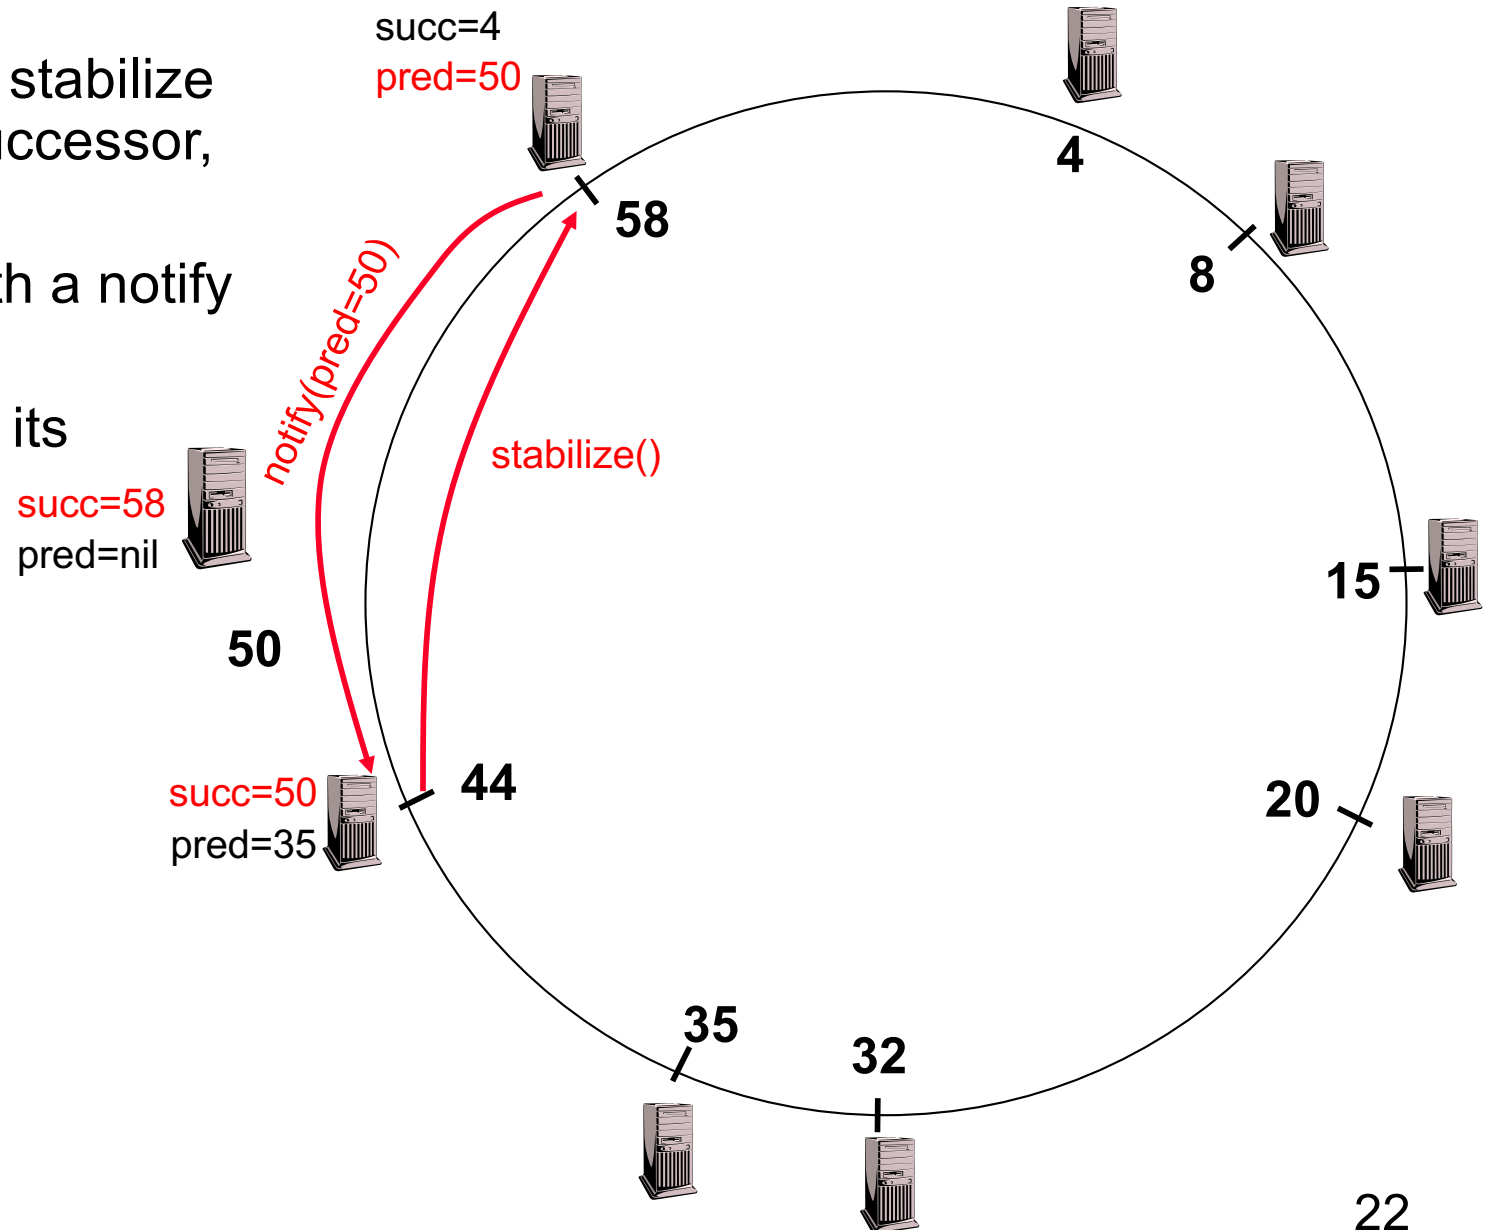
Joining Operation

- Node 50: send stabilize() to node 58
- Node 58:
 - update predecessor to 50
 - send notify() back



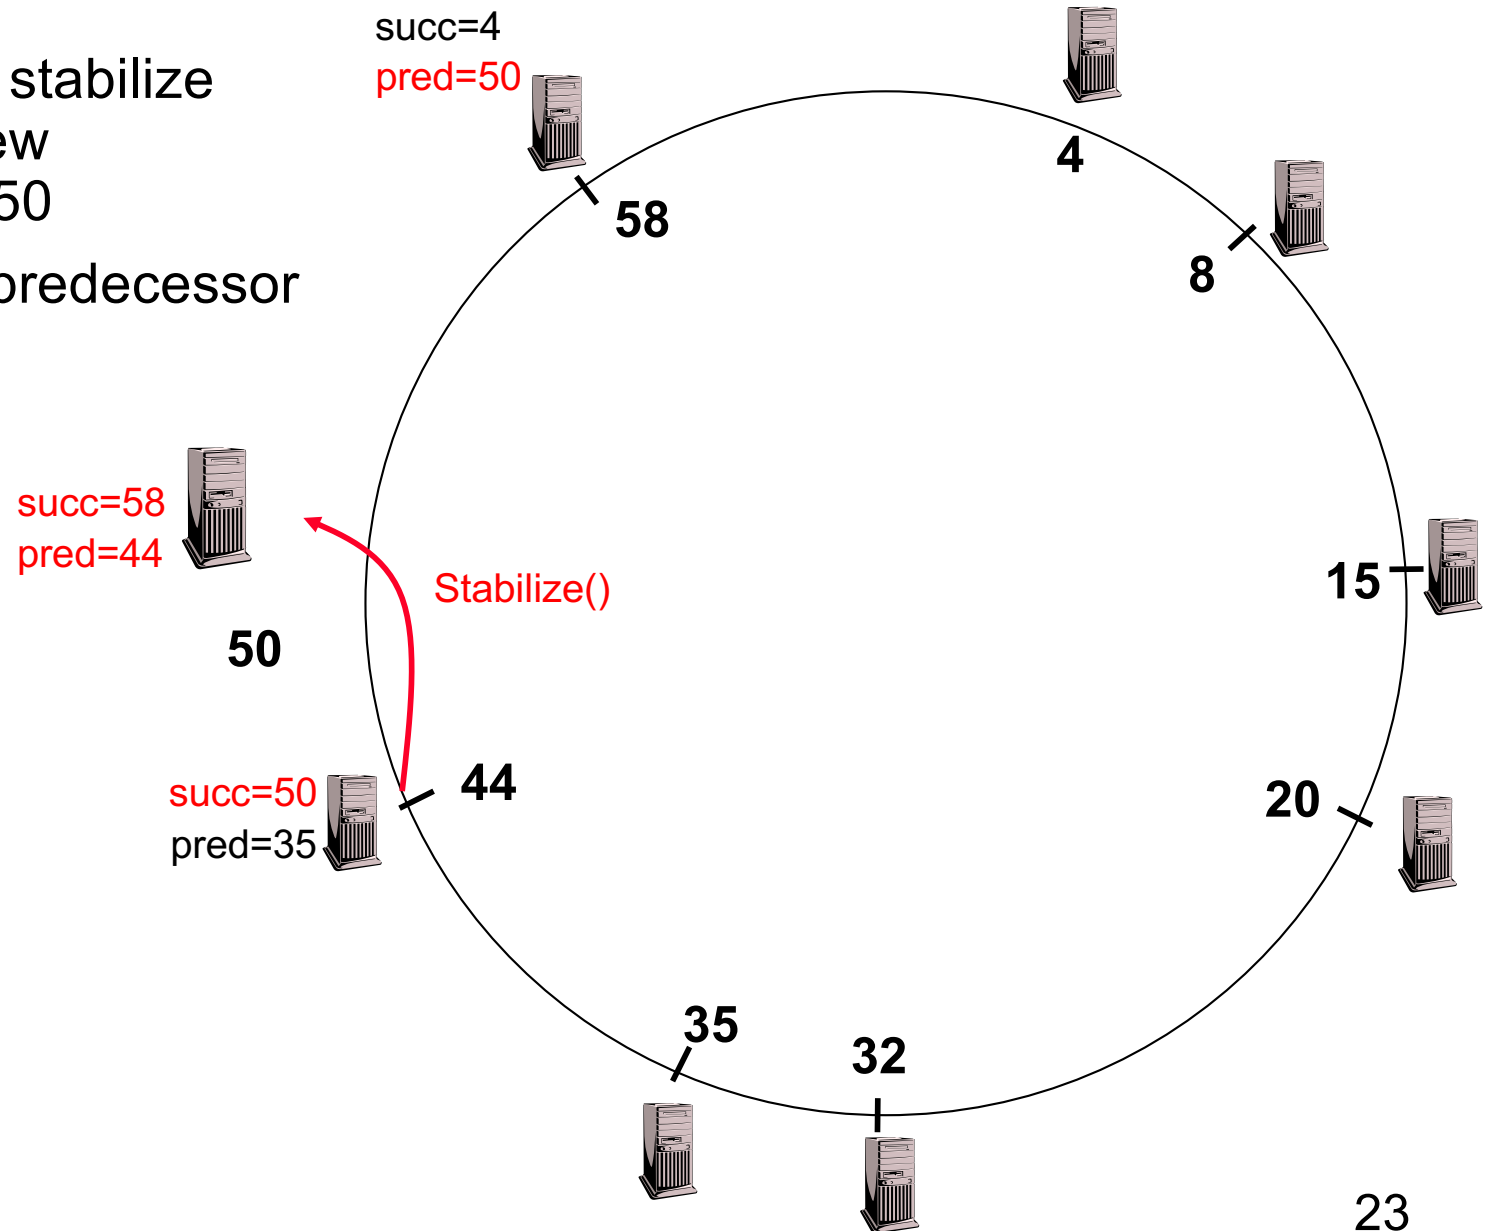
Joining Operation (cont' d)

- Node 44 sends a stabilize message to its successor, node 58
- Node 58 reply with a notify message
- Node 44 updates its successor to 50




Joining Operation (cont' d)

- Node 44 sends a stabilize message to its new successor, node 50
- Node 50 sets its predecessor to node 44



Joining Operation (cont'd)

- This completes the joining operation!



Achieving Robustness

- To improve robustness each node maintains the k (> 1) immediate successors instead of only one successor
- In the notify() message, node A can send its $k-1$ successors to its predecessor B
- Upon receiving notify() message, B can update its successor list by concatenating the successor list received from A with A itself

Improve Performance

- Overlay hops \neq Network Latency
 - Chose finger that reduces expected time to reach destination
 - Chose the closest node from range $[N+2^{i-1}, N+2^i)$ as successor
- Accommodate heterogeneous systems
 - Multiple virtual nodes per physical node
 - Better load balancing, but more complex.

Conclusions

- Distributed Hash Tables are a key component of scalable and robust overlay networks
 - Chord: $O(\log n)$ state, $O(\log n)$ distance
- A single function: find an item given a key (exact match)
- What can DHT do for us?
 - Distributed object lookup
 - Decentralized file system
 - Application layer multicast
 - Persistent storage
 - ...