

# CSC 525: Computer Networks

# Overlay Networks

- Overlay network
  - A network built on top of an existing network
  - Add an additional layer of indirection/virtualization to implement a new service
  - Rely on the underlay network for basic services
- Example: Internet is an overlay network
  - goal: connect local area networks
  - built on local area networks (e.g., Ethernet), phone lines
  - add an Internet Protocol header to all packets
  - Between two IP routers, there can be multiple layer-2 devices, but IP routers are not aware of them.
- Today overlay networks often refer to application level networks over the IP network.

# Benefits

- Easier deployment
  - Don't have to modify underlying hardware or software
  - Only need to deploy new software on top of existing software
  - e.g., adding IP on top of Ethernet does not require modifying Ethernet protocol or driver
- Better separation of functionality

# Costs

- Overhead
  - Adds a layer in protocol stack
    - Additional packet headers, processing
  - The overlay doesn't have enough knowledge of underlying network to optimize performance.
  - Sometimes, certain work is redundant across the layers.
- Complexity
  - Layering does not eliminate complexity, it only manages it
  - More layers of functionality → more possible unintended interaction between layers

# Applications

- File sharing, Multicast, Unicast, Security and Privacy, Mobility, etc.
- We'll cover a number of systems. Pay attention to their differences and commonality.
  - E.g., what overlay topology does each system use, and how do they build and maintain such topologies?

# Overlay Multicast


- Provide multicast functionality above the IP layer: **application layer multicast**
- Challenge: do this efficiently and scalable
- Narada
  - Support multi-source multicast
  - Involve only end hosts
  - Small groups  $\leq$  a few hundreds of nodes

# Concerns with IP Multicast

- Scalability with the number of groups
  - Routers maintain per-group state
  - Aggregation of multicast addresses is not very effective.
- Supporting higher level functionality is difficult
  - IP Multicast: best-effort multi-point delivery service
  - End systems responsible for handling higher level functionality
    - Reliability, congestion control, security, address allocation, charging and accounting, etc.
- Deployment is difficult and stagnant (at inter-domain).

# End System Multicast


- Hosts forward packets, without router support
- Build a smarter forwarding tree
  - Repetitive Unicast doesn't scale
  - Duplicate transmission and extra delay inevitable, but should be minimized.





# End System Multicast

- Hosts forward packets, without router support
- Build a smarter forwarding tree
  - Repetitive Unicast doesn't scale
  - Duplicate transmission and extra delay inevitable, but should be minimized.





# Potential Benefits



- Routers don't maintain per-group state
  - End systems do, but they participate in much fewer groups
- Much easier to deploy
- Potentially simplify support for higher level functionality
  - Leverage computation and storage of end systems
  - For example, for buffering packets, transcoding, ACK aggregation
  - Leverage solutions for unicast congestion control and reliability

# What is a good multicast overlay?



- Conflicting goals: shorter delay and small node degree
- Challenge: End hosts don't know the underlying topology



- Cannot directly run a conventional multicast routing protocol



- Need trim the complete graph first



# Narada Design

**“Mesh”**: Richer overlay topology that includes all group members, but not a complete graph.

## Step 1

- Members have low degrees
- Shortest path delay between any pair of members along mesh is small

## Step 2

- Construct per-source trees using well known multicast routing algorithms (reverse path checking)
  - Members have low degrees
  - Small delay from source to receivers

# Narada Components

- Mesh Management:
  - Ensure mesh remains connected in face of membership changes
- Mesh Optimization:
  - Distributed heuristics for ensuring path delay between members along the mesh is small
- Forwarding tree construction:
  - Routing algorithms for constructing data-delivery trees
  - Distance vector routing, and reverse path forwarding


# Mesh Management

- Every node maintains a list of **all** the group members
- Periodically exchange the list with neighbors
- A new node gets a partial list of existing members by an out-of-band mechanism.
- Need to remember dead nodes for a while.



# Repairing Mesh Partition

- Partition is detected when some member entries are not refreshed as expected
- Probe these members by probability
  - Either they're dead, or establish another virtual link to repair the partition.



# Improve Mesh Quality

- Initially a new node randomly chooses some existing members to connect to.
- Over time, add or drop virtual links to improve the mesh quality.
- Keep the node degree under a threshold.

# Add Overlay Links

- Node  $i$  adds link  $i$ - $j$  if  $utility(j)$  is greater than a threshold.

---

**EvaluateUtility** ( $j$ ) begin

$utility = 0$

for each member  $m$  ( $m$  not  $i$ ) begin

$CL$  = current latency between  $i$  and  $m$  along mesh

$NL$  = new latency between  $i$  and  $m$  along mesh

if edge  $i$ - $j$  were added

if ( $NL < CL$ ) then begin

$utility += \frac{CL - NL}{CL}$

end

end

return  $utility$

---

# Drop Overlay Links

- Node  $i$  drops link  $i$ - $j$  if *consensus cost* ( $j$ ) is less than a threshold.

---

```
EvaluateConsensusCost( $j$ ) begin
   $Cost_{ij}$  = number of members for which  $i$  uses  $j$  as
             next hop for forwarding packets.
   $Cost_{ji}$  = number of members for which  $j$  uses  $i$  as
             next hop for forwarding packets.
  return  $\max(Cost_{ij}, Cost_{ji})$ 
end
```


---

# Multicast Routing

- Similar to DVMRP
  - Run distance vector protocol on top of the mesh
  - Use RPF to derive the per-source forwarding tree
- Leverage existing routing protocol.


# Desirable properties of heuristics

- **Stability:** A dropped link will not be immediately added back
- **Partition Avoidance:** A partition of the mesh is unlikely to be caused as a result of any single link being dropped




Delay improves to Stan1, CMU  
but marginally.

**Do not add link!**




Delay improves to CMU, Gatech1  
and significantly.

**Add link!**



Used by Berk1 to reach only Gatech2 and vice versa.


Drop!!



An improved mesh

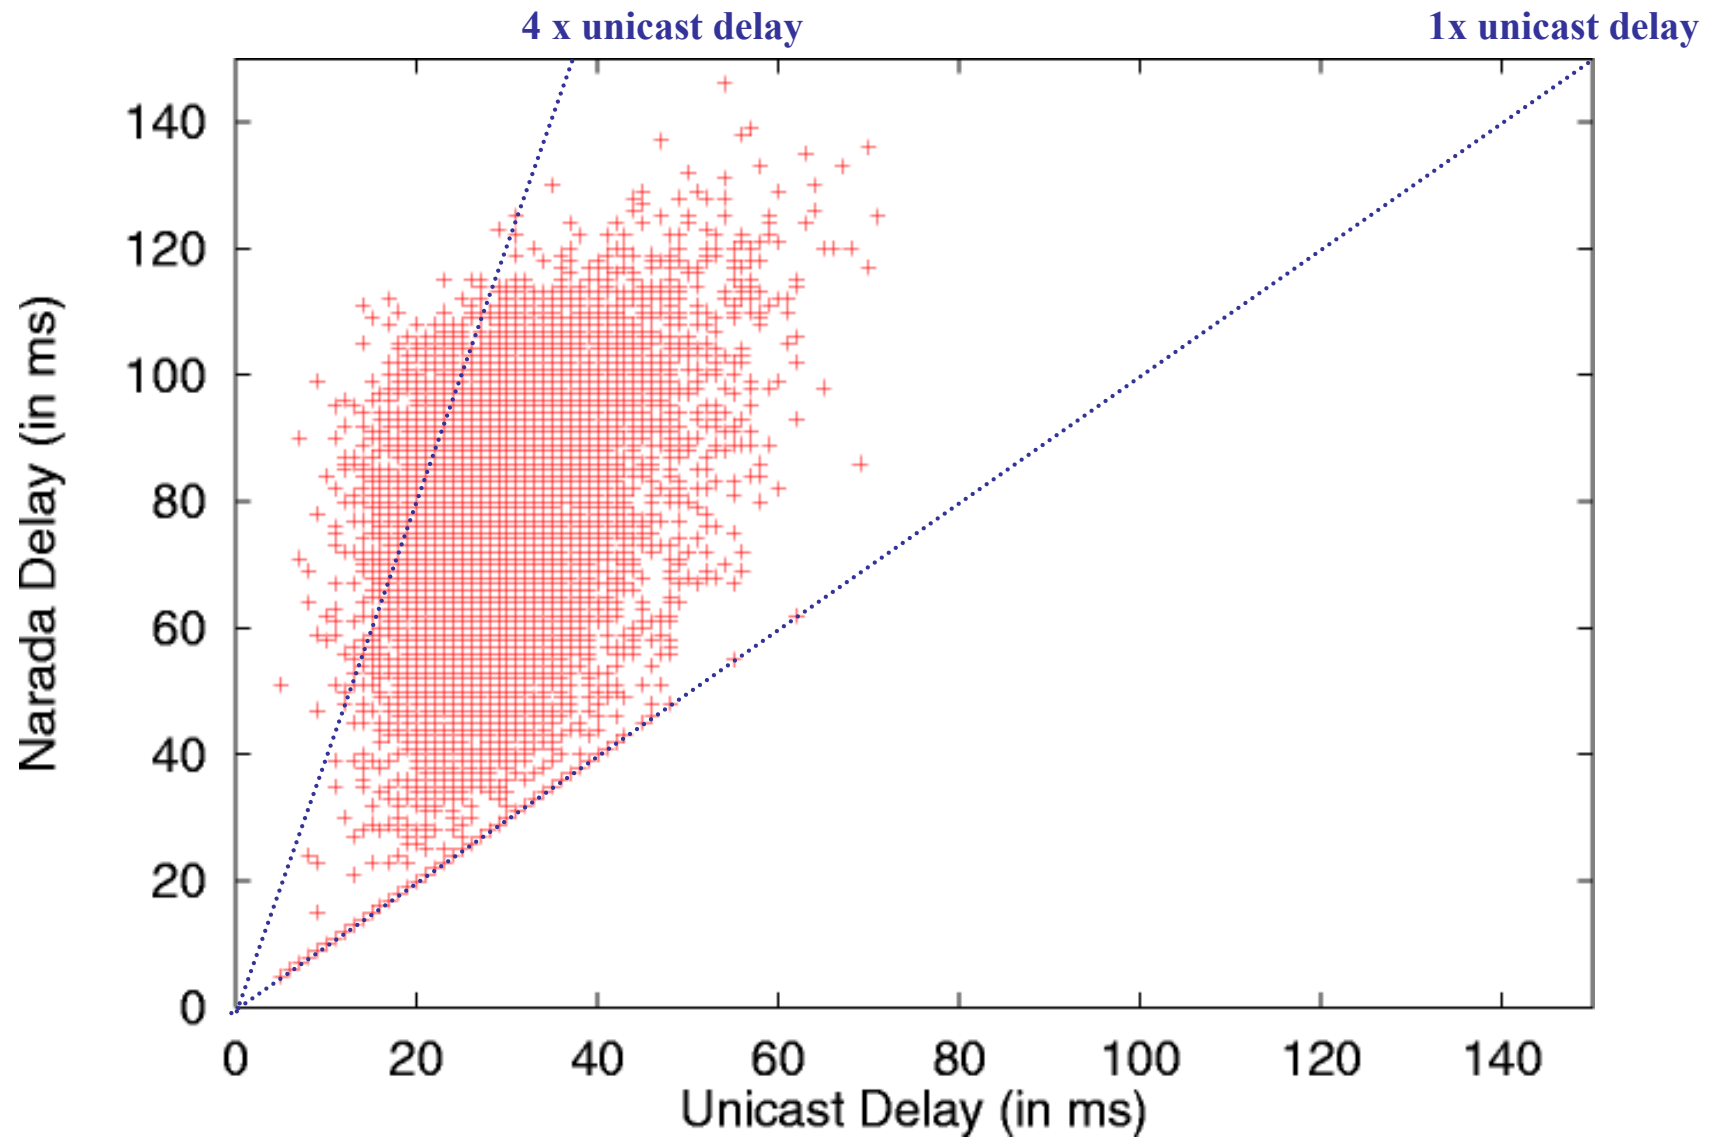
# Member Join

- Mesh improves over time
- But still changes after 40 minutes






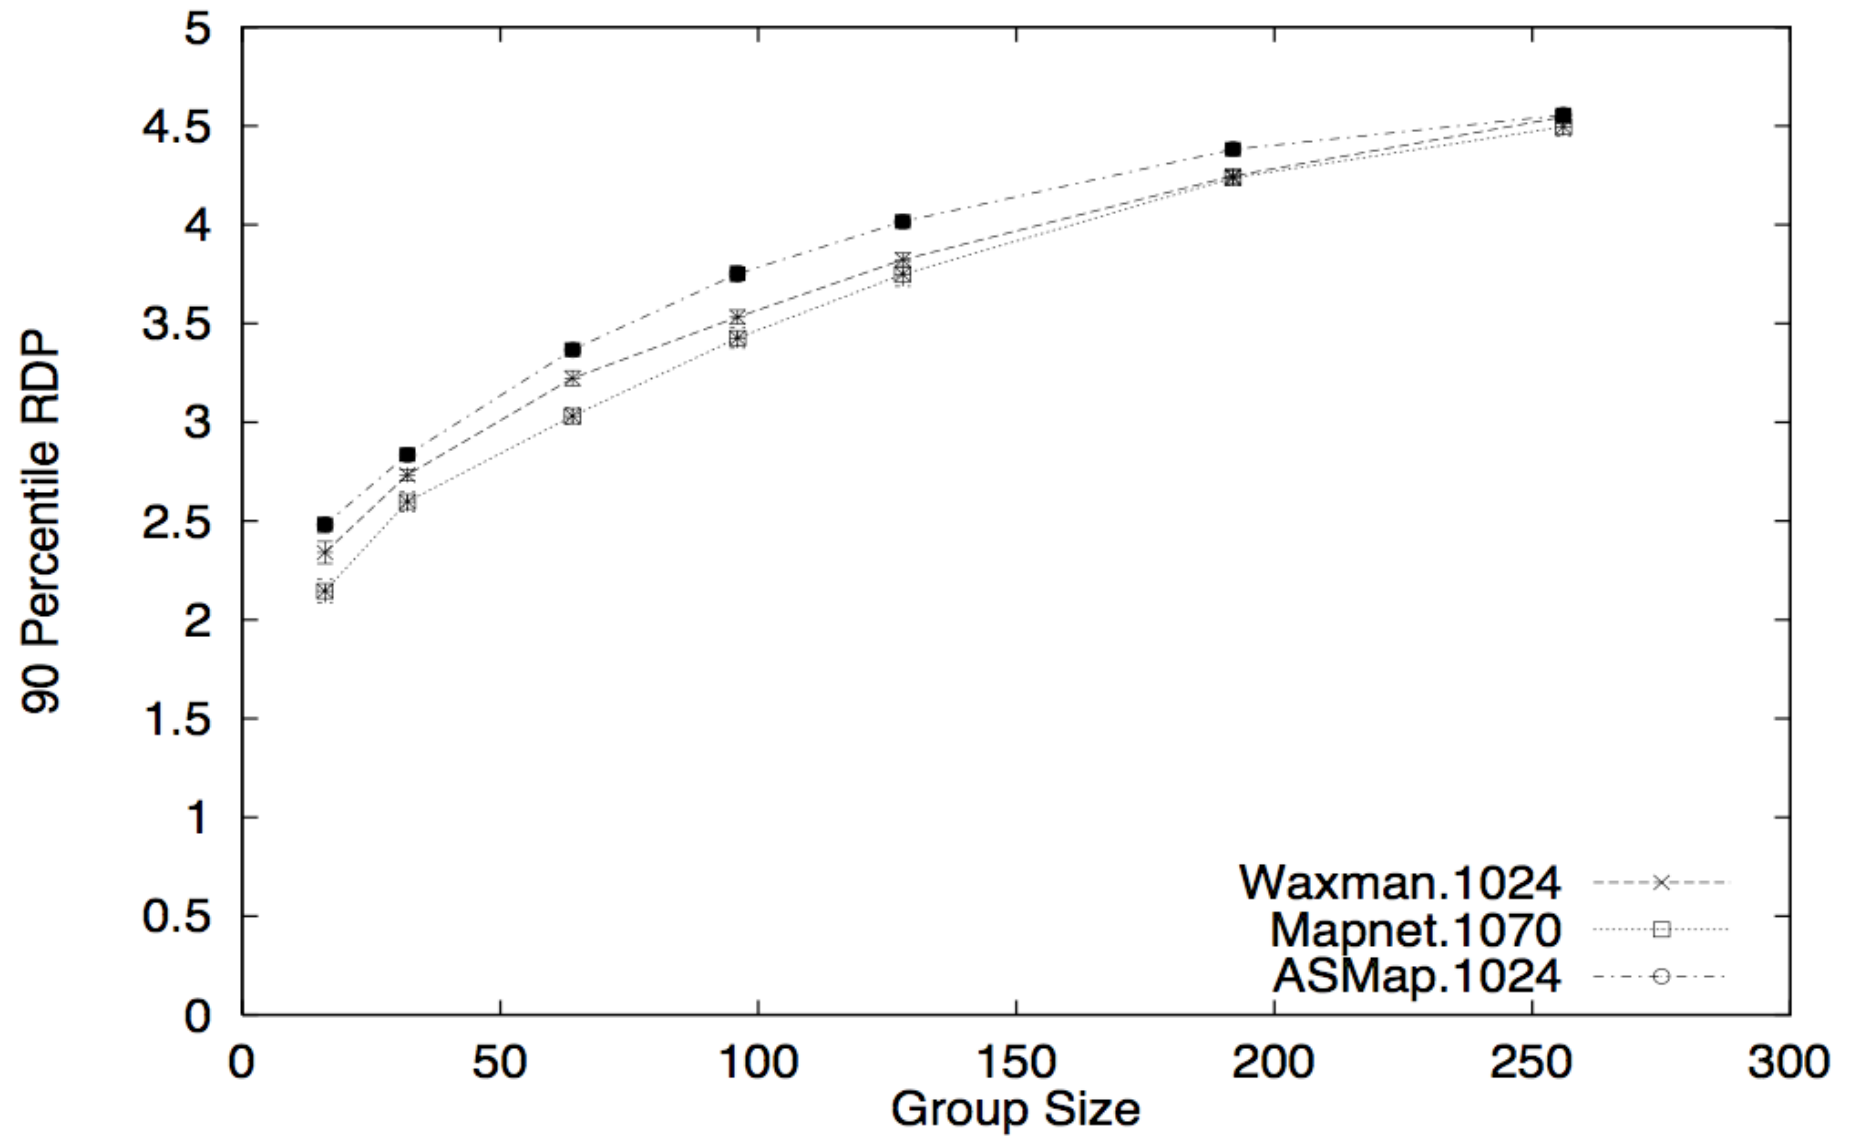
# Delay



# Link Stress




# Group Size



# Internet Evaluation

- 13 hosts, all join the group at about the same time
- No further change in group membership
- Each member tries to maintain 2-4 neighbors in the mesh
- Host at CMU designated source



# Conclusions

- For small-sized groups, an end-system overlay approach
  - is feasible
  - has a low performance penalty compared to IP Multicast
  - has the potential to simplify support for higher layer functionality
  - allows for application-specific customizations
- Where to implement multicast
  - IP layer or Application layer or both ?